

Engineering Simulations in the Age of Data

Gianluca Iaccarino

Mechanical Engineering & Institute for Computational Mathematical Engineering
Stanford University

jops@stanford.edu



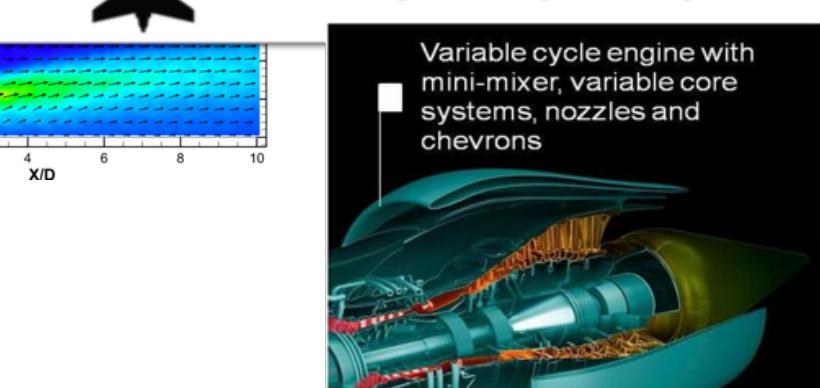
The Age of Data



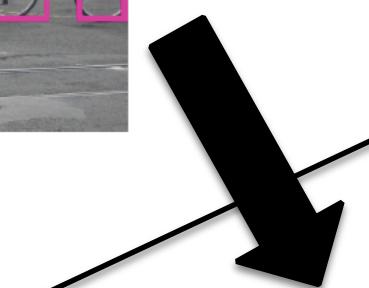
NASA CRM, AoA = 18°
 $M_{inf} = 0.25$, $Re = 11M$
Instantaneous near-wall velocity
DLR companion expts.

Sensor data from a cross-country flight

20 TB \times 2 \times 6 \times 28,537 \times 365
20 iterates of information per engine every hour
twin-engine Boeing 737
six-hour, cross-country flight from New York to Los Angeles
of commercial flights in the sky in the United States on any given day.
 $= 2,499,841,200 \text{ TB}$



Data



Software for Data

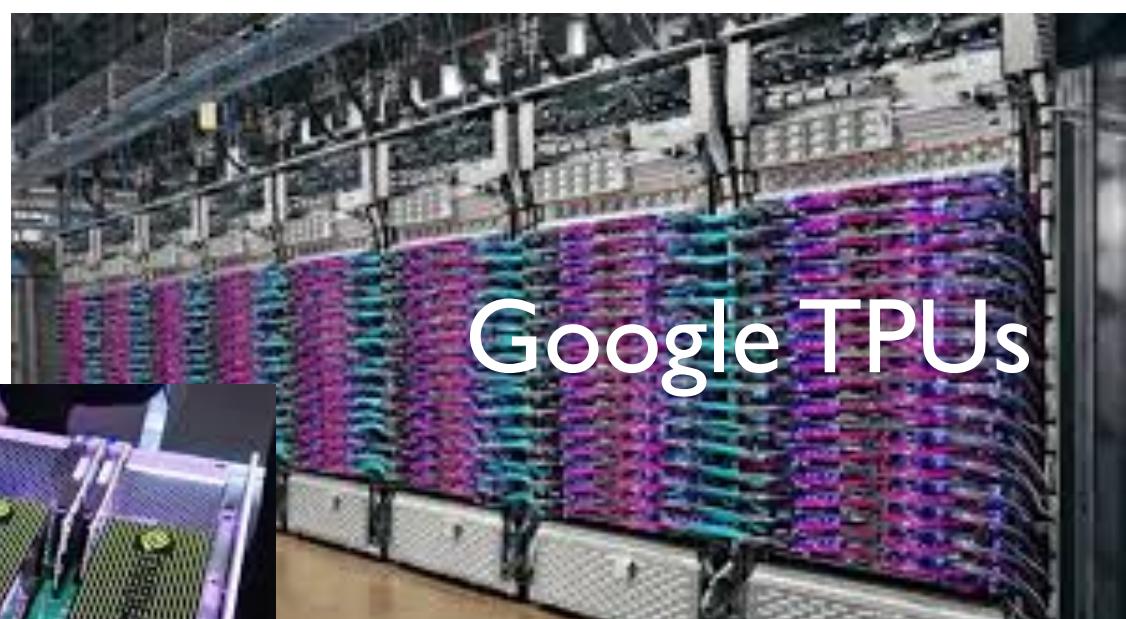


TensorFlow

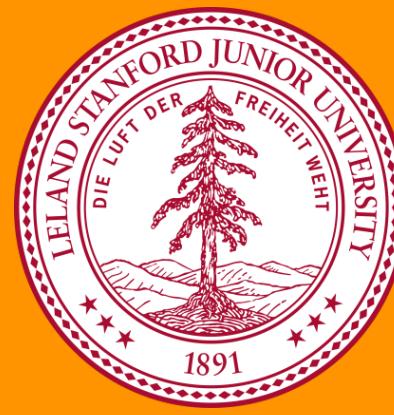
K Keras

PYTORCH

Computing for Data



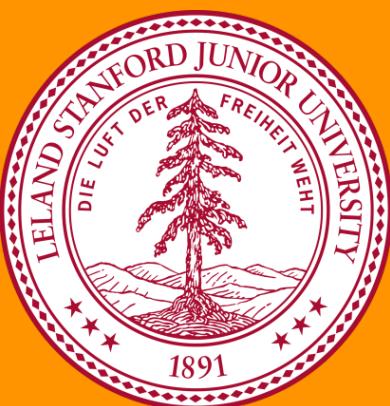
Google TPUs



Engineering Simulations

...or...

the Science of Predictions



Weather Predictions

C. Abbe, *Physical Basis of Long Range Weather Forecast, 1901*

DECEMBER, 1901.

MONTHLY WEATHER REVIEW.

551

24th, both preceded by heavy swell and followed by low dew-point. The north wind evidently precipitated the terrific downpour on north Hilo. There was lightning reported from Hawaii for the 8th, 9th, 13th, 24th, and 25th, and Maui on the 8th. Snow fell on Mauna Kea and Mauna Loa on the 8th and 24th; on Haleakala on the 8th or 9th. Earthquake reported at Hilo, 7:30 p. m. on the 2d.

AN AURORAL-LUNAR HALO DISPLAY.

IN BROECK, Braidentown, Fla., dated December 29, 1901.

t of the 28th I observed an auroral display with banding about west-northwest to east-southeast. At the time an unusually brilliant halo around the moon was 15° in diameter. The upper half was fringed on its outer edge by rays about 3° or 4° long, a few much longer, radiating from the moon (the center of the halo), but from it about east-southeast on the horizon. Some

THE PHYSICAL BASIS OF LONG-RANGE WEATHER FORECASTS.¹

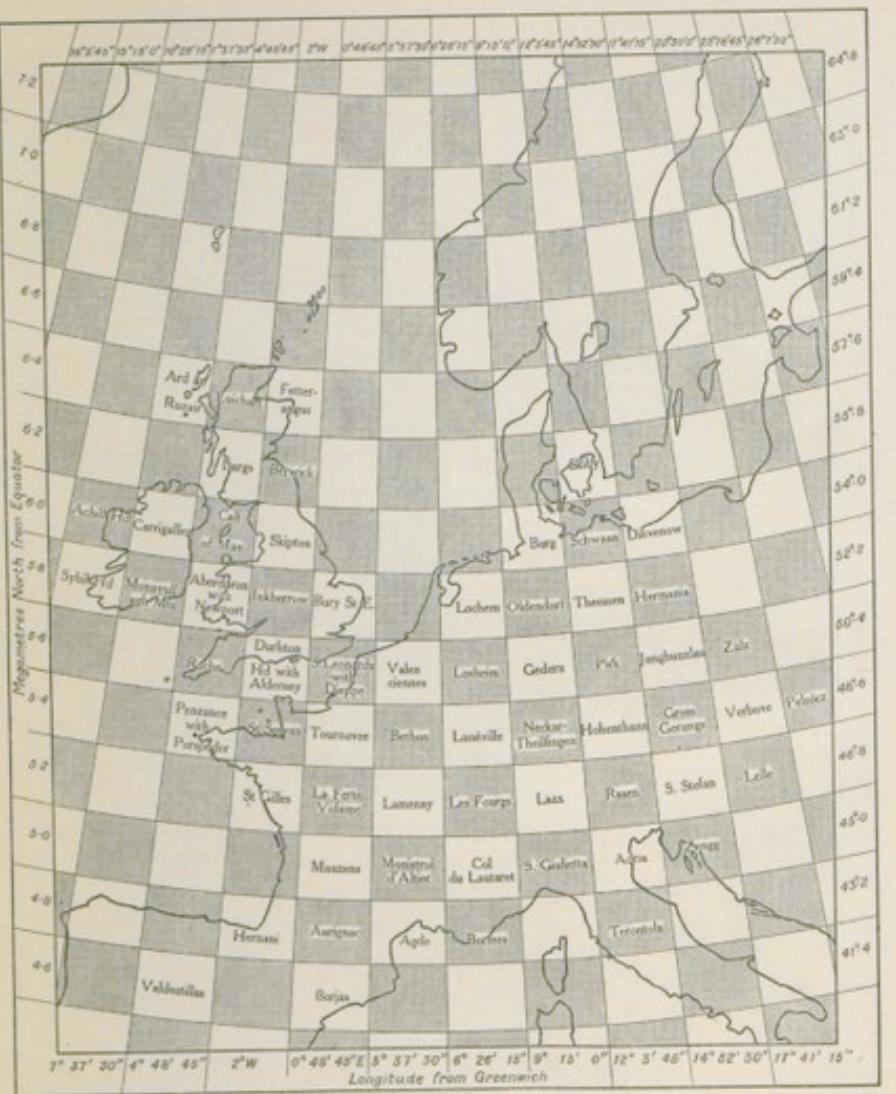
By PROF. CLEVELAND ABBE.

The expression "long range" must not be misunderstood. It refers only to the length of time intervening between the date of making a weather prediction and the date when we expect it to be fulfilled. At the present time, by the help of the daily weather map, the official weather forecasters of this country, and indeed of every civilized nation on the globe, publish forecasts, in detail, of approaching weather changes, and especially storms, for one and two, or possibly occasionally three days in advance. These predictions all relate to comparatively minute details for regions that have been charted and studied daily for many years. They merely represent the direct teachings of experience; they are generalizations based upon observations but into which physical theories have as yet entered in only a superficial manner if at all. They

WEATHER PREDICTION BY NUMERICAL PROCESS

LEWIS F. RICHARDSON, B.A., F.R.MET.SOC., F.I.NST.P.
FORMERLY SUPERINTENDENT OF ESKDALEMUIK OBSERVATORY
LECTURER ON PHYSICS AT WESTMINSTER TRAINING COLLEGE

CAMBRIDGE
AT THE UNIVERSITY PRESS
1922



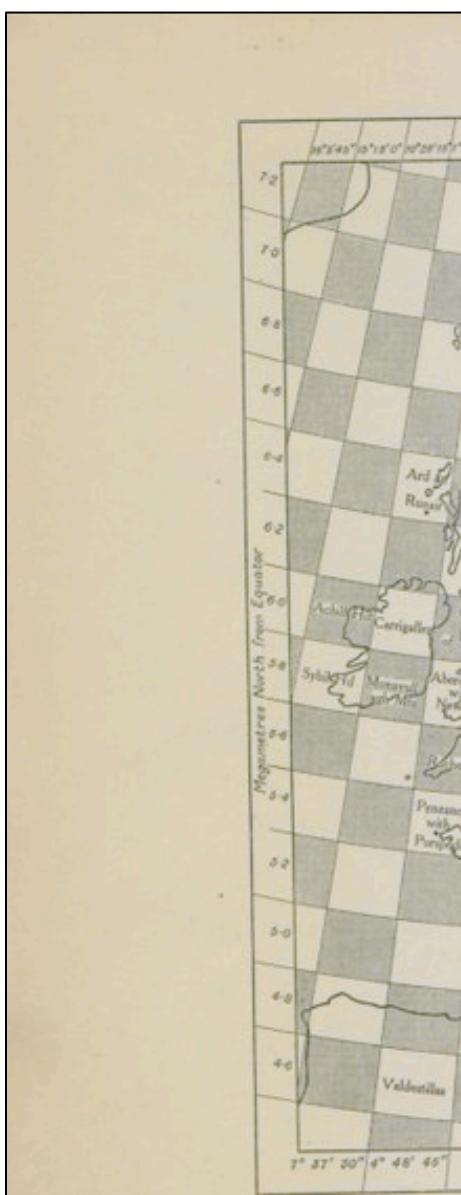
An arrangement of meteorological stations designed to fit with the chief mechanical properties of the atmosphere. Other considerations have been here disregarded. Pressure to be observed at the centre of each shaded chequer, velocity at the centre of each white chequer. The numerical coordinates refer to these centres as also do the names, although as to the latter there may be errors of 5 or 10 km. The word "with" in "St Leonards with Dieppe" etc. is intended to suggest an interpolation between observations made at the two places. See page 9, and Chapters 3 and 7. Contrast the existing arrangement shown on p. 184.

L. F. Richardson, *Weather Predictions by Numerical Process* 1922



Weather Predictions

C. Abbe, Physical Basis of Long Range Weather Forecast, 1901



An arrangement of meteorological stations designed to fit with the chief mechanical properties of the atmosphere. Other considerations have been here disregarded. Pressure to be observed at the centre of each shaded chequer, velocity at the centre of each white chequer. The numerical coordinates refer to these centres as also do the names, although as to the latter there may be errors of 5 or 10 km. The word "with" in "St Leonards with Dieppe" etc. is intended to suggest an interpolation between observations made at the two places. See page 9, and Chapters 3 and 7. Contrast the existing arrangement shown on p. 184.

DECEMBER, 1901.

MONTHLY WEATHER REVIEW.

551

24th, both preceded by heavy swell and followed by low dew-point. The north wind evidently precipitated the terrific downpour on north Hilo. There was lightning reported from Hawaii for the 8th, 9th, 13th, 24th, and 25th, and Maui on the 8th. Snow fell on Mauna Kea and Mauna Loa on the 8th and 24th; on Haleakala on the 8th or 9th. Earthquake

THE PHYSICAL BASIS OF LONG-RANGE WEATHER FORECASTS.¹

By PROF. CLEVELAND ABBE.

The expression "long range" must not be misunderstood. It refers only to the length of time intervening between the

"Despite the advances made by Richardson, it took him, working alone, several months to produce a wildly inaccurate six-hour forecast for an area near Munich, Germany. In fact, some of the changes predicted in Richardson's forecast could never occur under any known terrestrial conditions. Adding to the failure of this effort, a six-hour forecast is not particularly useful if it takes weeks to produce."

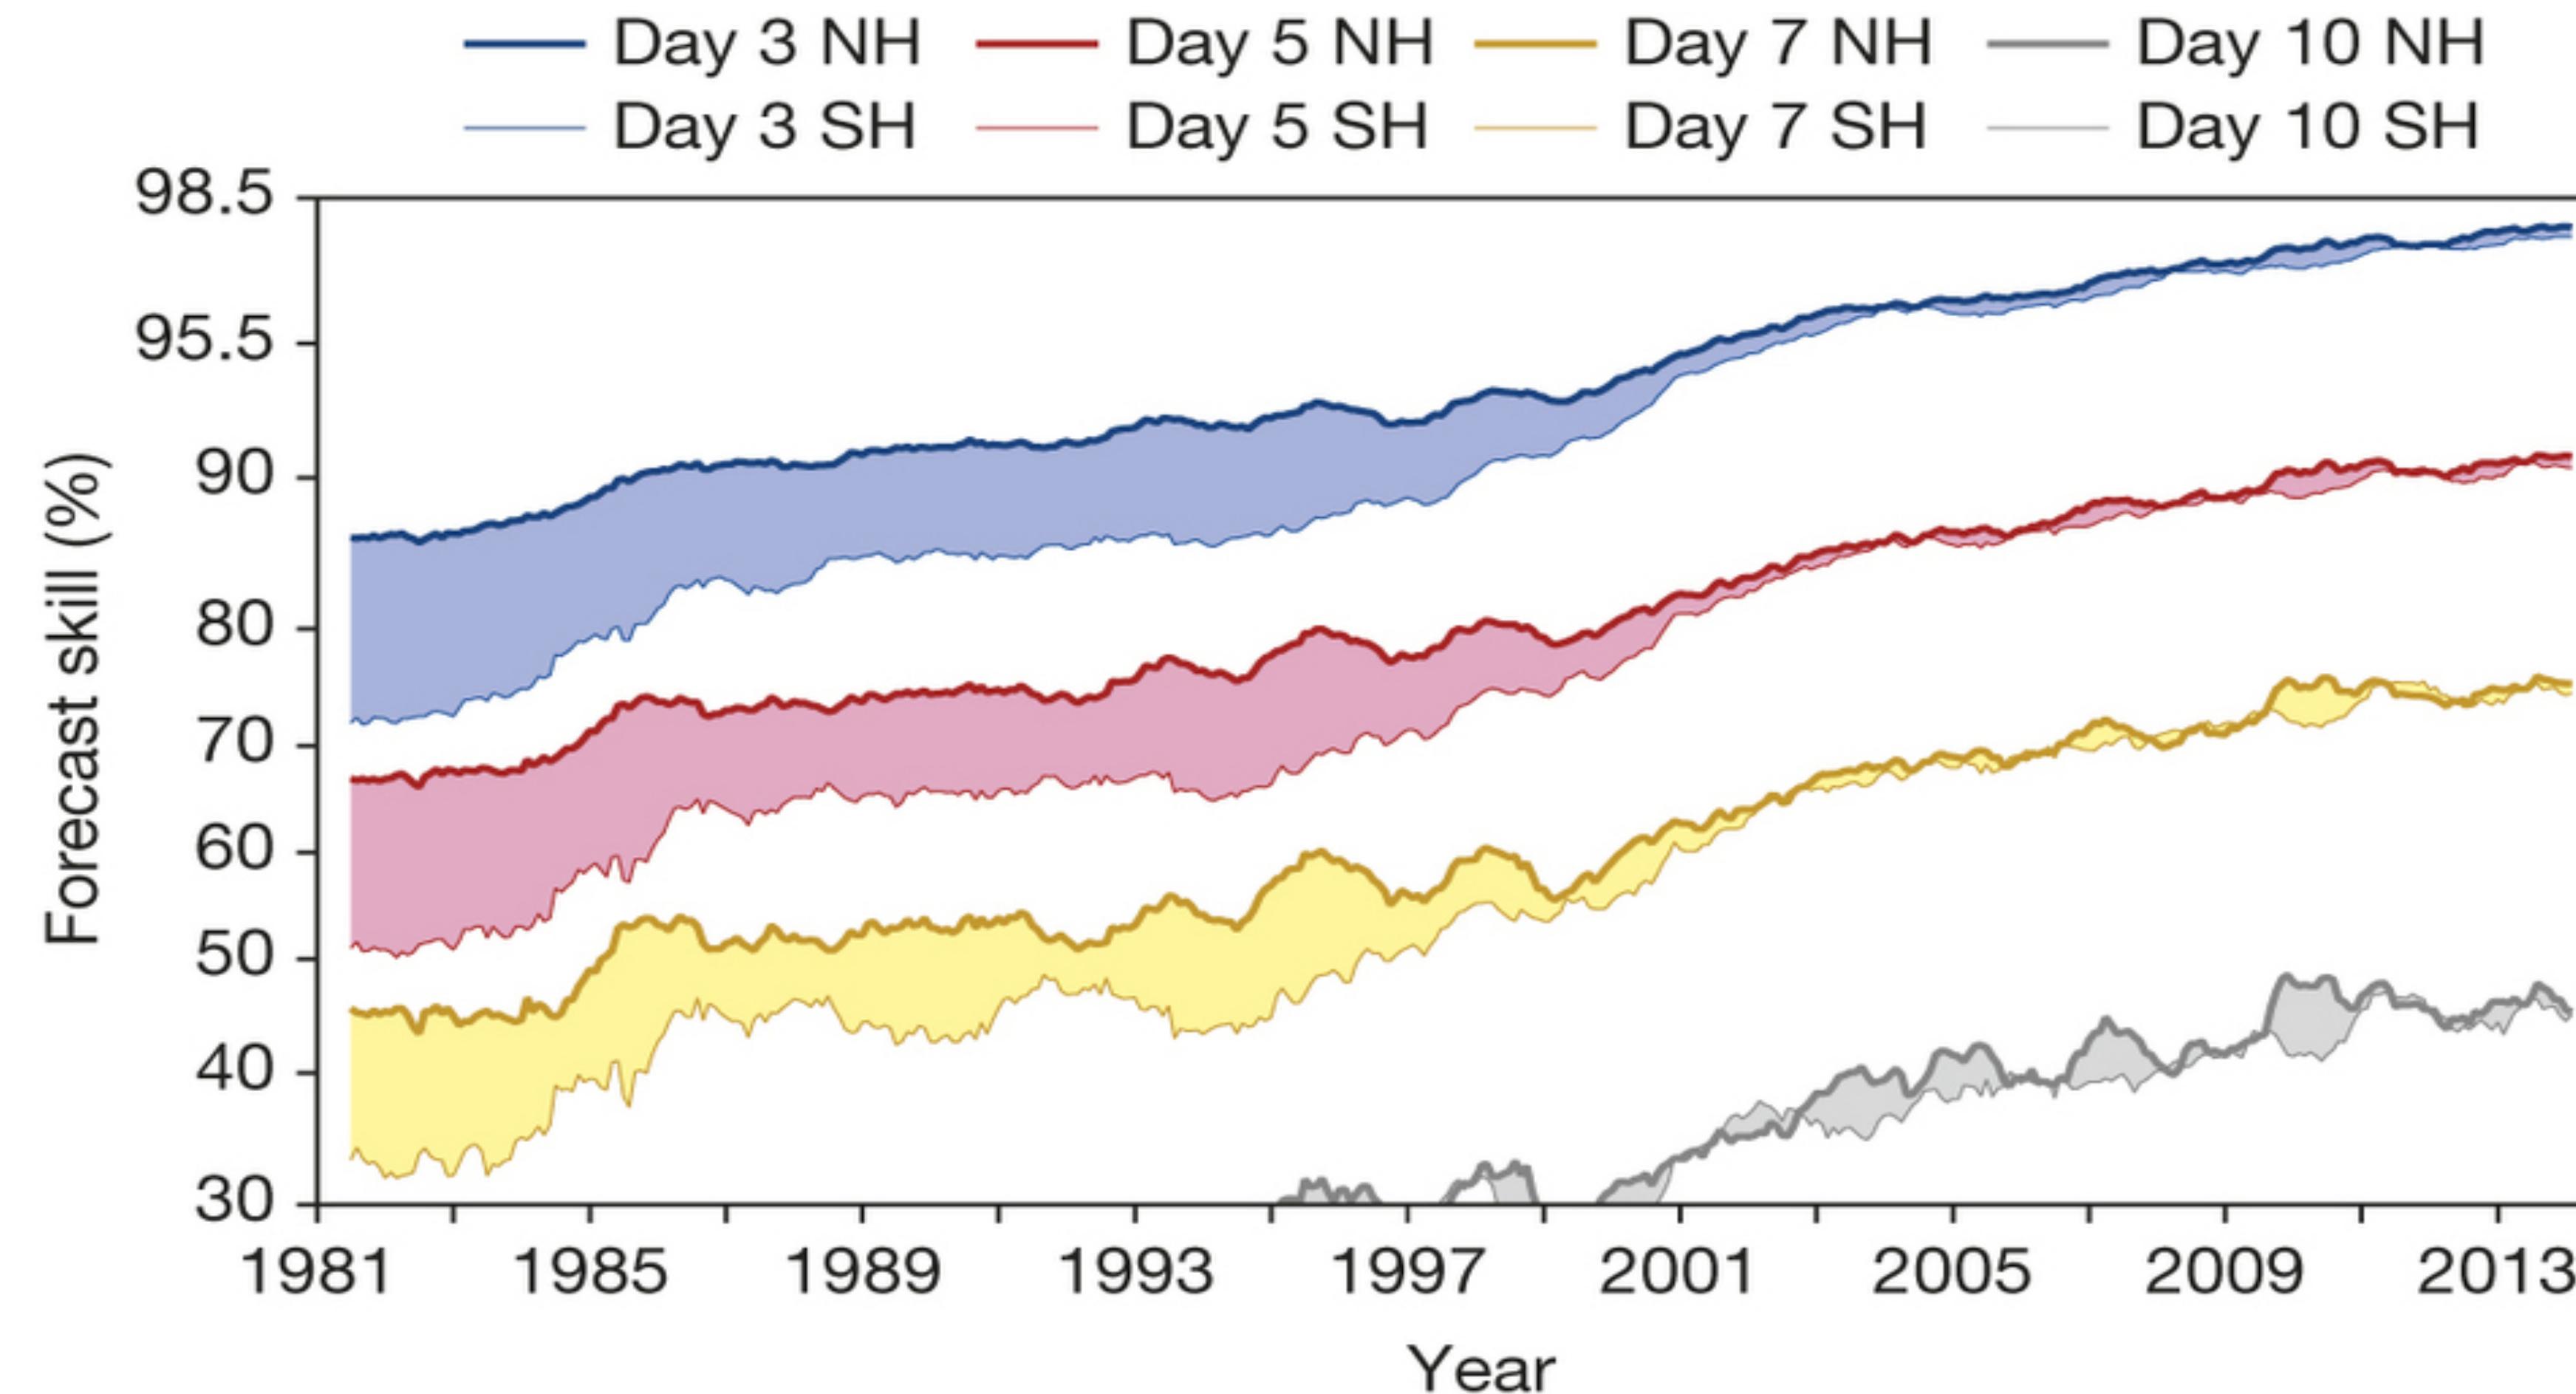
"Weather Forecast through the Ages", 2007

CAMBRIDGE
AT THE UNIVERSITY PRESS
1922

L. F. Richardson, Weather Predictions by Numerical Process 1922



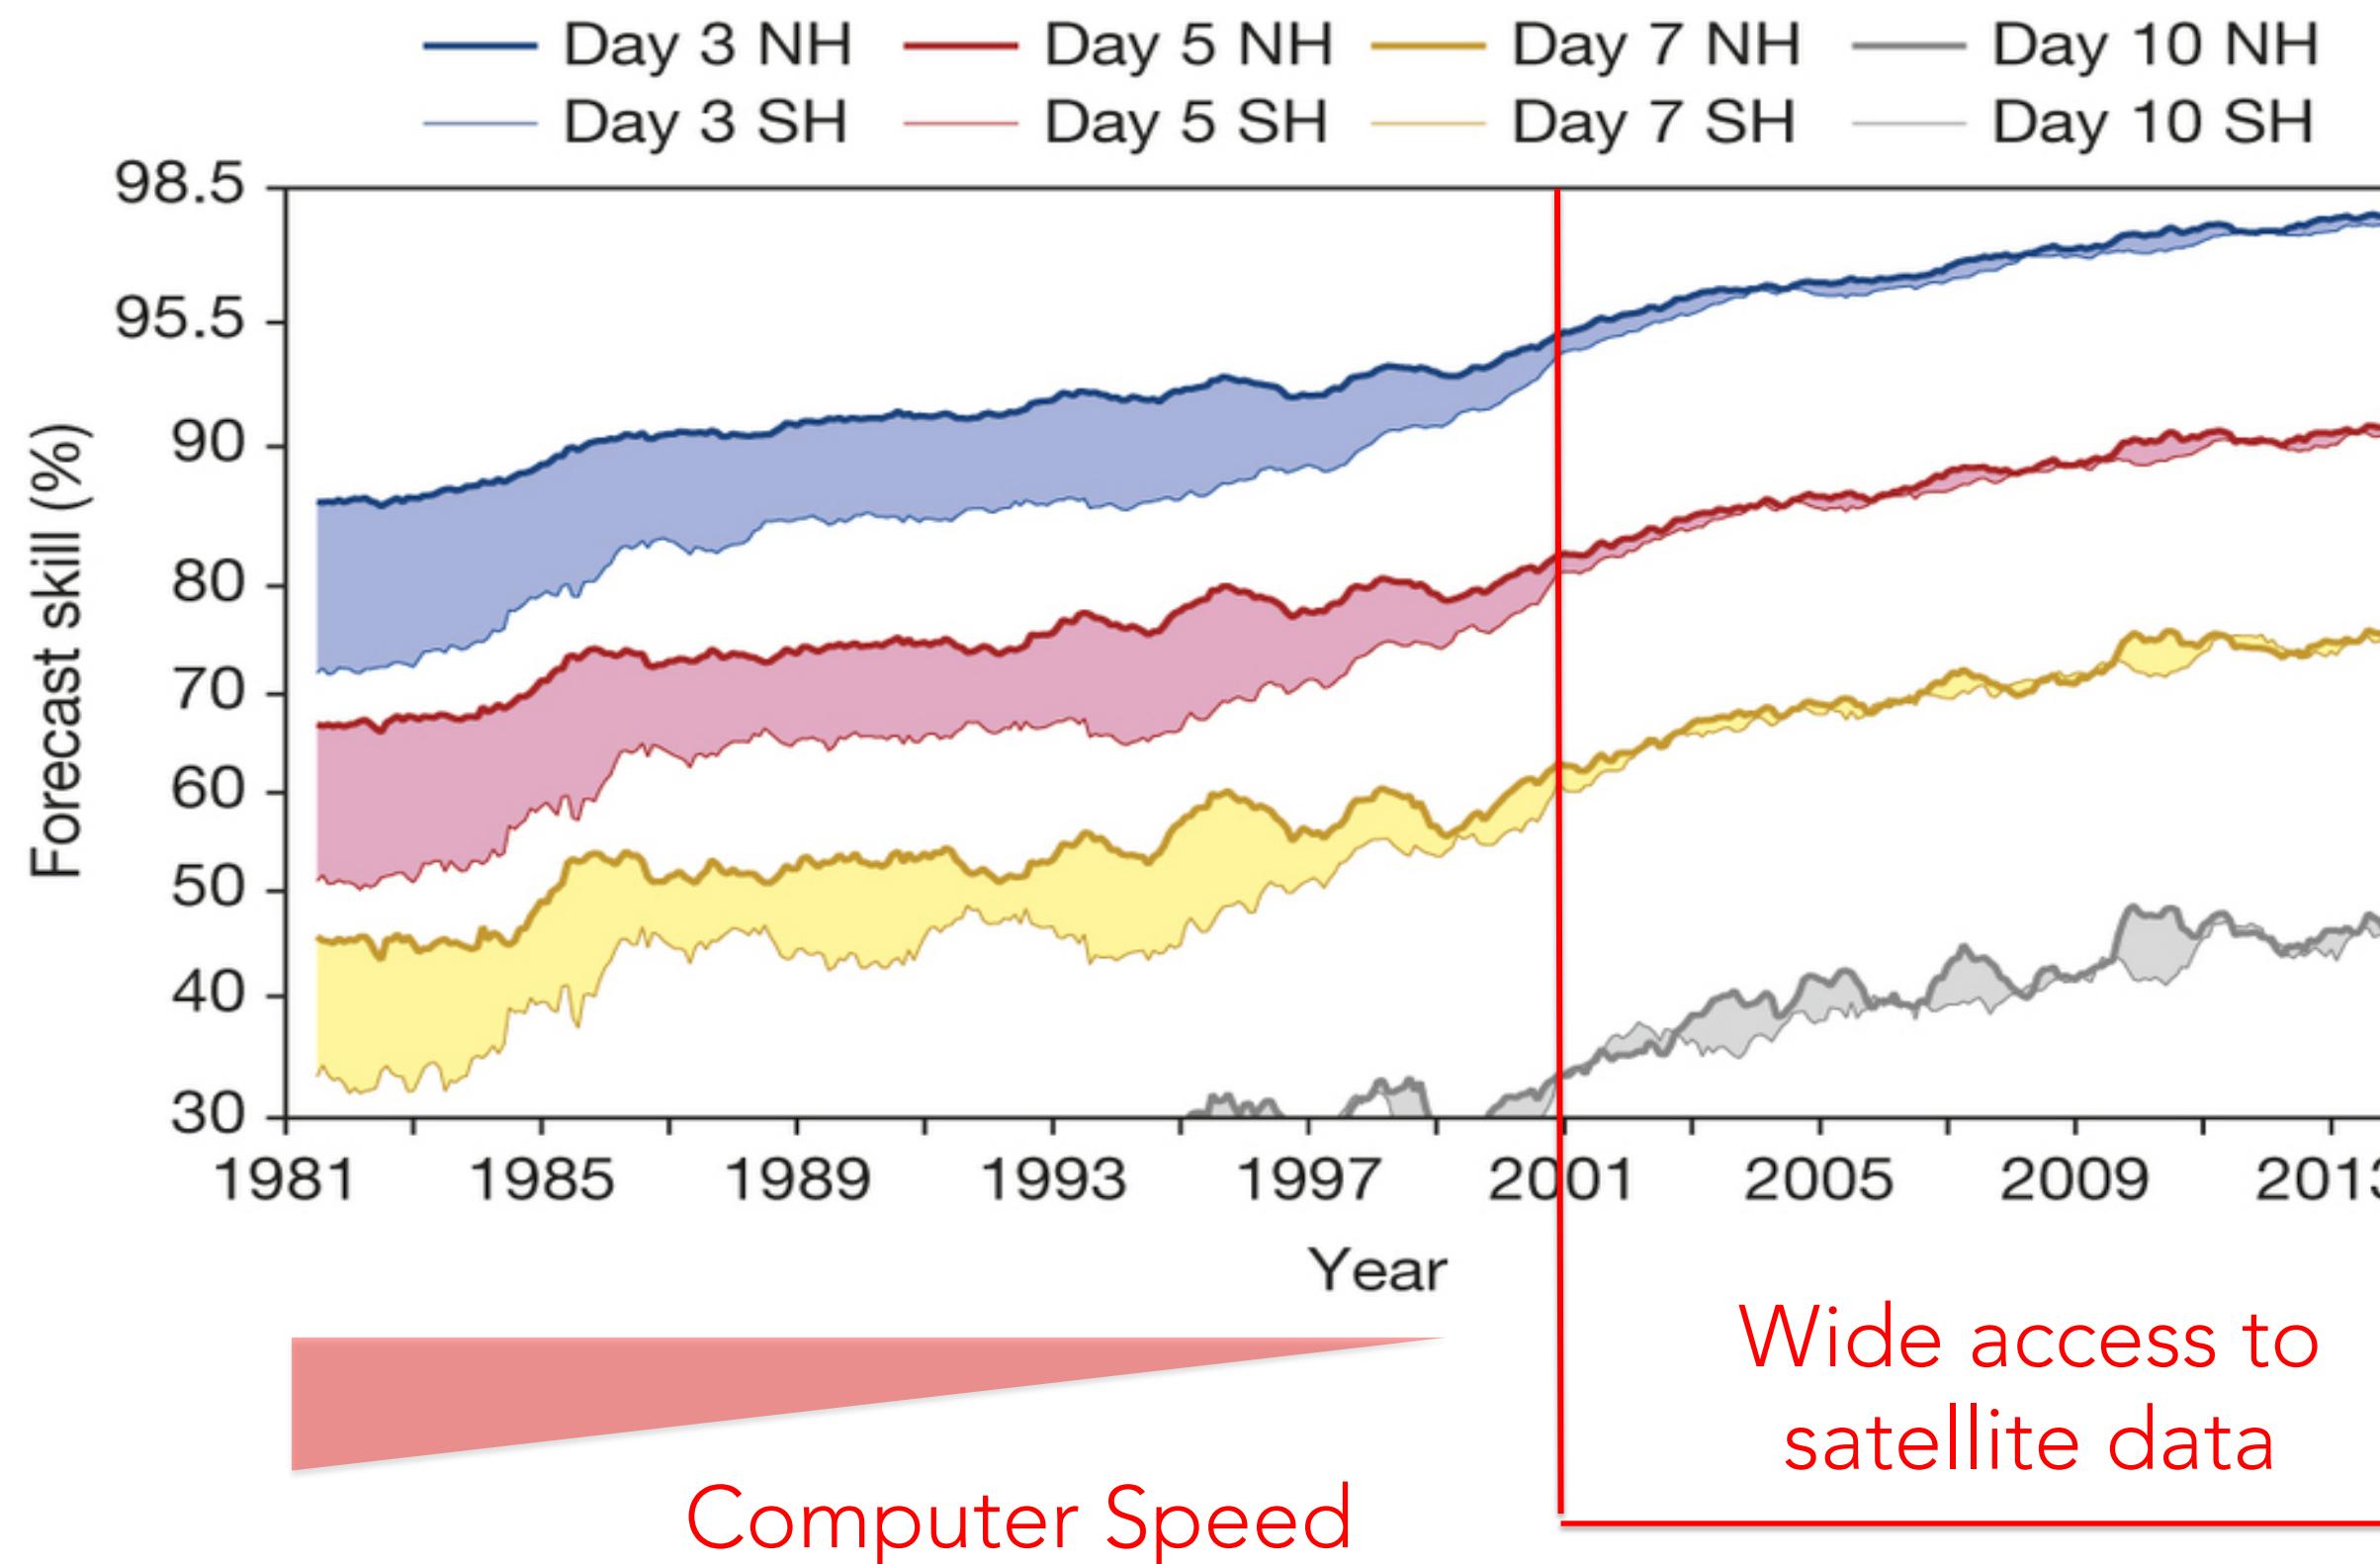
Are weather predictions predictive?



*The quiet revolution of numerical weather prediction
(Bauer et al., Nature 2015)*



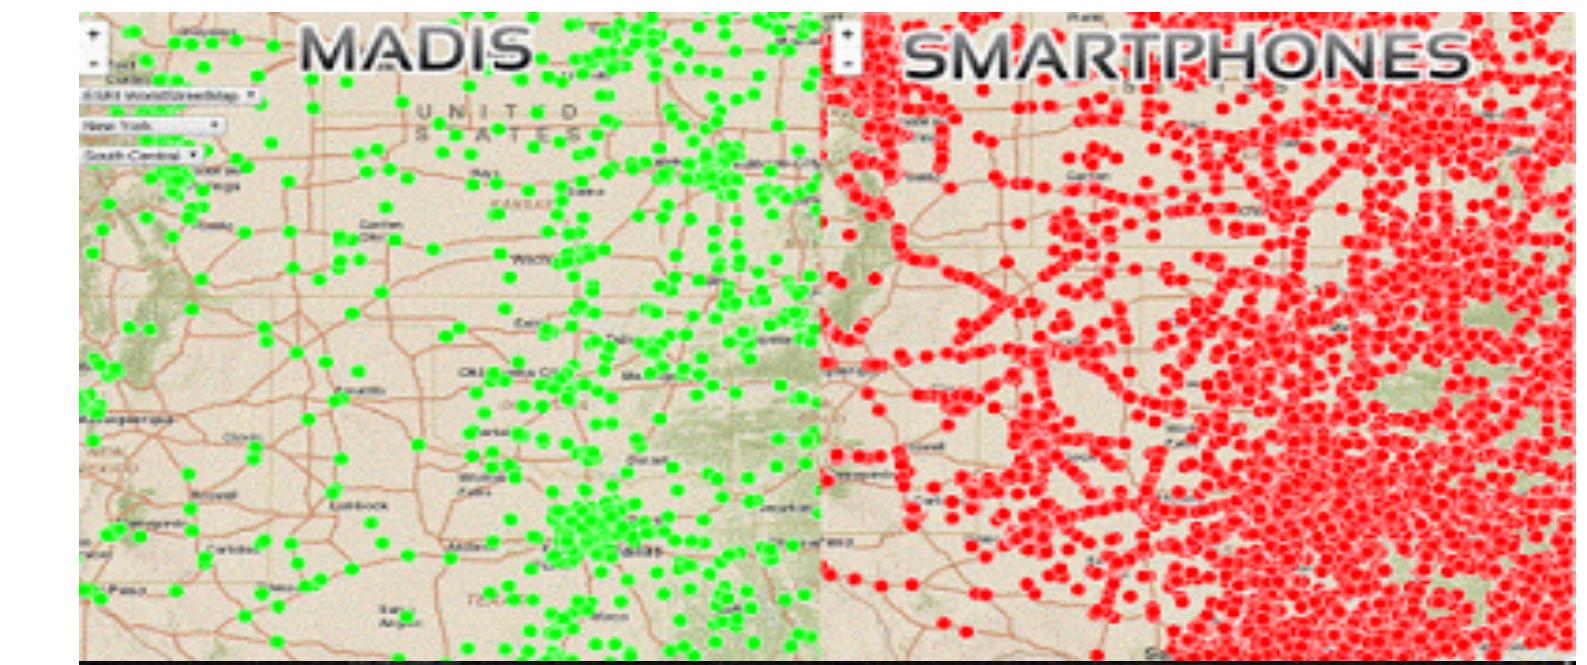
Are weather predictions predictive?



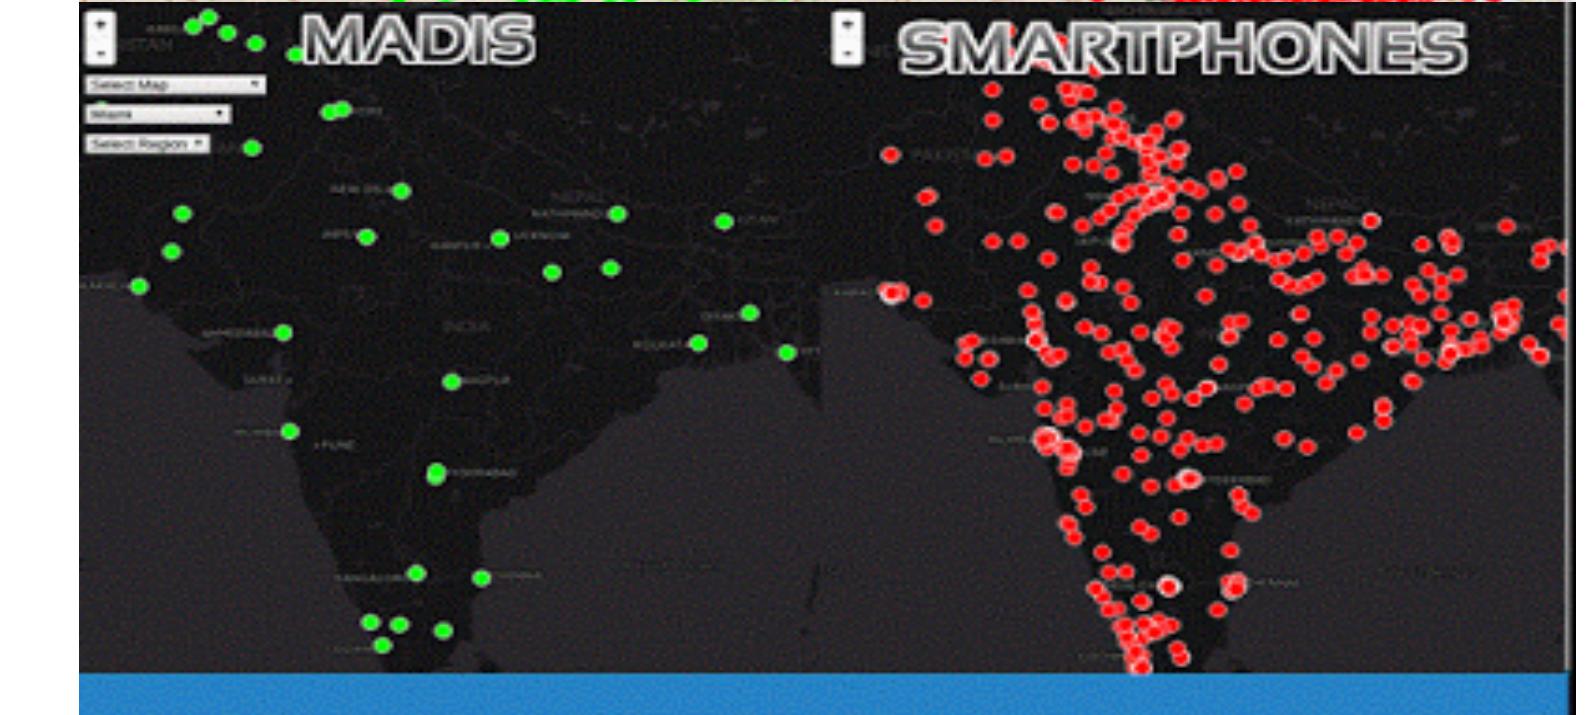
*The quiet revolution of numerical weather prediction
(Bauer et al., Nature 2015)*

Next revolution

Smartphone Crowdsourcing?
(Big Data)



US



India

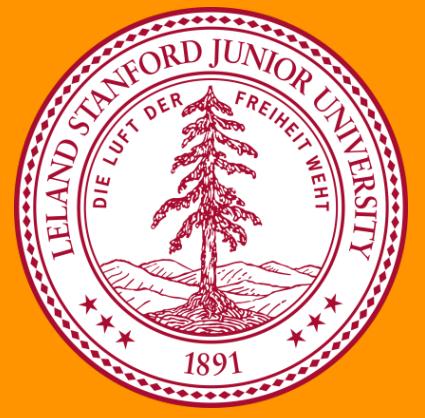
MADIS: Meteorological Assimilation
Data Ingest System



**(Engineering)
Simulations**

+

Data



Simulations + Data

Validation

The “C” word

Statistical Inference

Machine Learning

Physics + Data



Simulations + Data



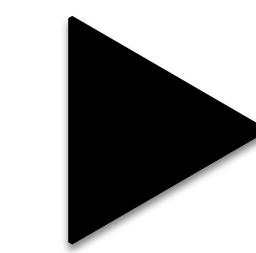
Validation

The “C” word

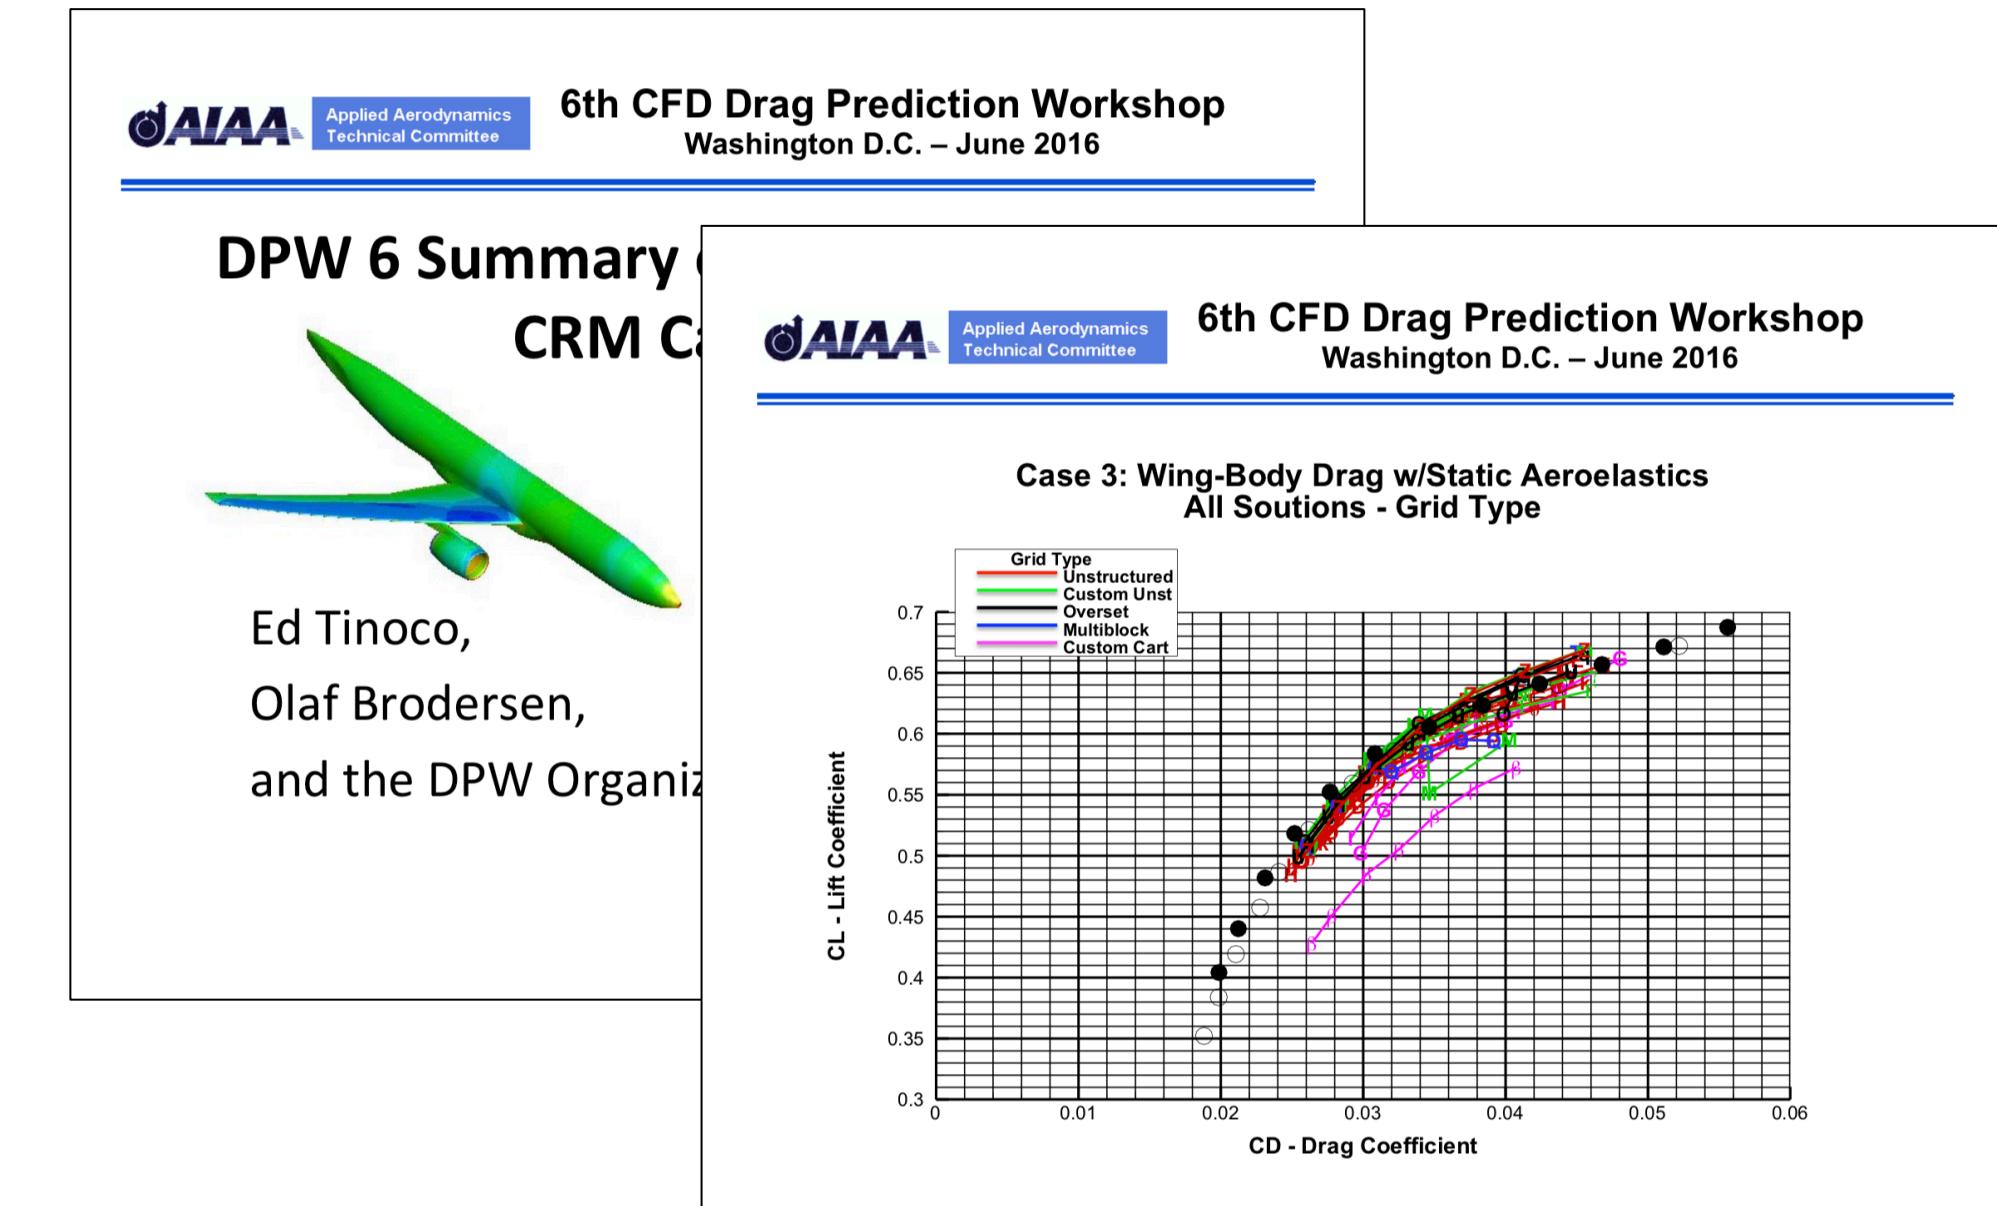
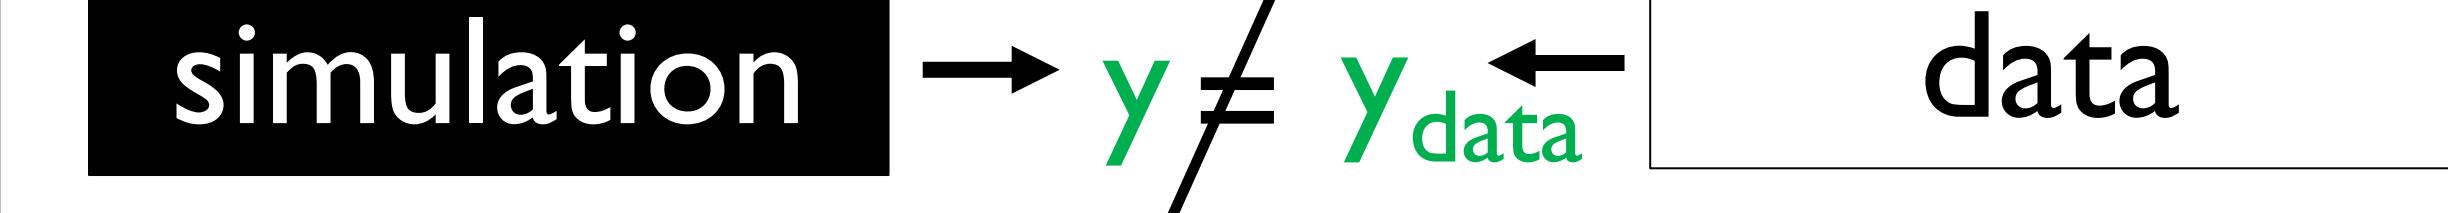
Statistical Inference

Machine Learning

Physics + Data



Blind Tests: Accuracy Evaluation





Simulations + Data

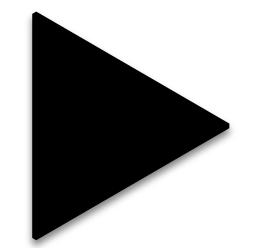
Validation

The “C” word

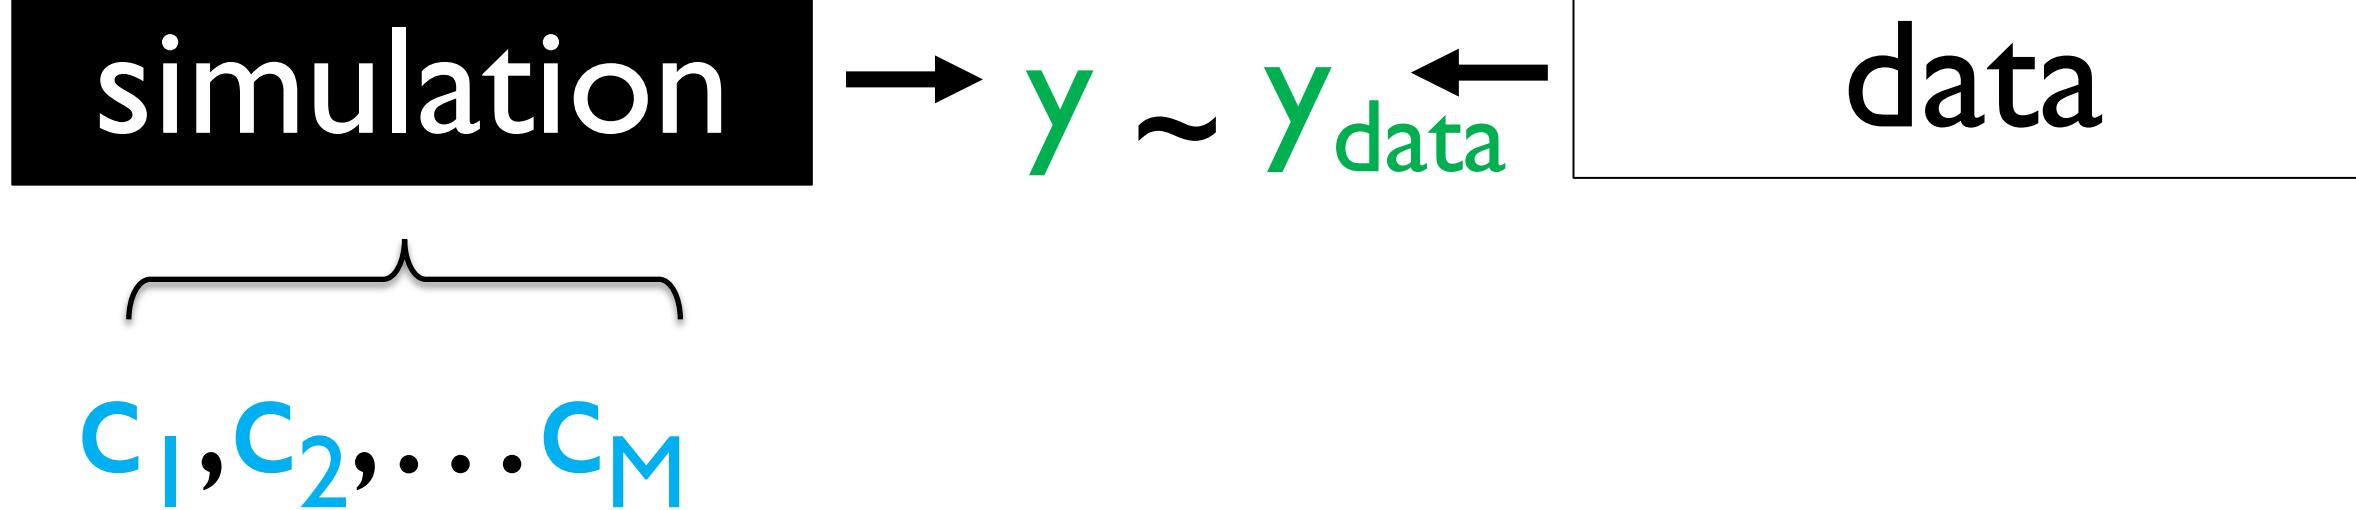
Statistical Inference

Machine Learning

Physics + Data



Calibration: adjust (tune) model parameters to match data



	Mixing Layer	Plane Jet	Round Jet	Radial Jet	Wake
Experiment ^{7, 78-81}	0.103-0.120	0.100-0.110	0.086-0.096	0.096-0.110	0.320-0.400
$k - \omega$ Wilcox ^{82*}	0.096	0.108	0.094	0.099	0.326
$k - \omega$ Wilcox ^{83*}	0.105	0.101	0.088	0.099	0.339
$k - \omega$ Wilcox ^{84*}	0.141	0.135	0.369	0.317	0.496
SST Menter ^{85†}	0.100	0.112	0.127	-	0.257
$k - \epsilon$ Launder & Sharma ^{77*} with Pope ^{86*}	0.098 0.098	0.109 0.109	0.120 0.086	0.094 0.040	0.256 0.256
$k - \zeta$ Robinson, et al. ^{87*}	0.112	0.115	0.091	0.097	0.313
$k - \tau$ Speziale, et al. ^{88*}	0.082	0.089	0.102	0.073	0.221
SA Spalart & Allmaras ^{76*}	0.109	0.157	0.248	0.166	0.341

Source: Yoder et al. NASA/TM 2013-218072



Simulations + Data



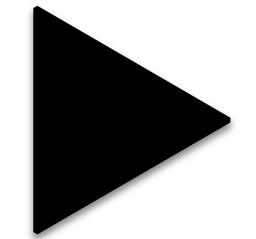
Validation

The “C” word

Statistical Inference

Machine Learning

Physics + Data



Principled Calibration: solve an optimization problem to find parameters

$$\min_{c_1, c_2, \dots, c_M} \| y - y_{\text{data}} \|^2$$



c_1, c_2, \dots, c_M

....assuming that data is noisy leads to
Bayesian inversion



Simulations + Data



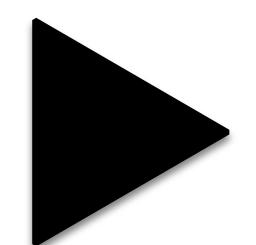
Validation

The “C” word

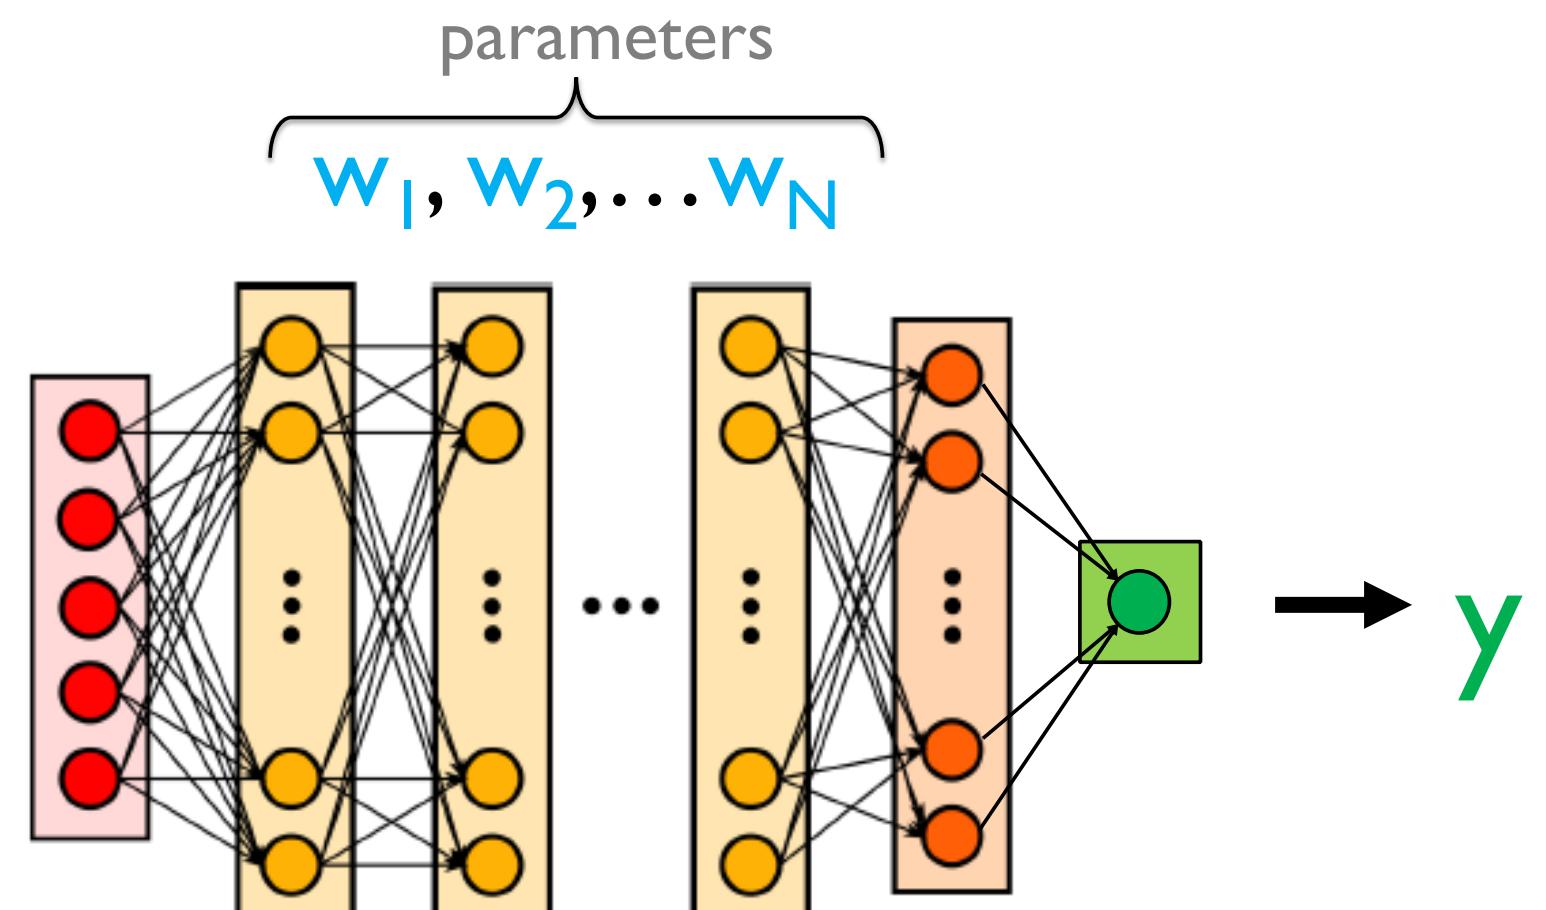
Statistical Inference

Machine Learning

Physics + Data



Deep Neural Network: a general input-output relationship



...solve an optimization problem to find the weights

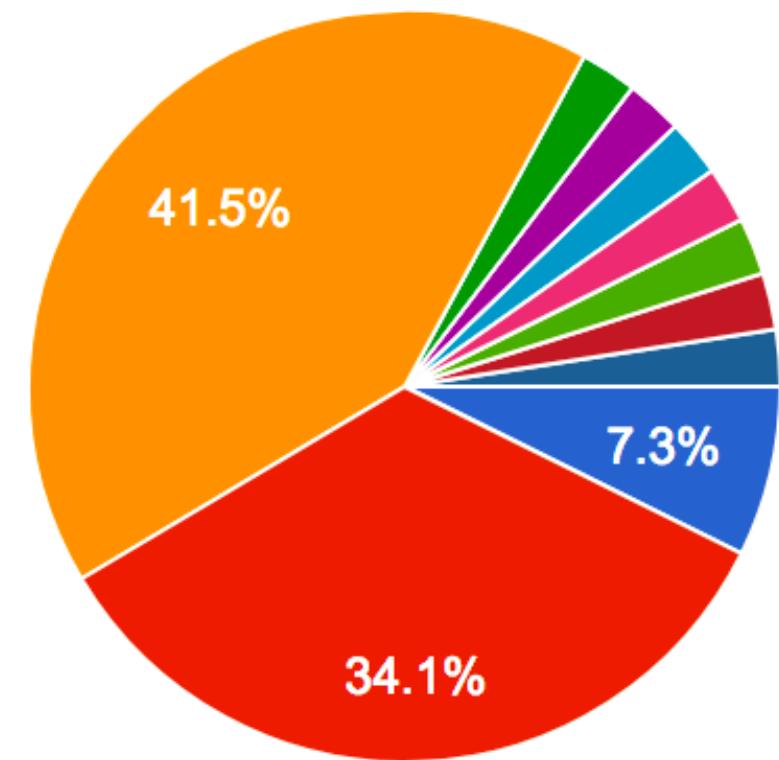
$$\min_{w_1, w_2, \dots, w_N} \| y - y_{\text{data}} \|^2$$



A Survey of ~80 Fluid Mechanics Researchers

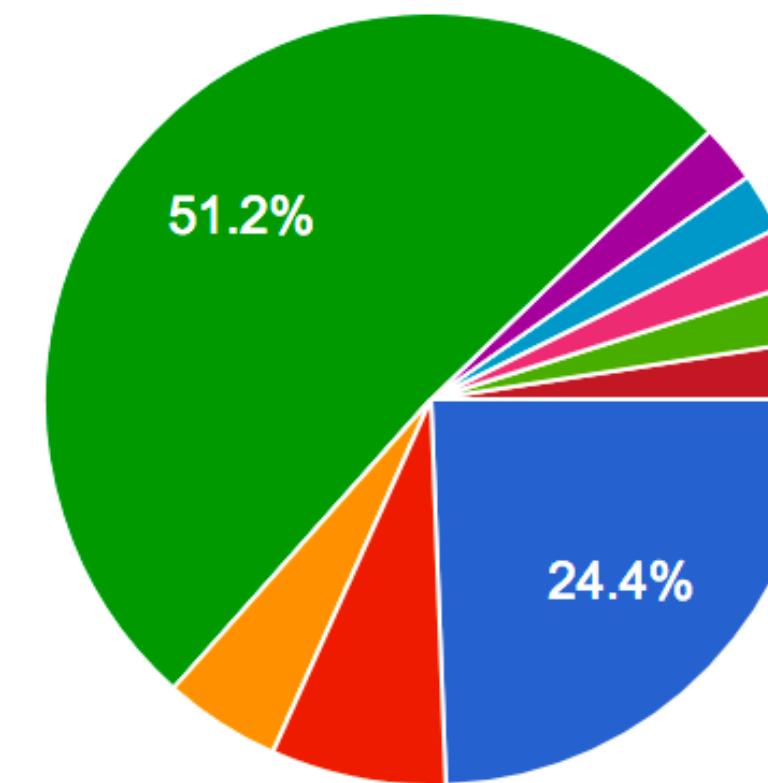
2018 Stanford Center for Turbulence Research

What is your main criticism of data-driven predictions?



- 82.9% ▶ [grey bar] ◀
- Our problems are much more comp...
 - Low confidence in the resulting bla...
 - We need to use physics-based equ...
 - It's ONLY buzz and hype
 - The higher purpose is sometimes m...
 - All of the above, but maybe similar t...
 - It is just glorified curvefitting that no...
 - it has the risk of giving more import...
- ▲ 1/2 ▼

What do you think is the most important ingredient for data-driven predictions?



- 51.2% ▶ [grey bar] ◀
- Using the "right" data (appropriate, relevant, etc.)
 - Having enough data
 - Employing the best learning strateg...
 - Satisfying the correct physical princi...
 - Thinking-driven, data-aided is better
 - Tough question, this probably "flow"...
 - depends on the field it is applied to
 - Using for its strengths, avoid cases...
 - all of the above



Simulations + Data

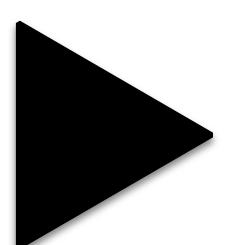
Validation

The “C” word

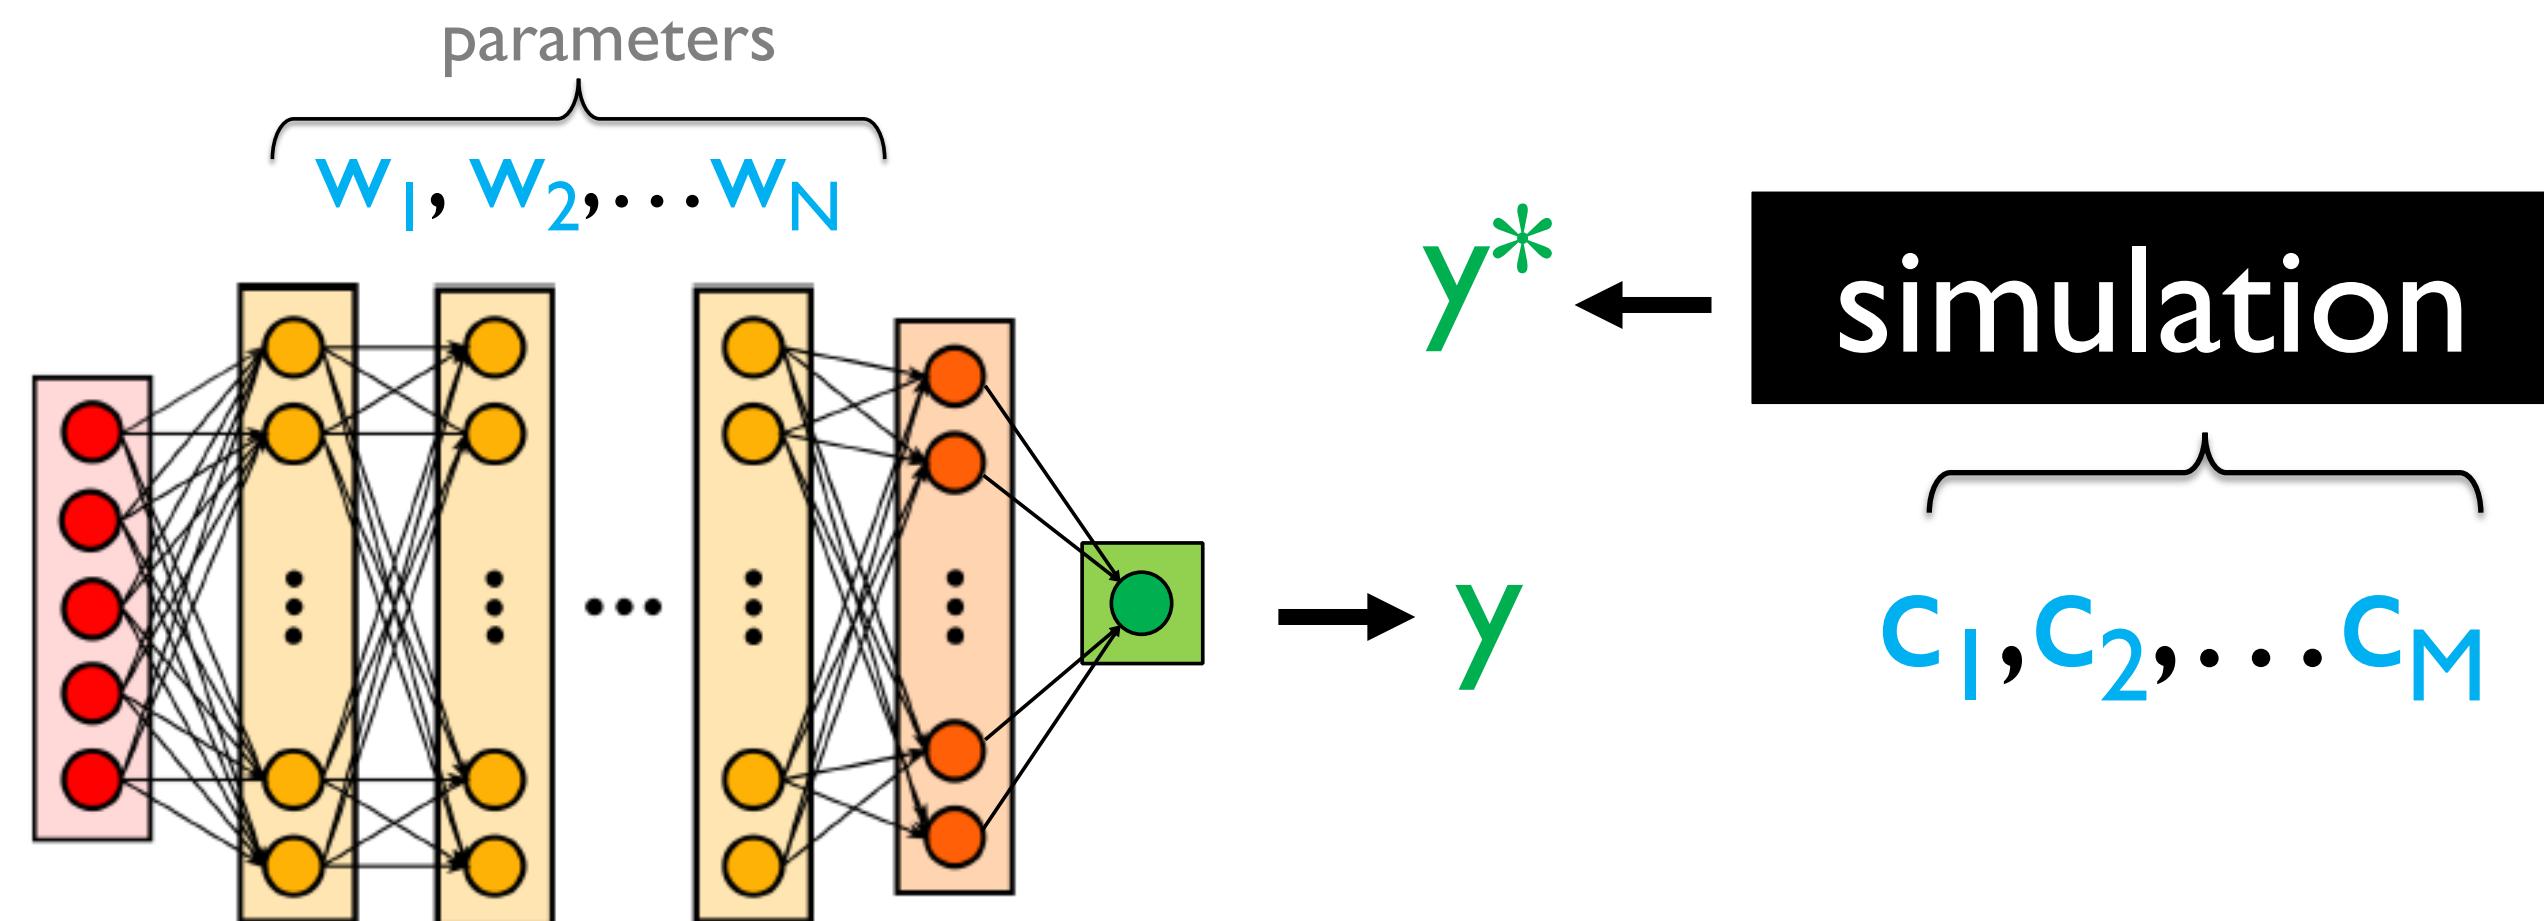
Statistical Inference

Machine Learning

Physics + Data



Deep Neural Network & Simulation



...solve an optimization problem to find the weights but constrain the acceptable outputs

$$\min_{w_1, w_2, \dots, w_N} \| y - y_{\text{data}} \|^2 \quad \& \quad \min_{c_1, c_2, \dots, c_M} \| y - y^* \|^2$$



Simulations + Data



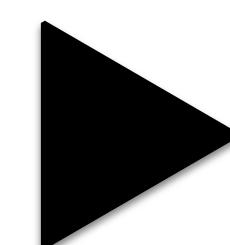
Validation

The “C” word

Statistical Inference

Machine Learning

Physics + Data



Deep Neural Network & Simulation

<https://github.com/lululxvi/deepxde>

arXiv.org > cs > arXiv:1907.04502

Search...

Help | Advanced

Computer Science > Machine Learning

DeepXDE: A deep learning library for solving differential equations

Lu Lu, Xuhui Meng, Zhiping Mao, George E. Karniadakis

(Submitted on 10 Jul 2019)

Deep learning has achieved remarkable success in diverse applications; however, its use in solving partial differential equations (PDEs) has emerged only recently. Here, we present an overview of physics-informed neural networks (PINNs), which embed a PDE into the loss of the neural network using automatic differentiation. The PINN algorithm is simple, and it can be applied to different types of PDEs, including integro-differential equations, fractional PDEs, and stochastic PDEs. Moreover, PINNs solve inverse problems as easily as forward problems. We propose a new residual-based adaptive refinement (RAR) method to improve the training efficiency of PINNs. For pedagogical reasons, we compare the PINN algorithm to a standard finite element method. We also present a Python library for PINNs, DeepXDE, which is designed to serve both as an education tool to be used in the classroom as well as a research tool for solving problems in computational science and engineering. DeepXDE supports complex-geometry domains based on the technique of constructive solid geometry, and enables the user code to be compact, resembling closely the mathematical formulation. We introduce the usage of DeepXDE and its customizability, and we also demonstrate the capability of PINNs and the user-friendliness of DeepXDE for five different examples. More broadly, DeepXDE contributes to the more rapid development of the emerging Scientific Machine Learning field.

<https://github.com/kailaix/ADCME.jl>

arXiv.org > math > arXiv:1905.12530

Search...

Help | Advanced

Mathematics > Numerical Analysis

Predictive Modeling with Learned Constitutive Relations from Indirect Observations

Daniel Z. Huang, Kailai Xu, Charbel Farhat, Eric Darve

(Submitted on 29 May 2019 (v1), last revised 15 Oct 2019 (this version, v2))

We present a new approach for predictive modeling and its uncertainty quantification for mechanical systems, where coarse-grained models such as constitutive relations are derived directly from observation data. We explore the use of neural networks to represent the unknown functions (e.g., constitutive relations). Its counterparts, like piecewise linear functions and radial basis functions, are compared, and the strength of neural networks is explored. The training and predictive processes in this framework seamlessly combine the finite element method, automatic differentiation, and neural networks (or its counterparts). Under mild assumptions, we establish convergence guarantees. This framework also allows uncertainty quantification analysis in the form of intervals of confidence. Numerical examples on a multiscale fiber-reinforced plate problem and a nonlinear rubbery membrane problem from solid mechanics demonstrate the effectiveness of the framework.

Comments: 32 pages, 20 figures

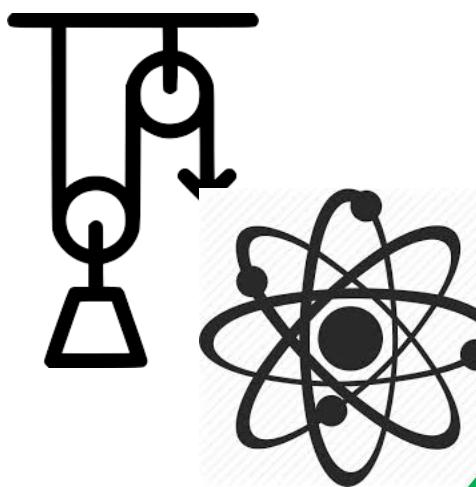
Subjects: Numerical Analysis (math.NA); Computational Physics (physics.comp-ph)



Simulations + Data

Physical Principles

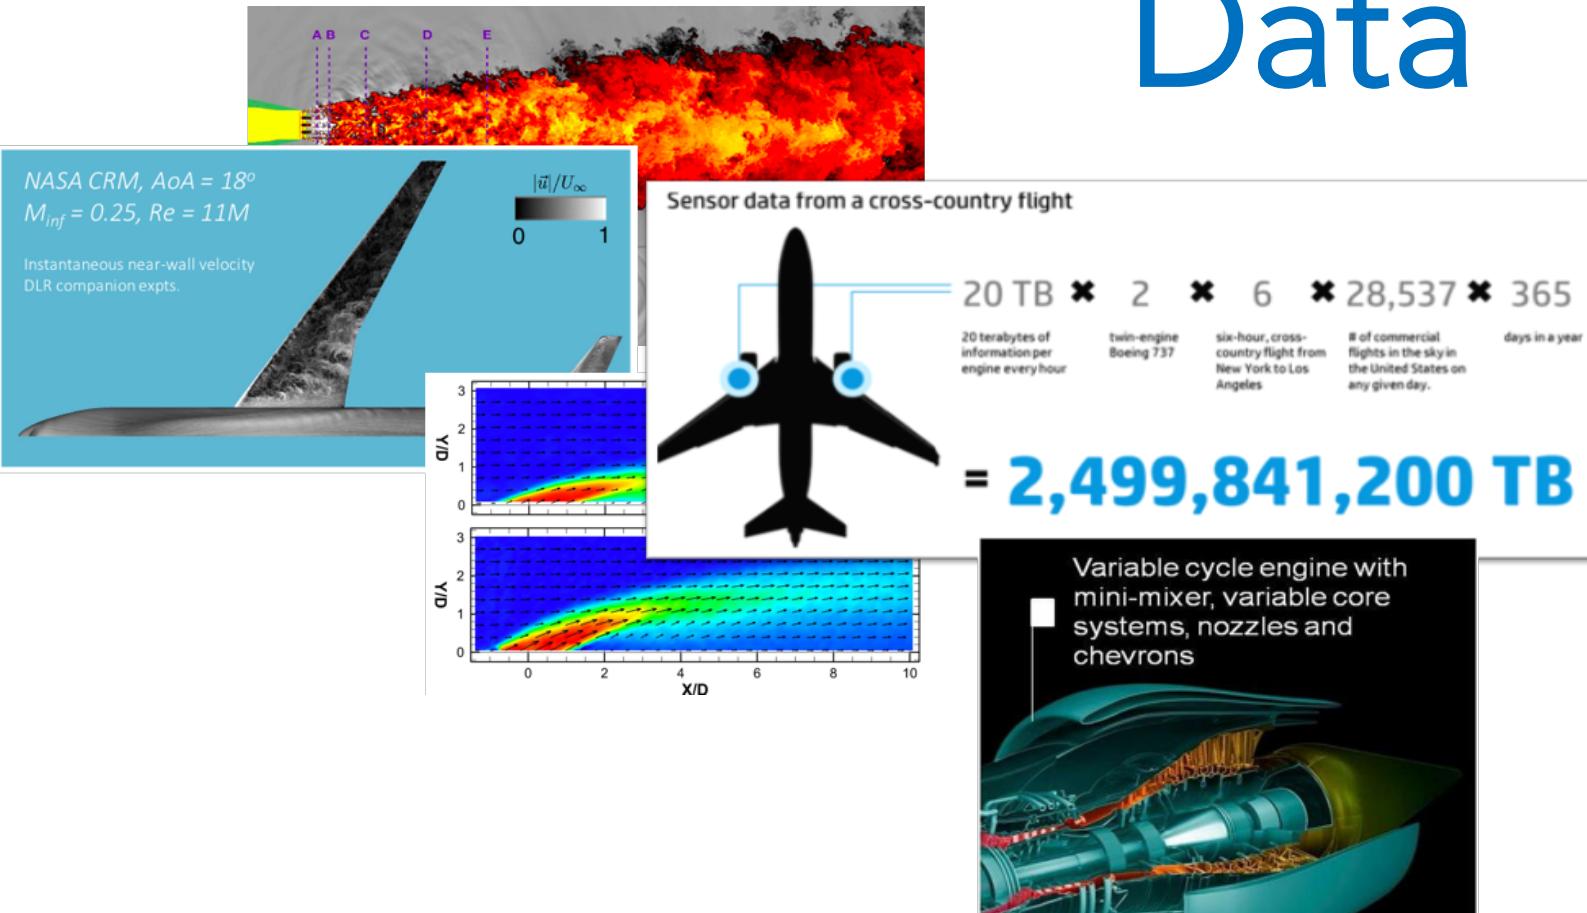
$$\begin{aligned} C(u) &= \sum_{k=1}^{\infty} P_k u^k & y = \phi(x) &= \frac{1}{\sqrt{\pi}} \int_0^{\infty} e^{-\frac{t^2}{2}} dt & S(a, t) &= \frac{2}{\pi t} \int_0^t \frac{\sin at}{t} dt & P(\eta_0 < \\ S_n &= A_n U T A_n & W_k &= \left(\frac{u}{k} \right) P_k^k (1-p)^{k-1} & P(\eta_0 y | f = x) &= \sup_{y' \leq y, y' \neq x} P(\eta_0 y' | f = x) \\ |A_n| &= \frac{u^n}{n!} & \int f(x) \log_2 \frac{1}{f(x)} dx & \leq \varepsilon & g^{-1} \cdot g &= e \\ \sum_{n=0}^{\infty} & \sum_{k=0}^{n+1} e^{-\frac{(k-n)^2}{2}} = H(n) & \prod_{k \leq b} \bigcup_{i=1}^{n-1} M_i; \bigcap_{i=0}^{\infty} X_n & & g &= \sqrt{\frac{2\mu}{\lambda}} \left(\frac{\eta_{2n} + \eta_{2n-2} \cdot \eta_{2n}}{\eta_{2n}} \right) \\ \int dG_n(x) &\geq \frac{1}{2} & n > \infty & & f_n(t) &= \frac{2^n u^{n-t} e^{-2t}}{(n-t)!} \\ f_{n-1}(t) &= \int f_n(u) f_n(t-u) du = \frac{2^{n+1} u^n e^{-2t}}{n!} & \lim_{t \rightarrow \infty} (g^t) &= 0 & \lim_{n \rightarrow \infty} \frac{S(n)}{n} &= P_k \\ \log \varphi(t) &= i \gamma t - c |t|^2 \left[1 + \beta \frac{t}{|t|} \cos \theta \right] B(n) & \sum_{k=1}^n \psi^*(b_k u) & C_{i,j} &= \sum_{j=1}^n a_{ij} b_j u \\ \int_{-\infty}^{\infty} e^{-\frac{u^2}{2}} du &= F(x) \left(\frac{1}{\sqrt{2\pi}} \right)^{-1} & \psi(n) &= \sum_{n=1}^{\infty} \frac{e^{i \omega n}}{n} \left(\frac{\sin(\omega n) - \log \frac{1}{q}}{1 - q} \right) C \\ \prod_{m=1}^n &= \prod_{l=1}^r l^{l m - r} & |\psi_s(n)| &= \left| \int_0^{\infty} e^{itx} dF(x) \right| \leq \int_0^{\infty} e^{-tx} dF(x) = \varphi_s(tx) & g^{-1} N_g &= \{ g^{-1} n_g | n \end{aligned}$$



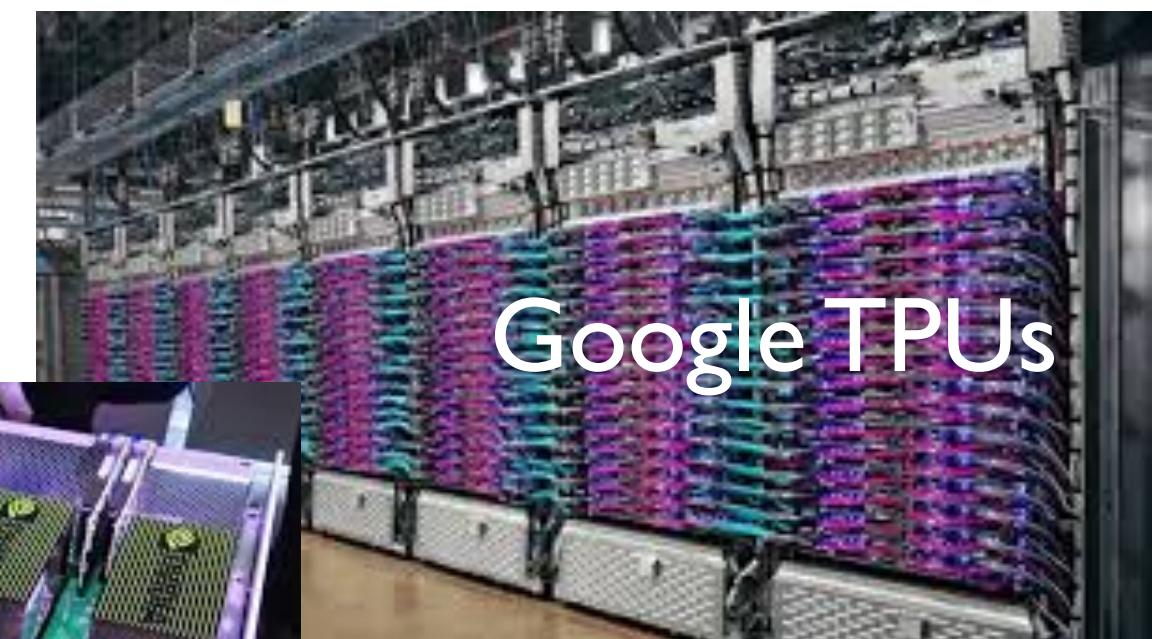
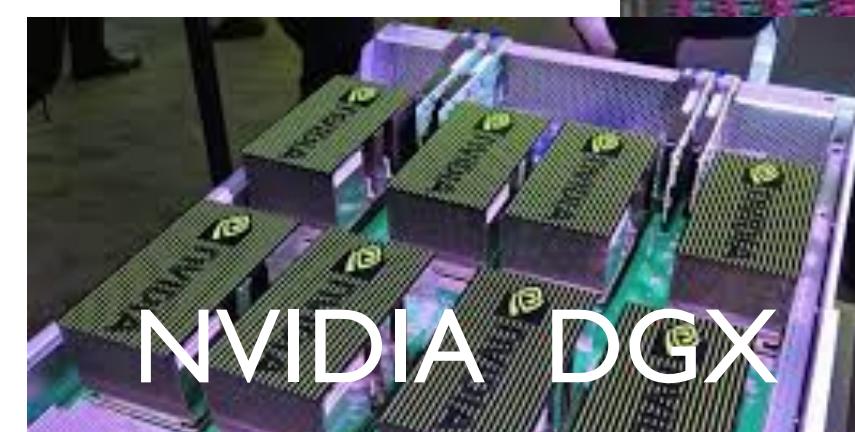
Software for Data



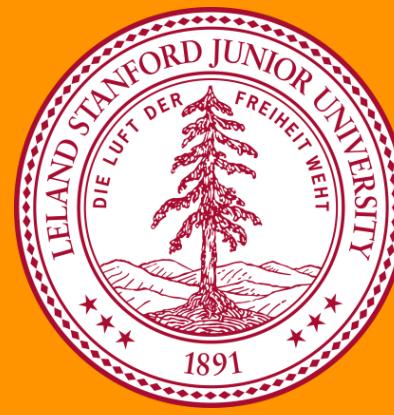
Data



Computing for Data



Google TPUs



Outline

1. Blind Experiment >> ML-driven Predictions
2. Physics First >> ML-driven Modeling
3. Reset >> ML-driven Numerical Methods



Outline

1. **Blind Experiment >> ML-driven Predictions**
2. **Physics First >> ML-driven Modeling**
3. **Reset >> ML-driven Numerical Methods**



ML@Google: Synthetic Faces (deepfakes)



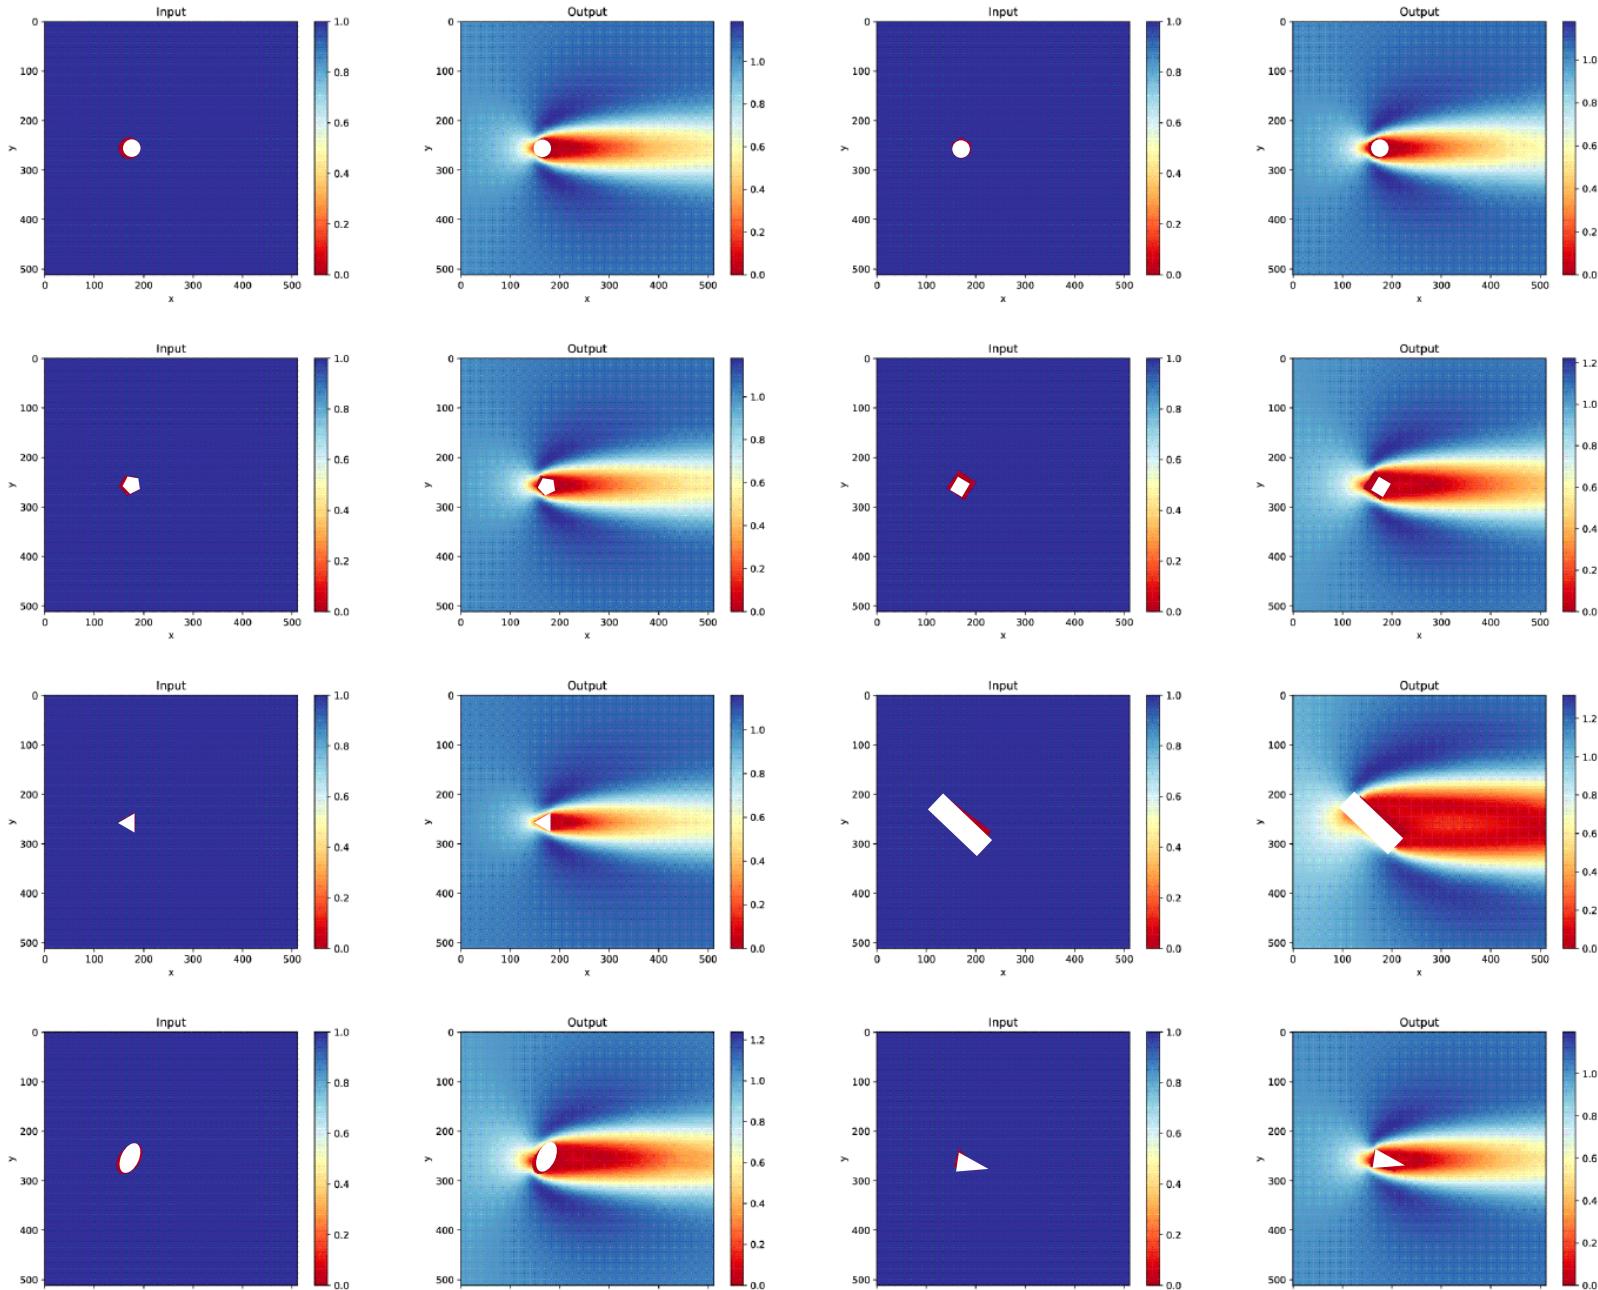
Increasingly realistic synthetic faces generated by ML@Google Goodfellow et al (2014), Radford et al. (2015), Liu and Tuzel (2016), Karras et al. (2017), Karras et al. (2018)



Synthetic Flow Simulations: flowfakes

Data

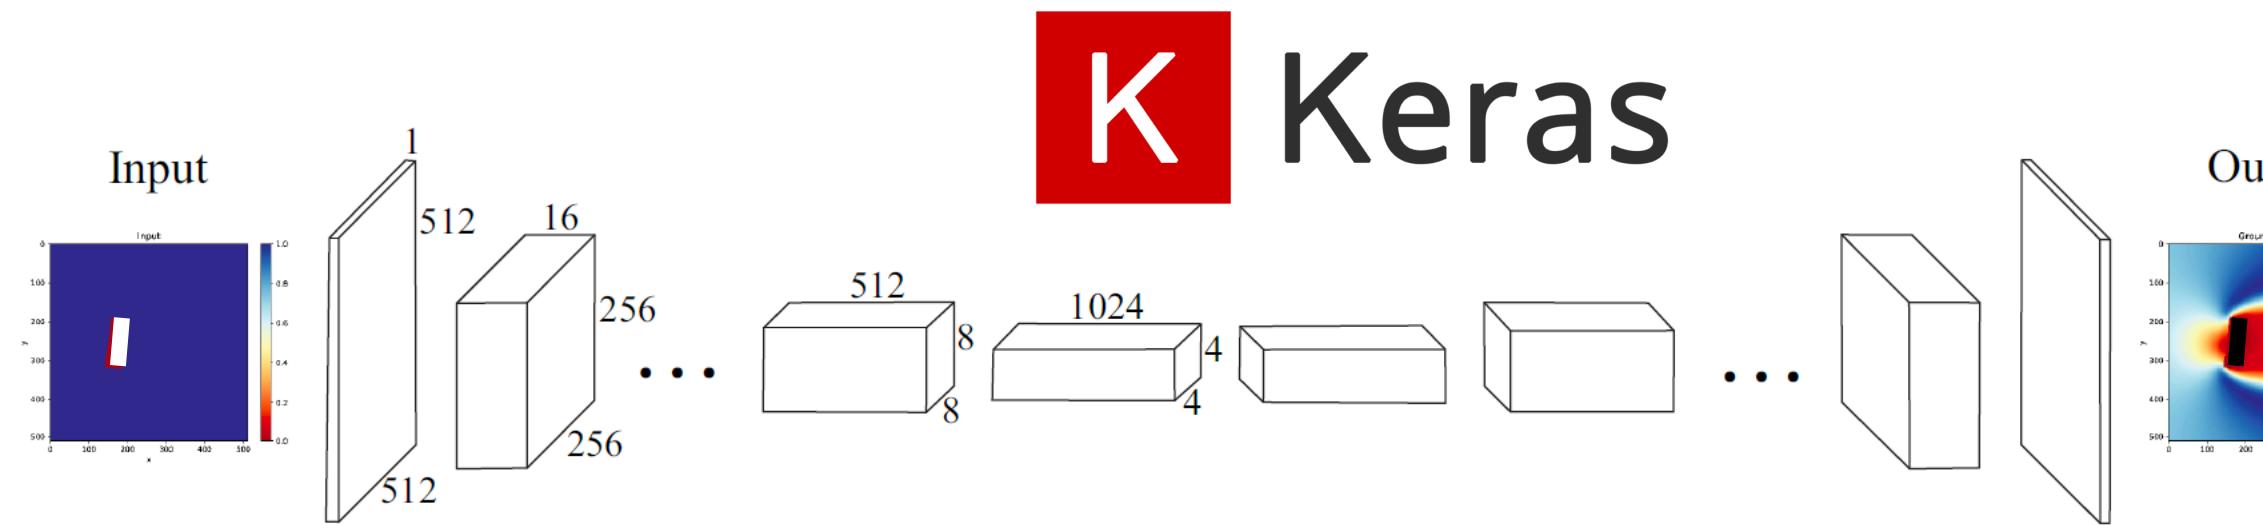
2D Geometries and Velocity Maps



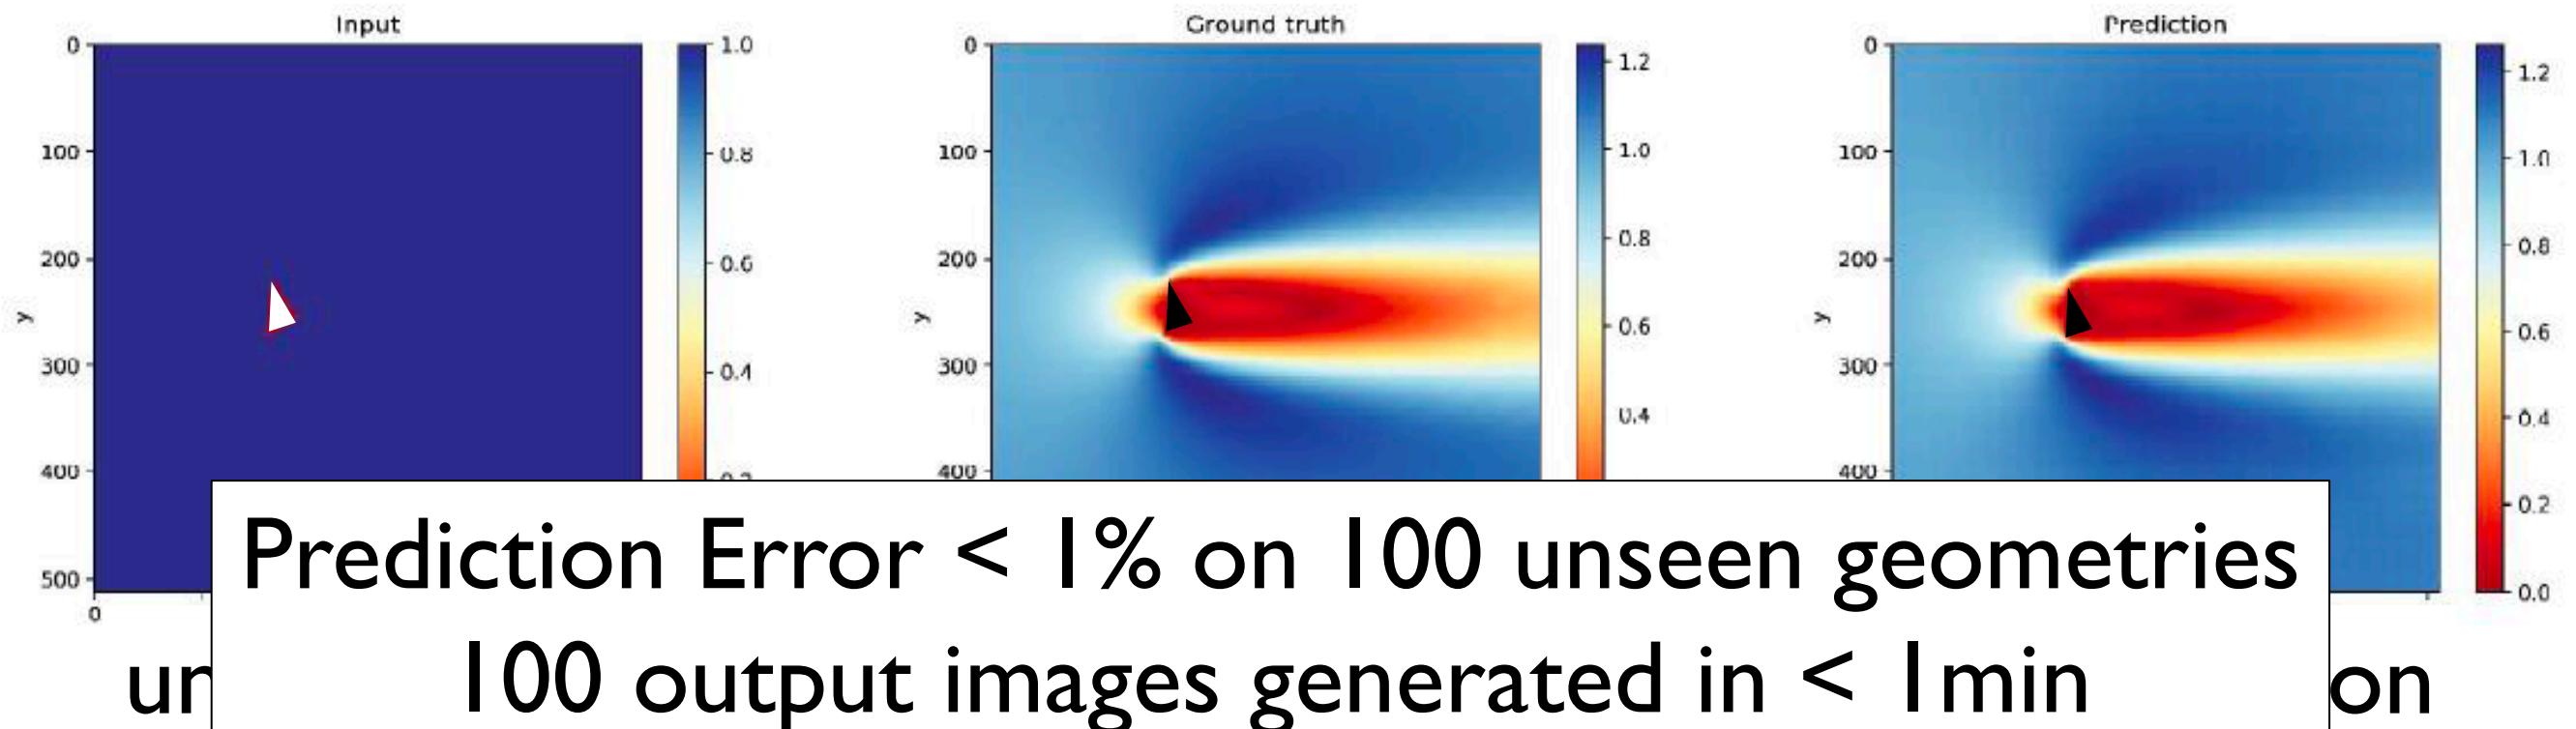
1000s of Fluent Computations

Machine Learning Training
(convolutional autoencoders)

K Keras

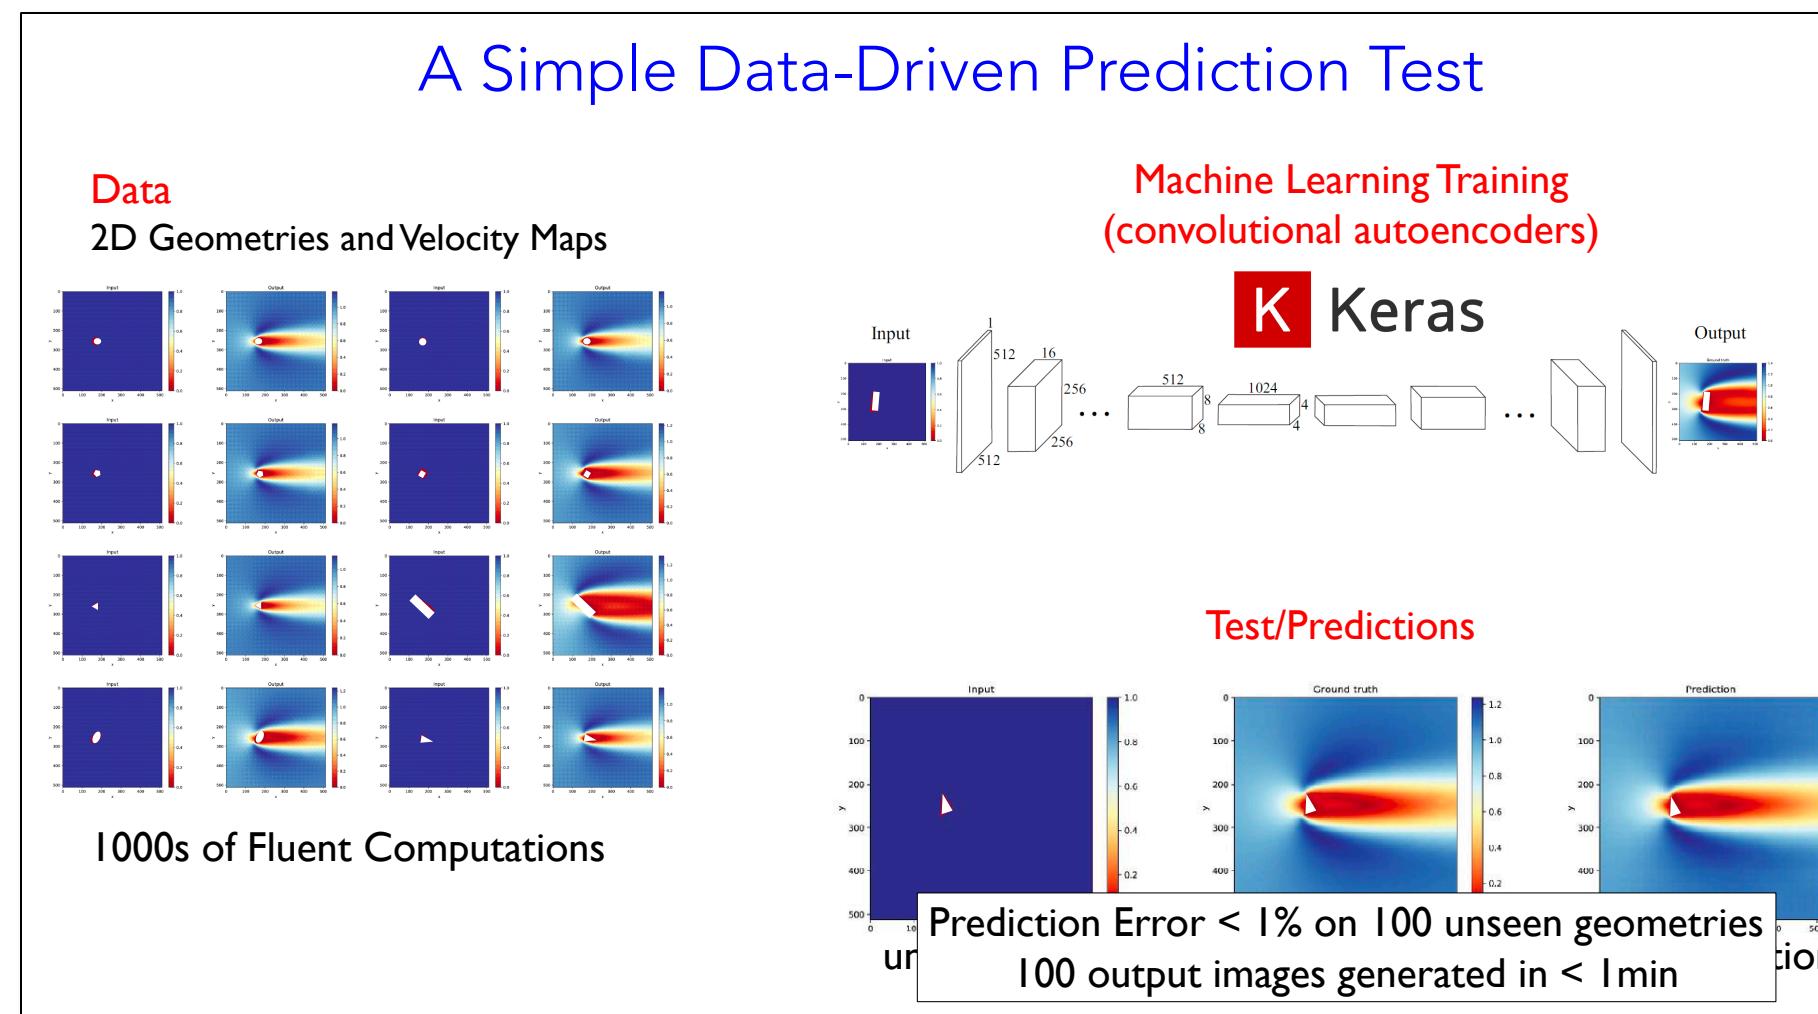


Test/Predictions

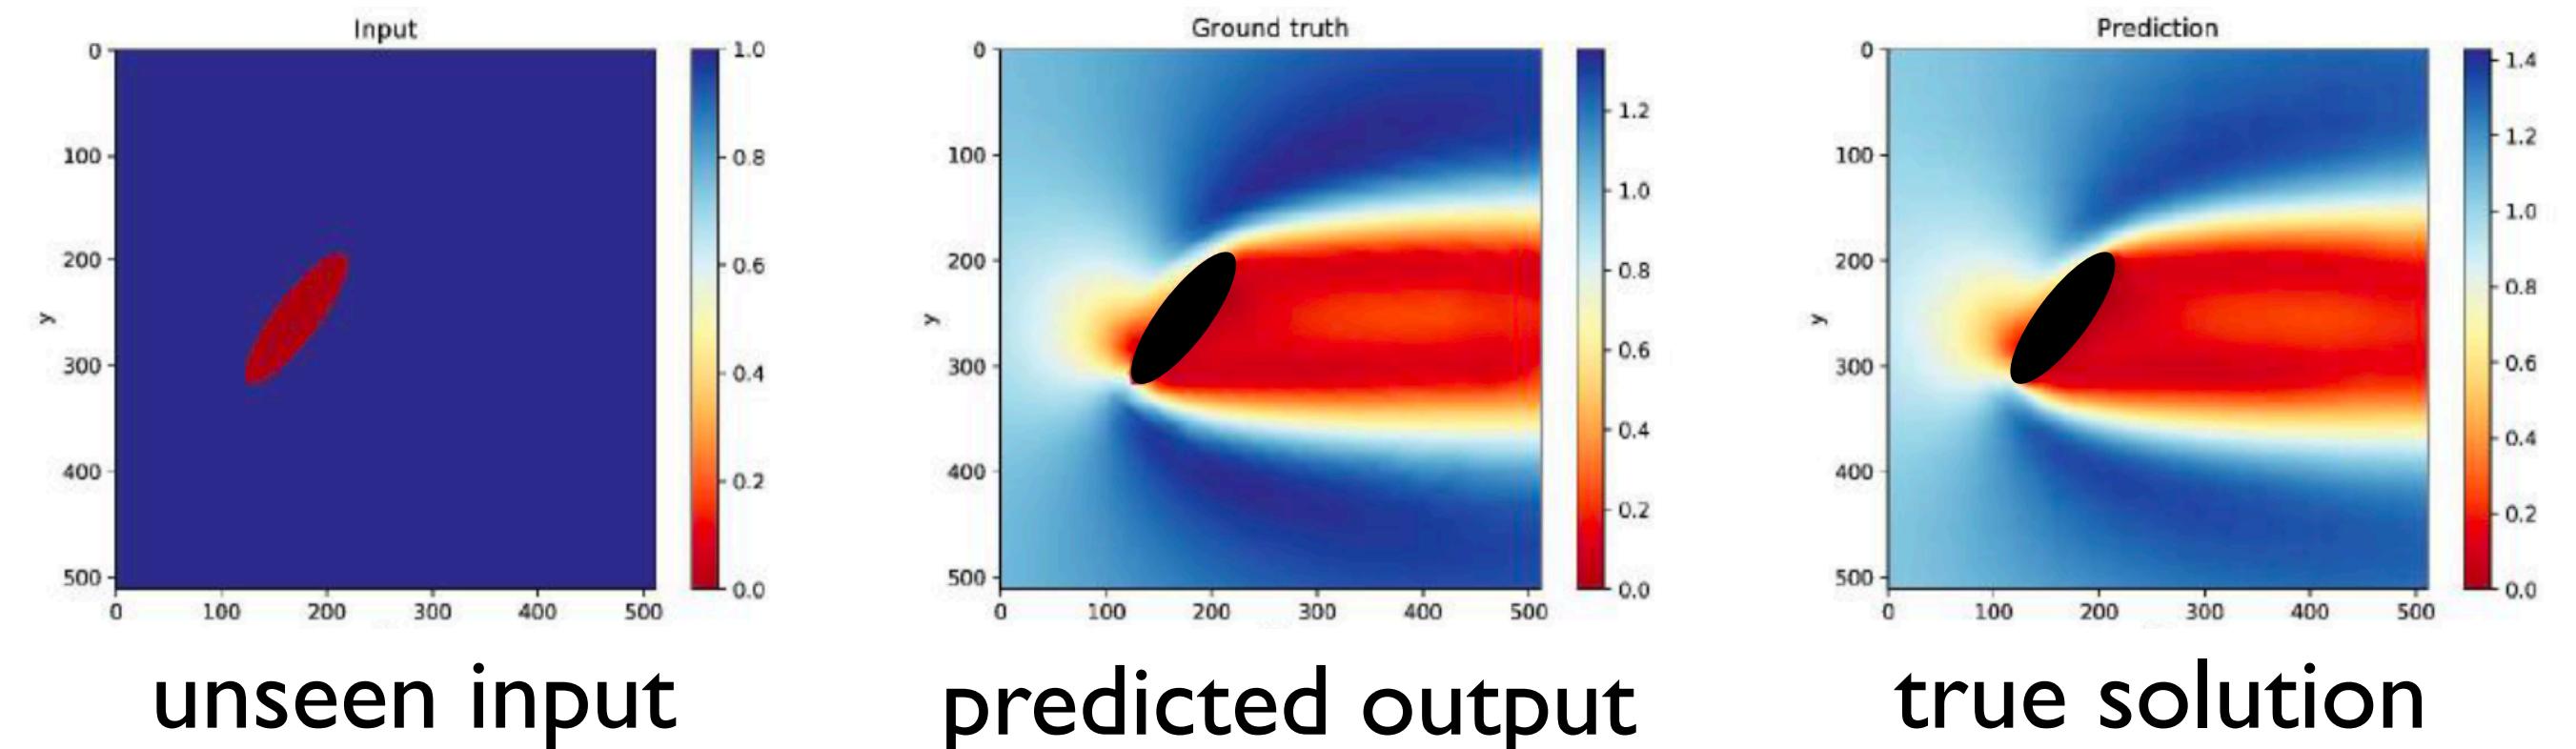




Synthetic Flow Simulations: flowfakes



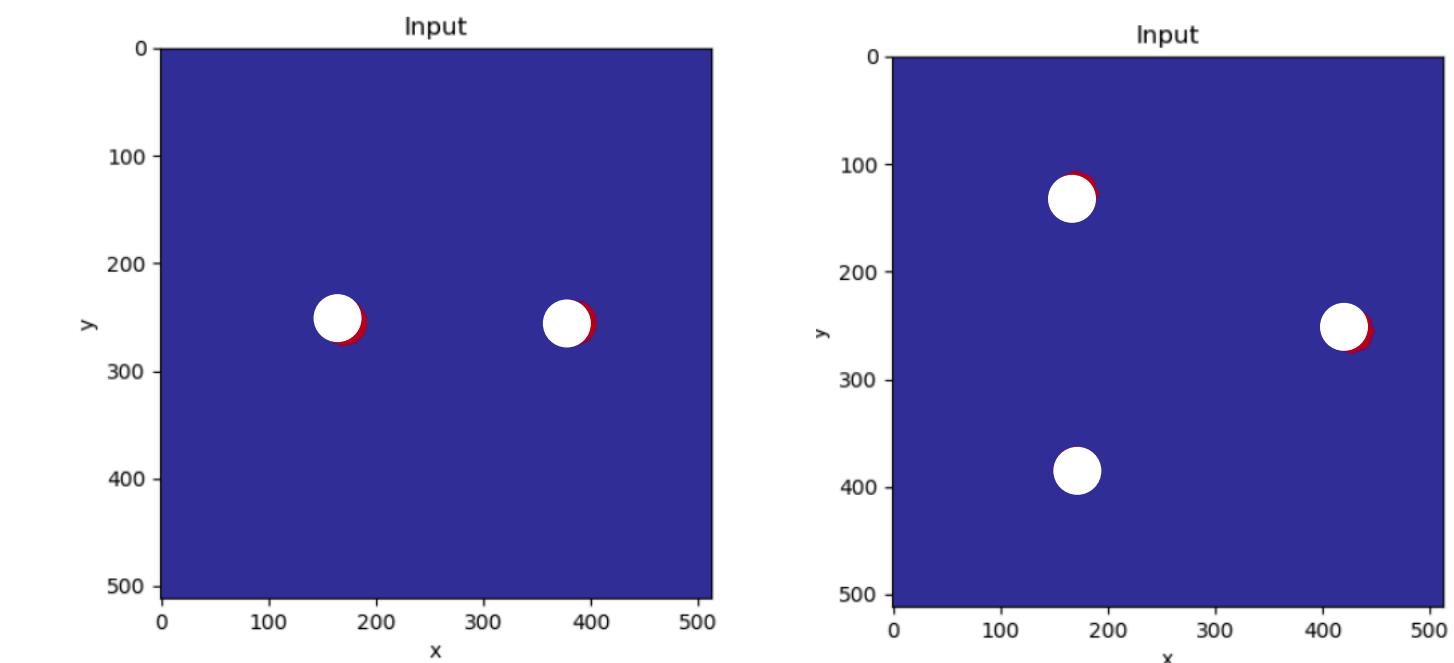
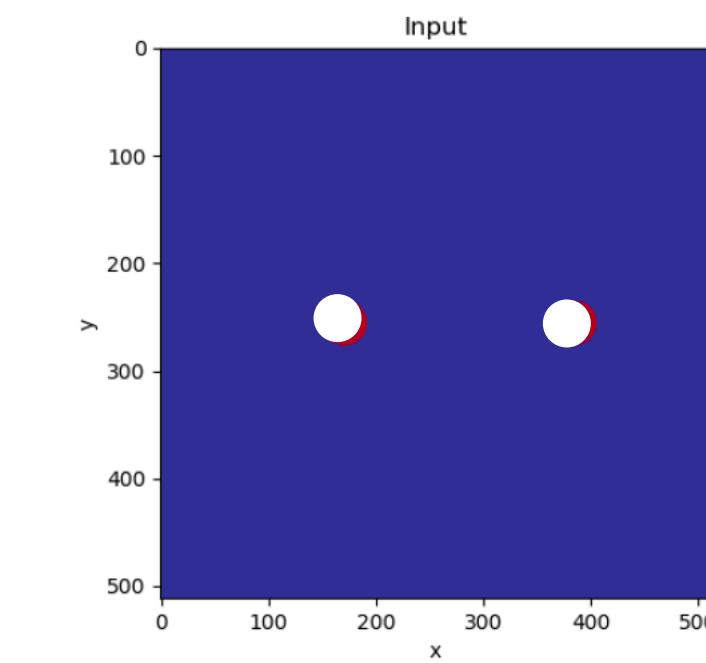
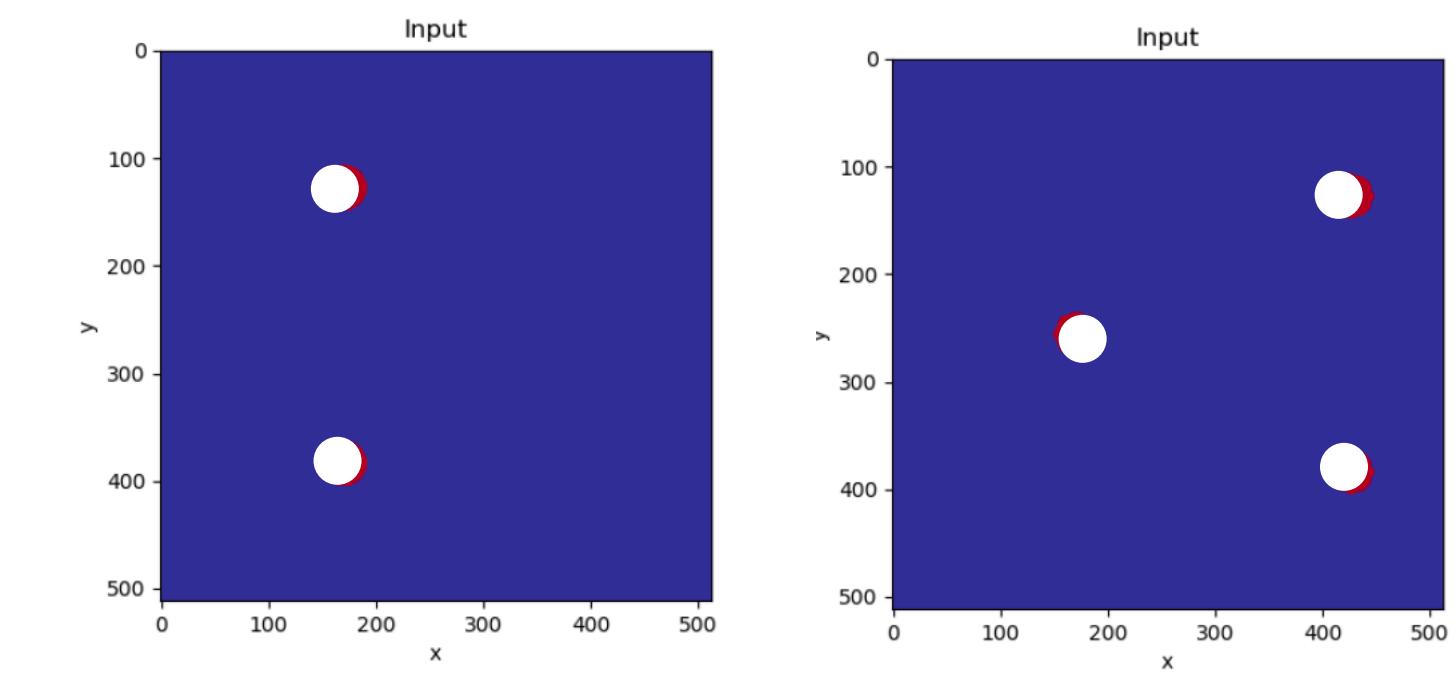
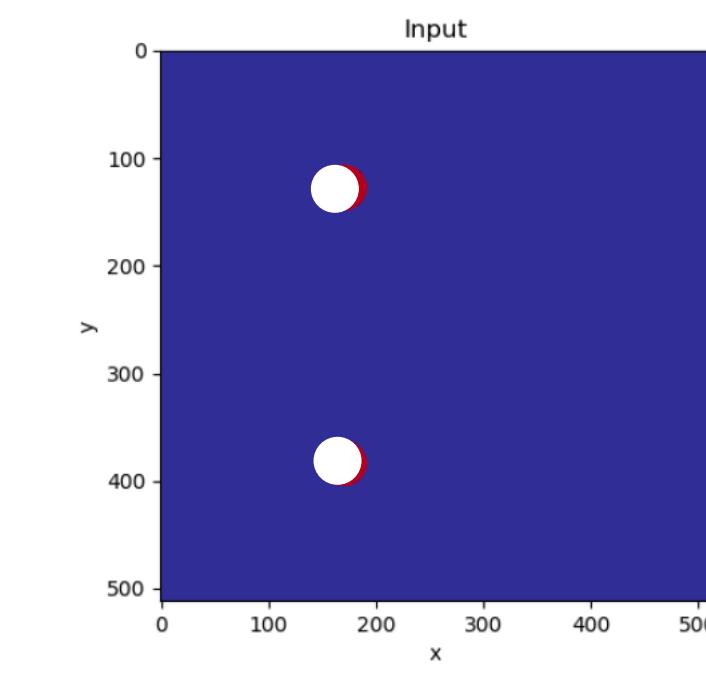
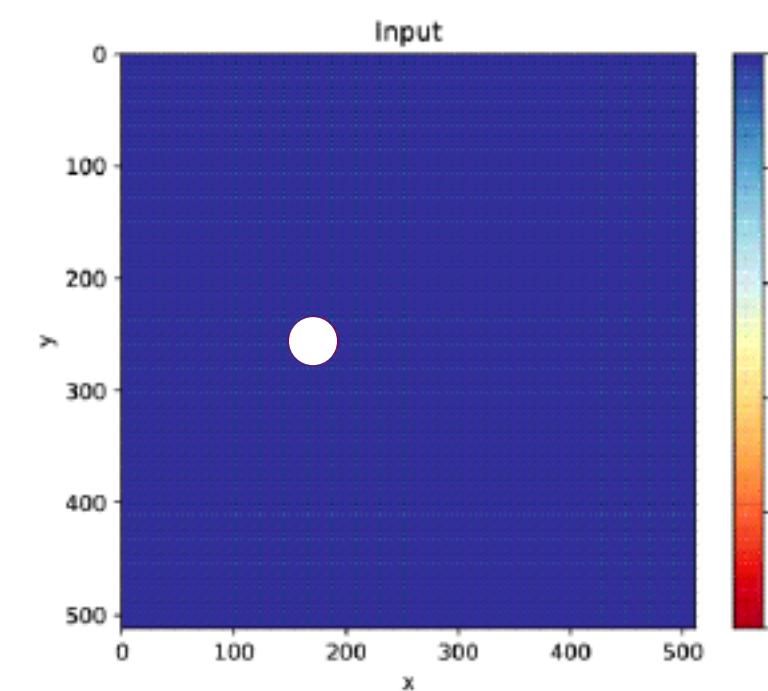
- Building the Dataset is **expensive** and requires extensive and robust automation
- Training the network is **expensive** and requires knowledge and expertise in selecting options
- Predictions accuracy is measured in average... in some cases the (local) **error** is as high as 5%





Synthetic Flow Simulations: **flowfakes**

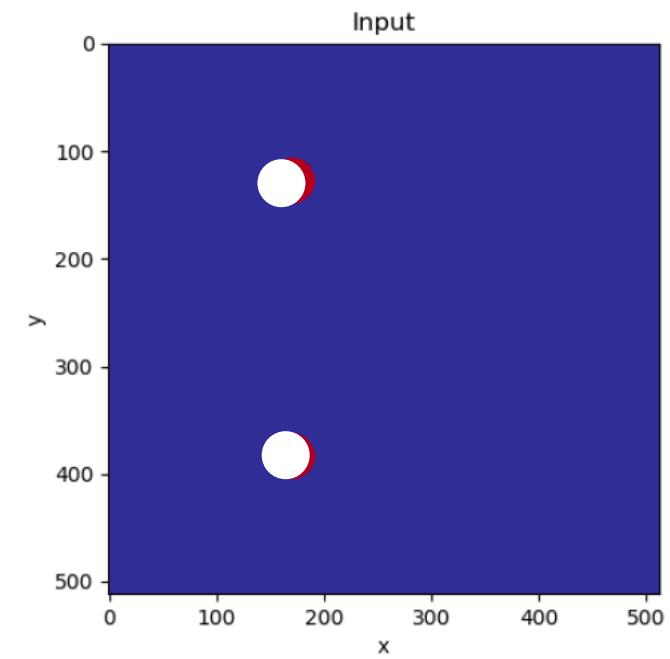
What is the accuracy expectation?



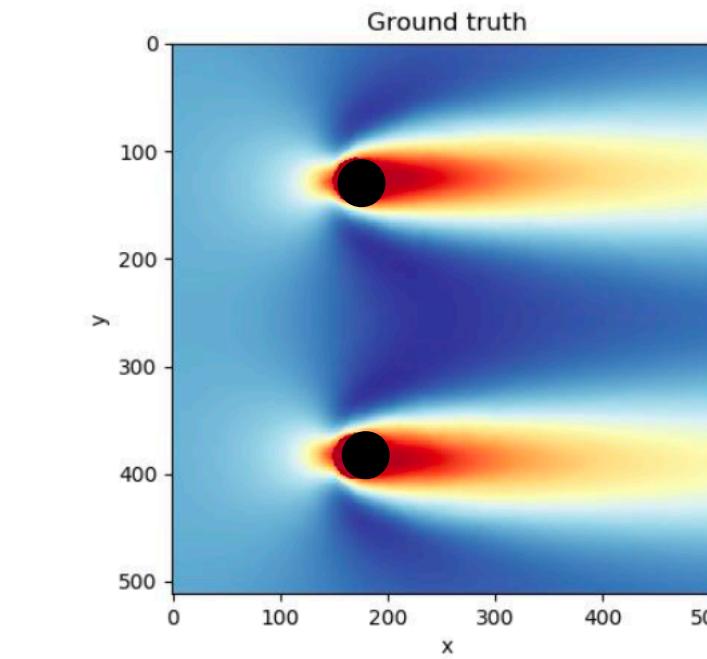


Synthetic Flow Simulations: flowfakes

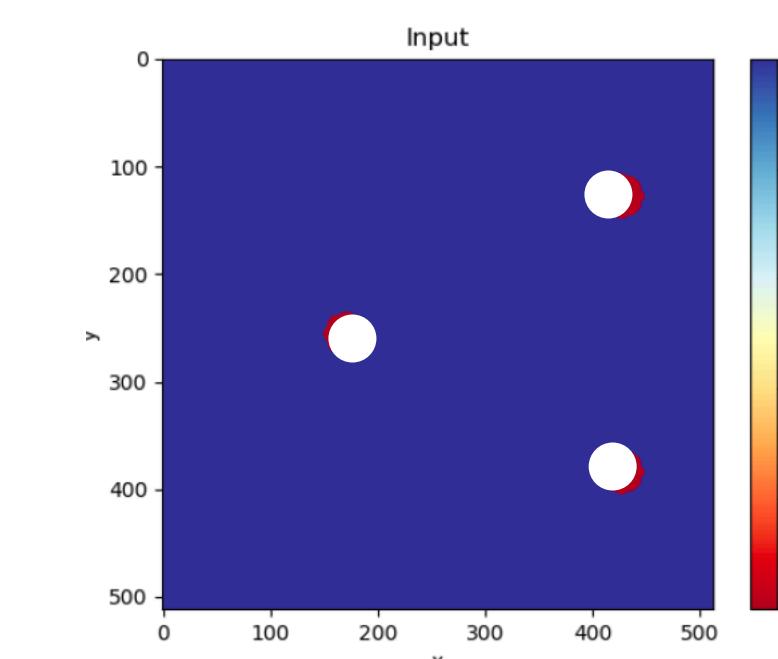
What is the accuracy expectation?



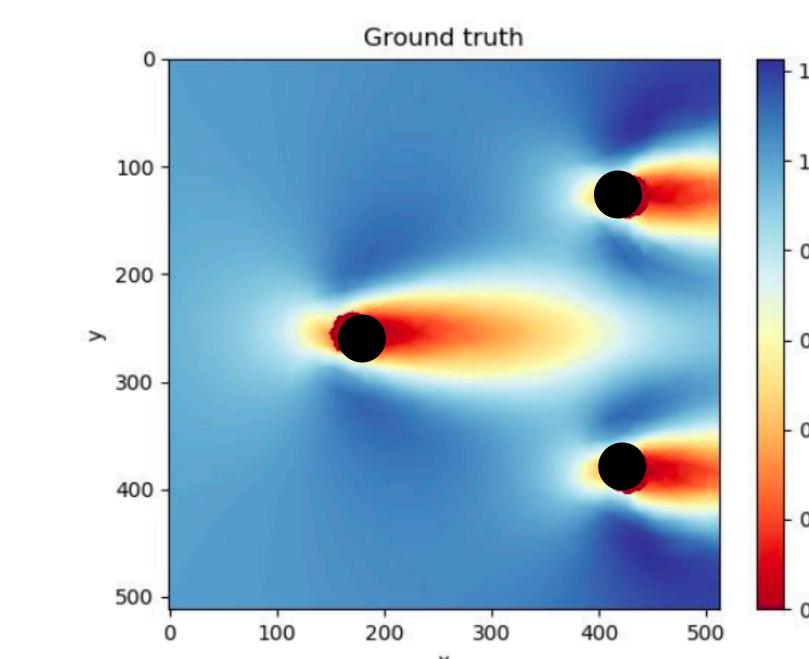
input



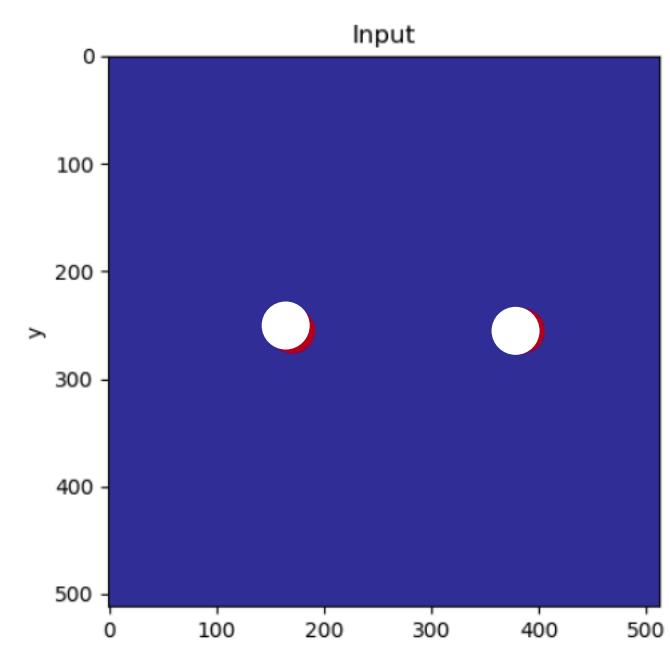
predicted output



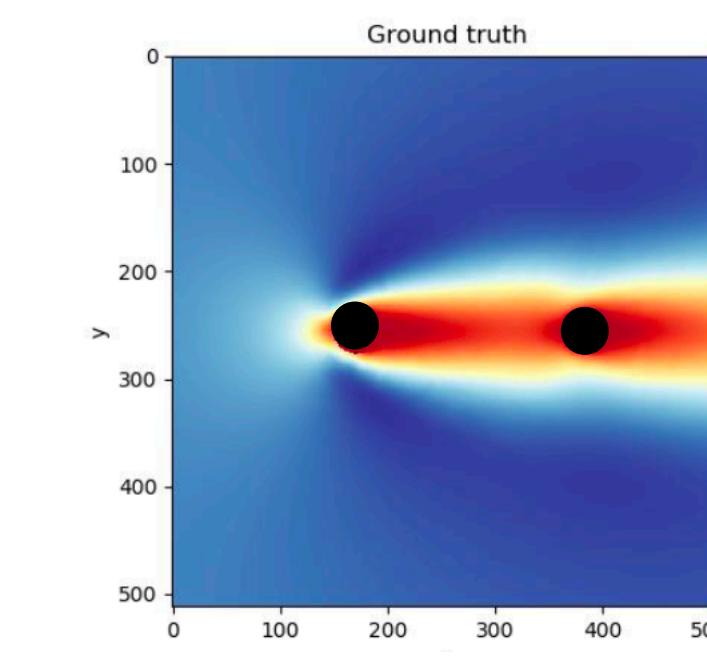
input



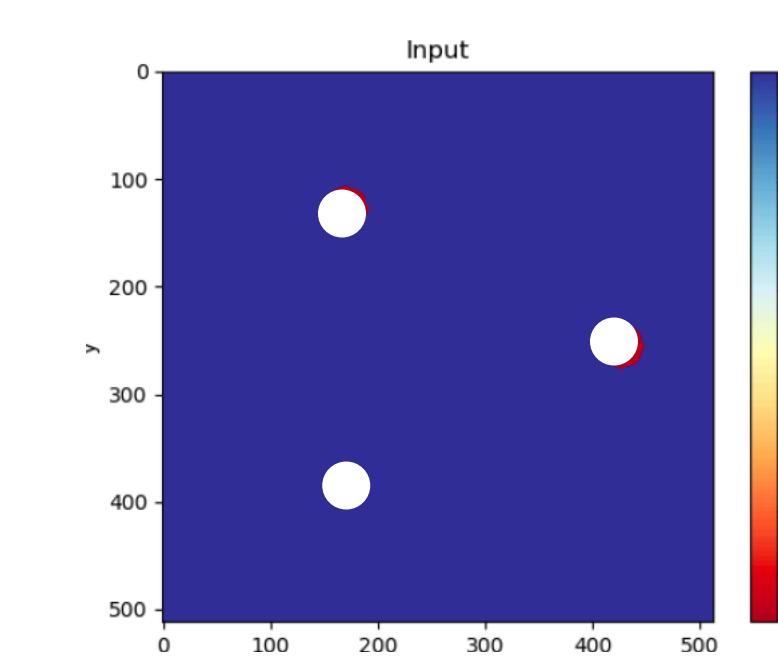
predicted output



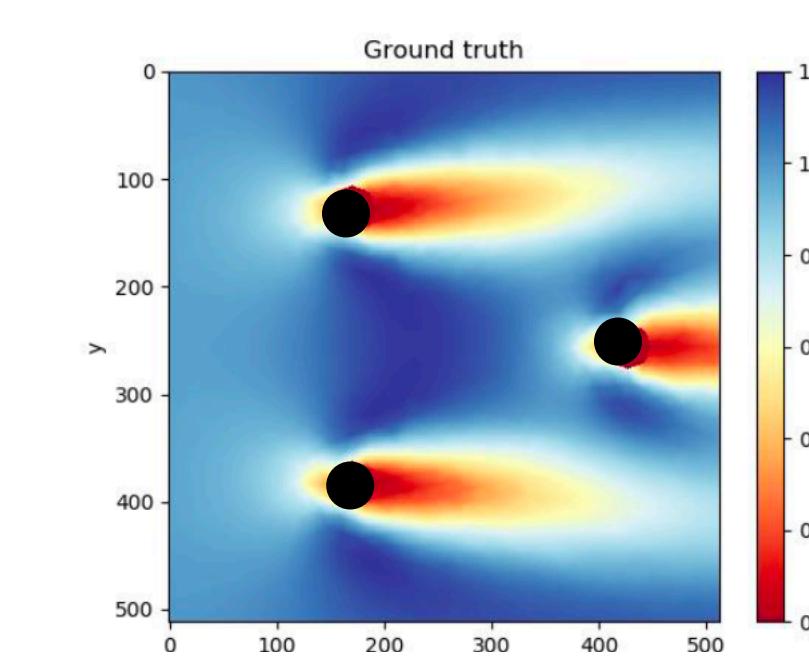
input



predicted output



input

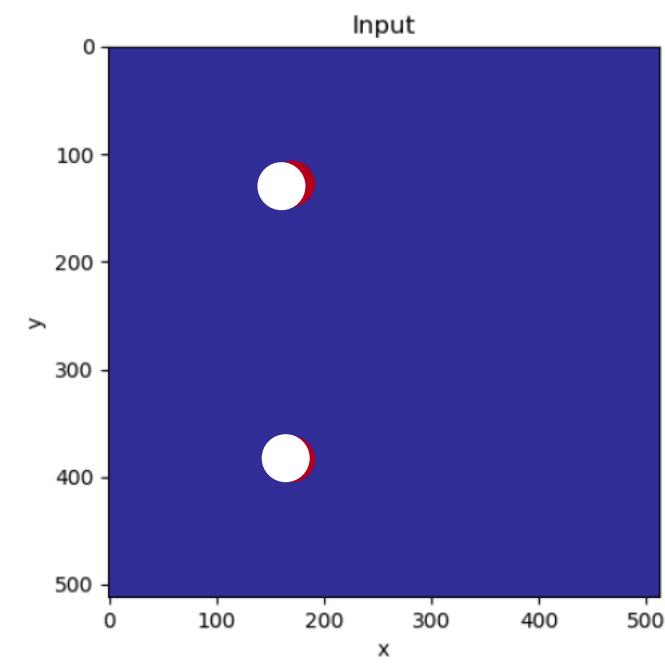


predicted output

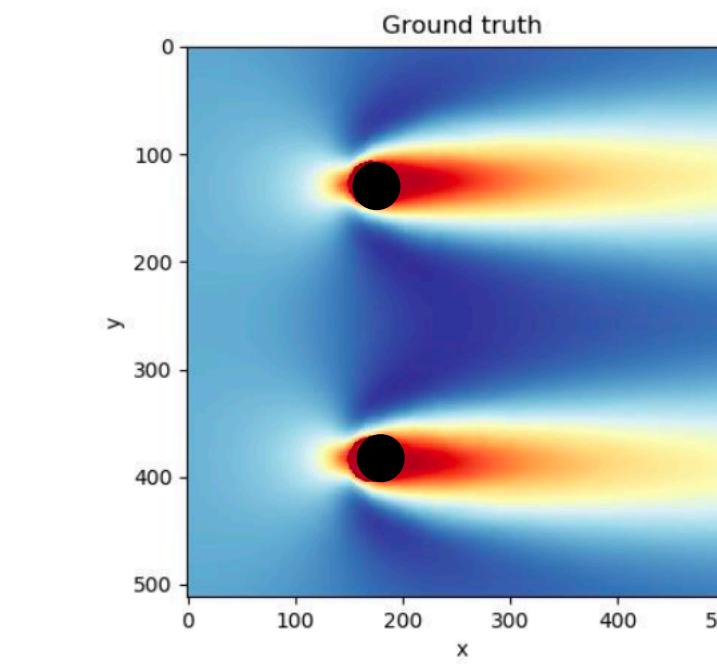


Synthetic Flow Simulations: flowfakes

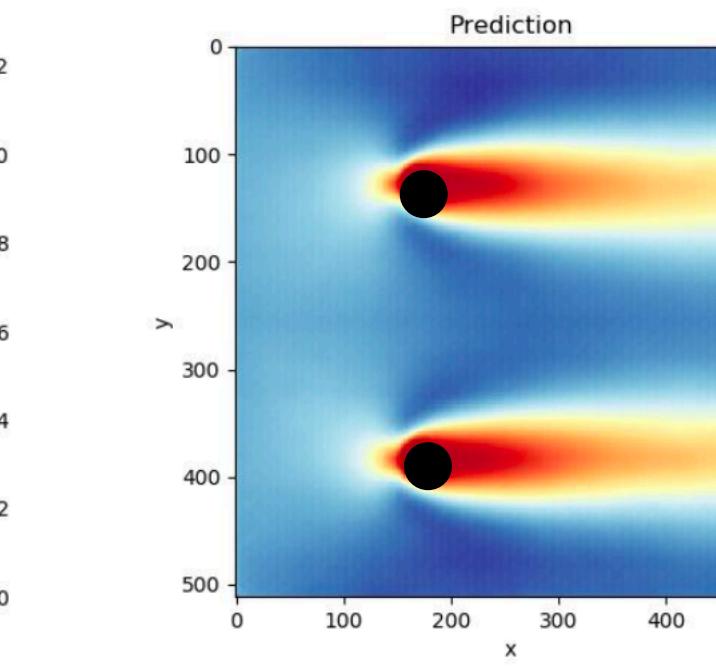
What is the accuracy expectation? ...difficult to say



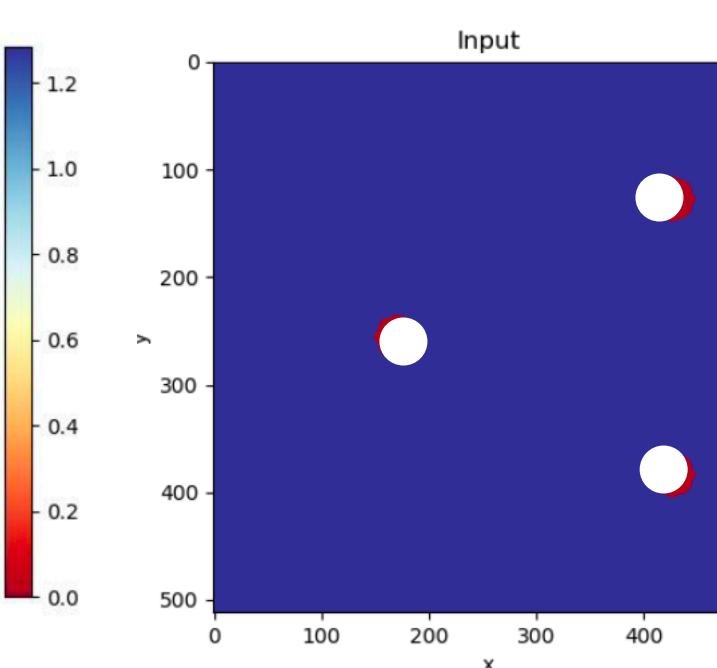
input



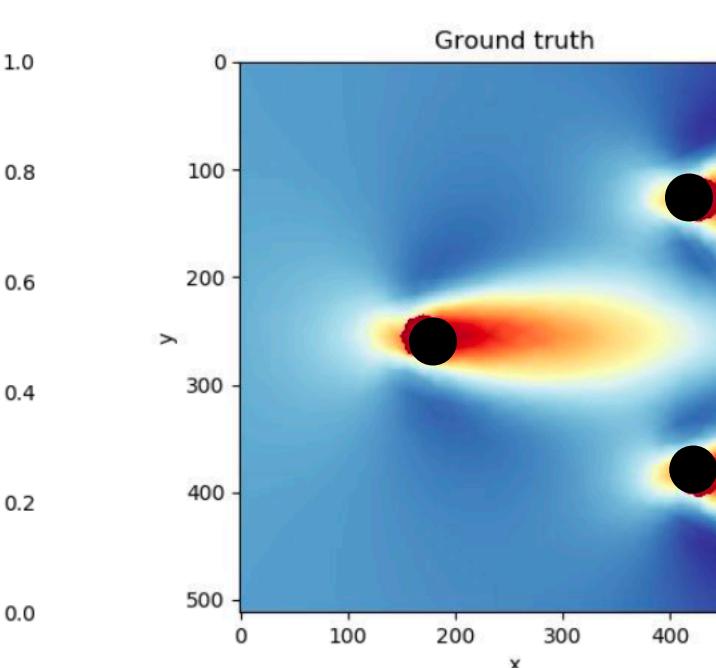
predicted output



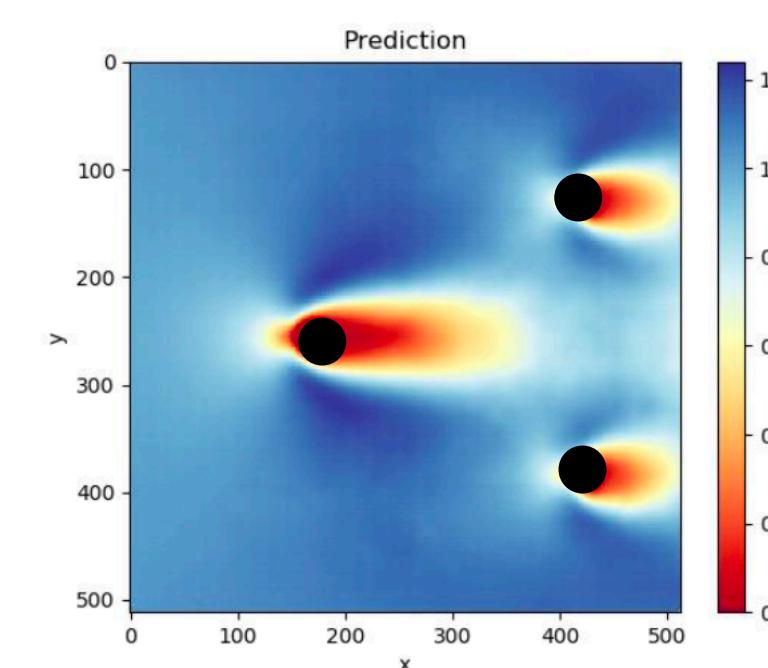
true solution



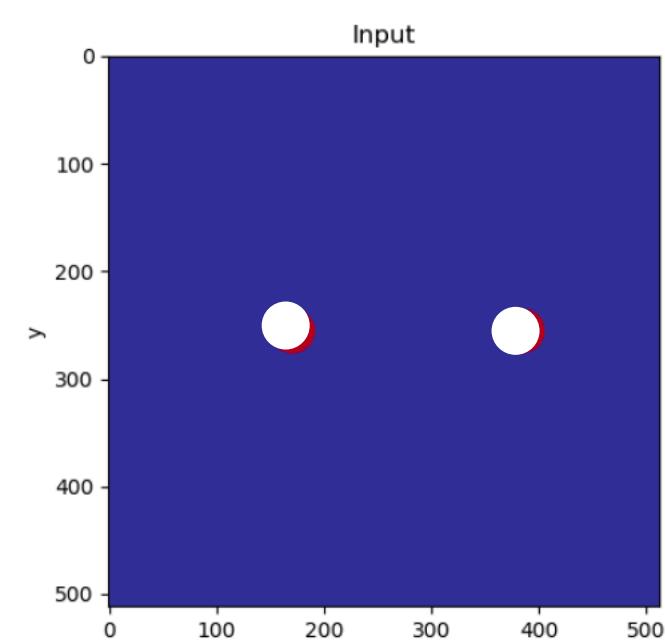
input



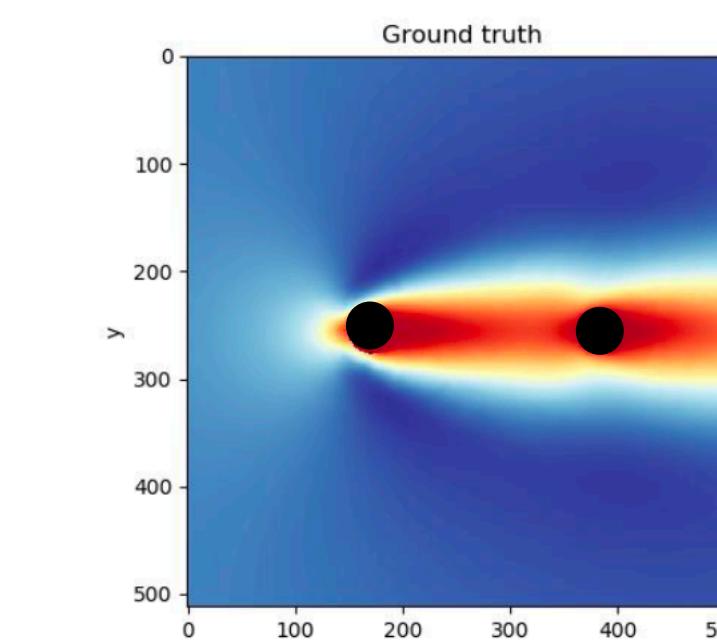
predicted output



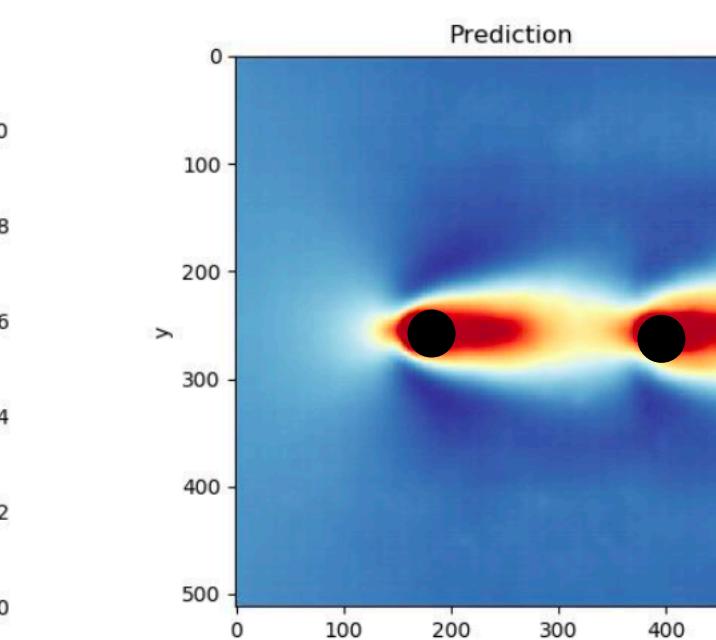
true solution



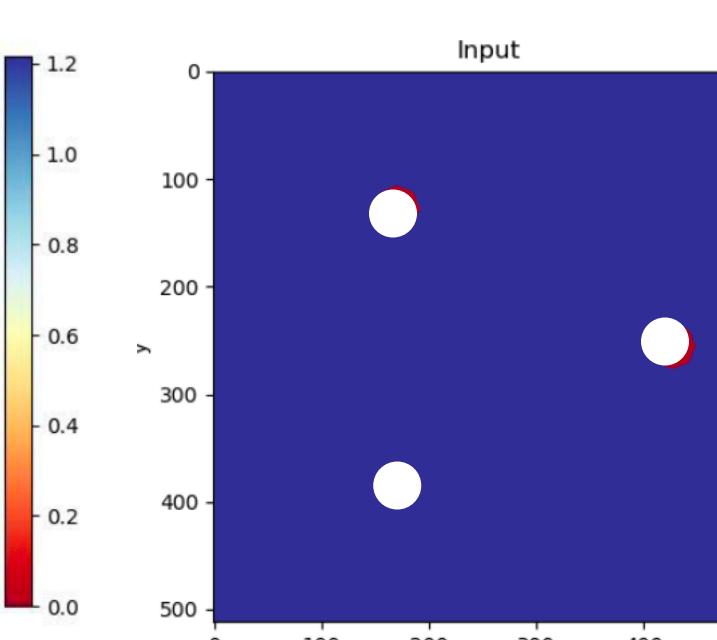
input



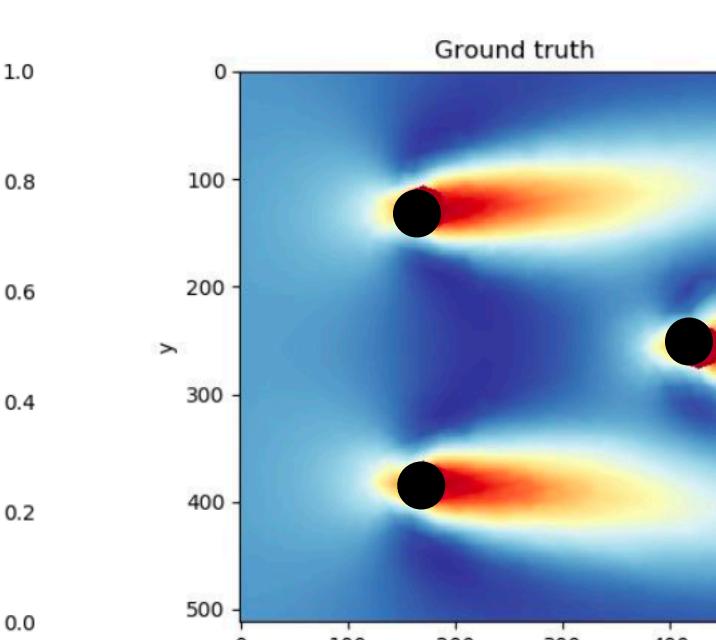
predicted output



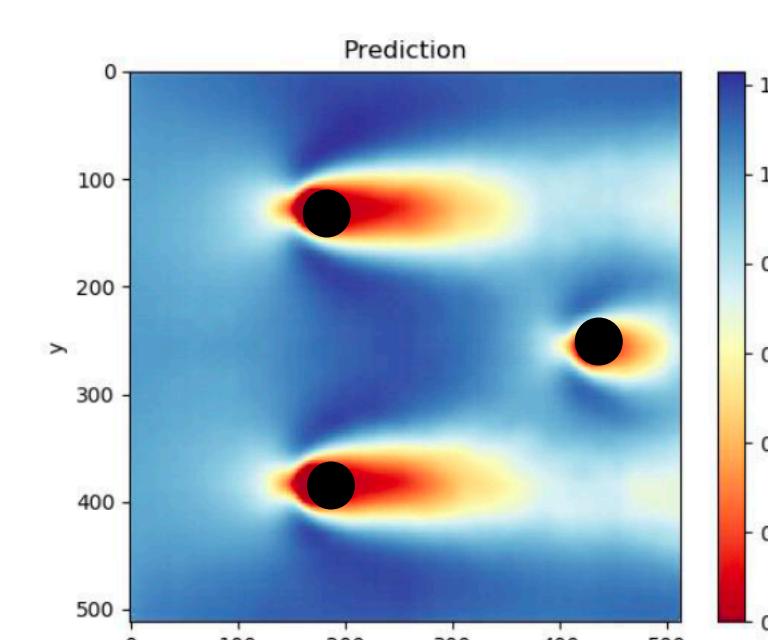
true solution



input



predicted output



true solution

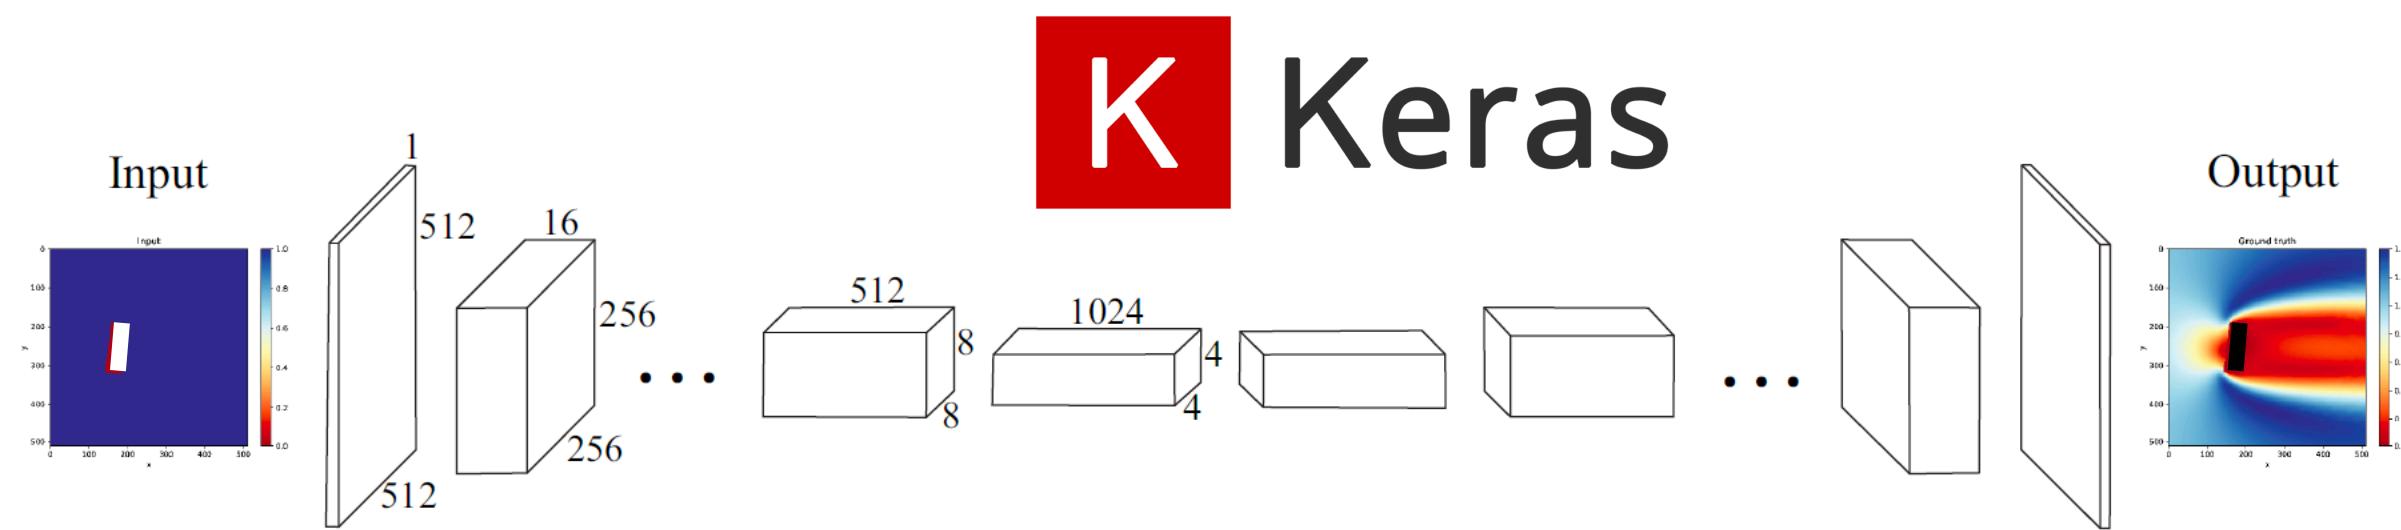
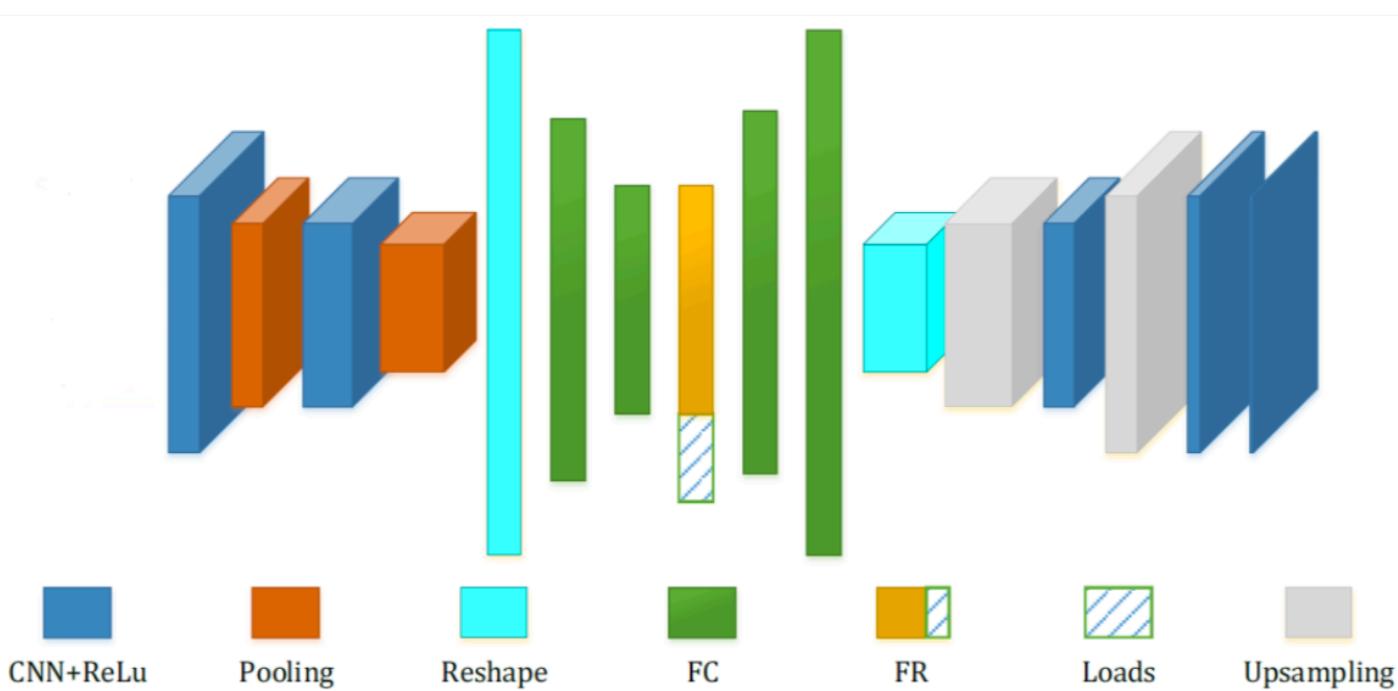


Synthetic Flow Simulations: flowfakes

CNN for Linear Stress Predictions

Nie, Jiang, Kara @ Carnegie Mellon, 2018

Machine Learning Training (convolutional autoencoders)



Some Details of the present DNN

- ~7000 solutions in dataset
- 39,158,080 hyper-parameters
- 1000 Epochs
- No fully connected layers
- No Max-pooling/Upsampling
- CNN with Leaky ReLU
- DCNN with ReLU

Layer#1 = Input
Layer#2 $C(f=16, k=4, s=2)$, Leaky ReLU
Layer#3 $C(f=32, k=4, s=2)$, Leaky ReLU, Batch Normalization
Layer#4 $C(f=64, k=4, s=2)$, Leaky ReLU, Batch Normalization
Layer#5 $C(f=128, k=4, s=2)$, Leaky ReLU, Batch Normalization
Layer#6 $C(f=256, k=4, s=2)$, Leaky ReLU, Batch Normalization
Layer#7 $C(f=512, k=4, s=2)$, Leaky ReLU, Batch Normalization
Layer#8 $C(f=1024, k=4, s=2)$, Leaky ReLU, Batch Normalization
Layer#9 $CT(f=1024, k=4, s=1)$, ReLU, Batch Normalization
Layer#10 $CT(f=512, k=4, s=2)$, ReLU, Batch Normalization
Layer#11 $CT(f=256, k=4, s=2)$, ReLU, Batch Normalization
Layer#12 $CT(f=128, k=4, s=2)$, ReLU, Batch Normalization
Layer#13 $CT(f=64, k=4, s=2)$, ReLU, Batch Normalization
Layer#14 $CT(f=32, k=4, s=2)$, ReLU, Batch Normalization
Layer#15 $CT(f=16, k=4, s=2)$, ReLU, Batch Normalization
Layer#16 Output = $CT(f=1, k=4, s=2)$, ReLU

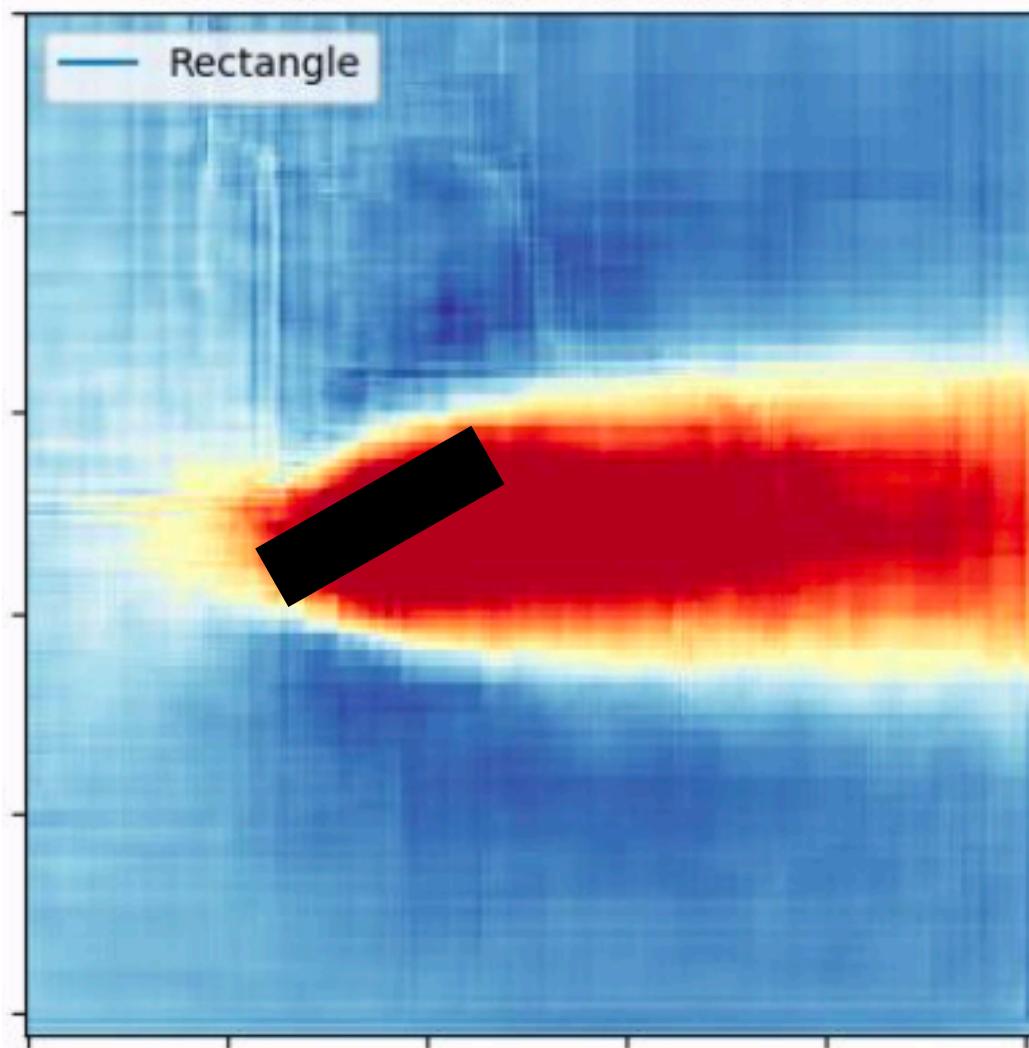
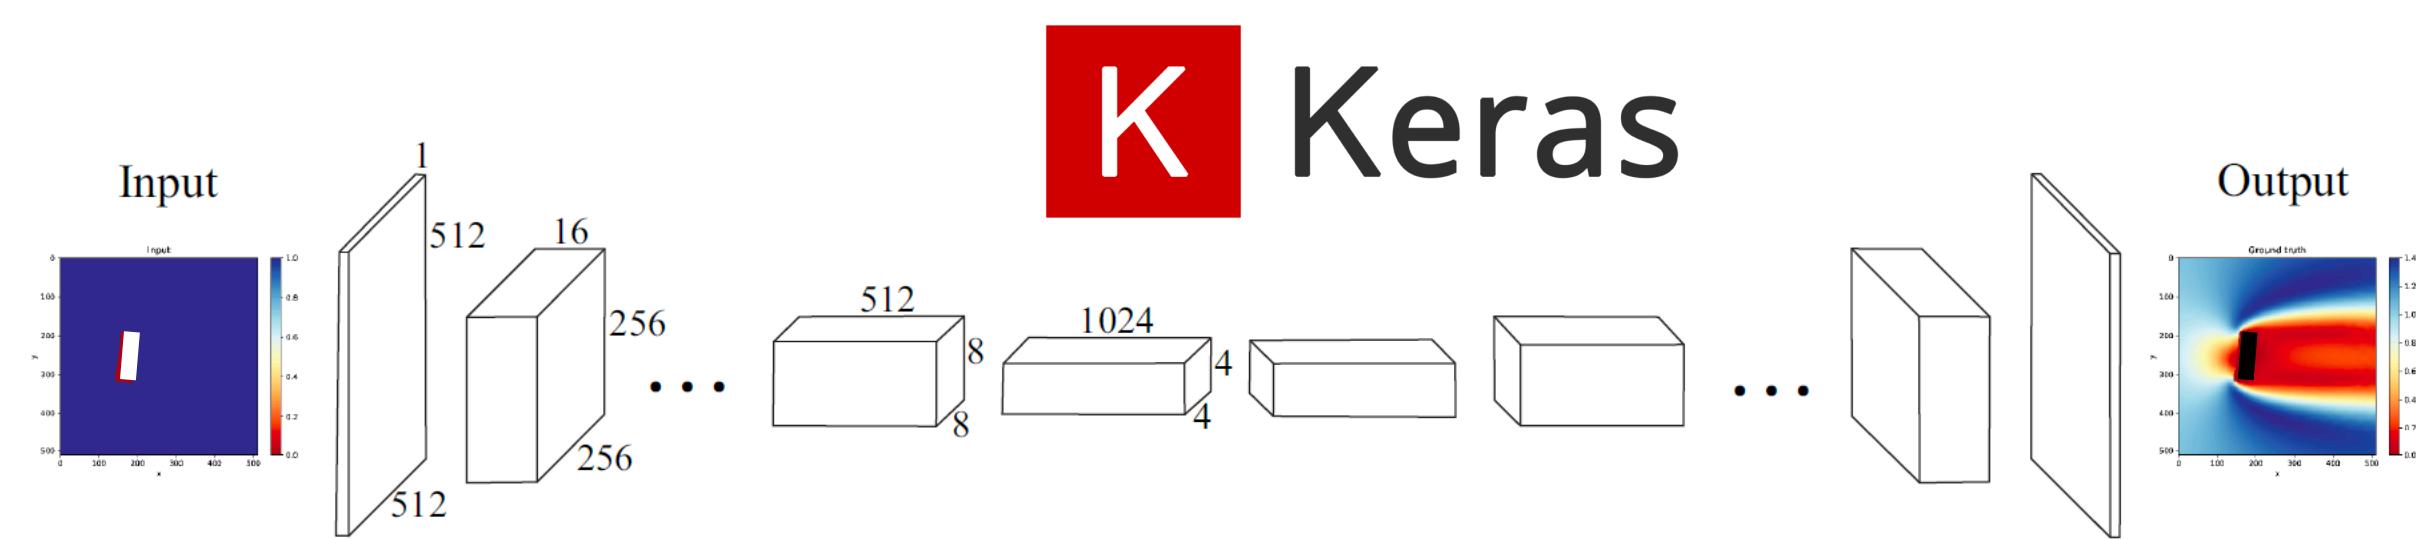
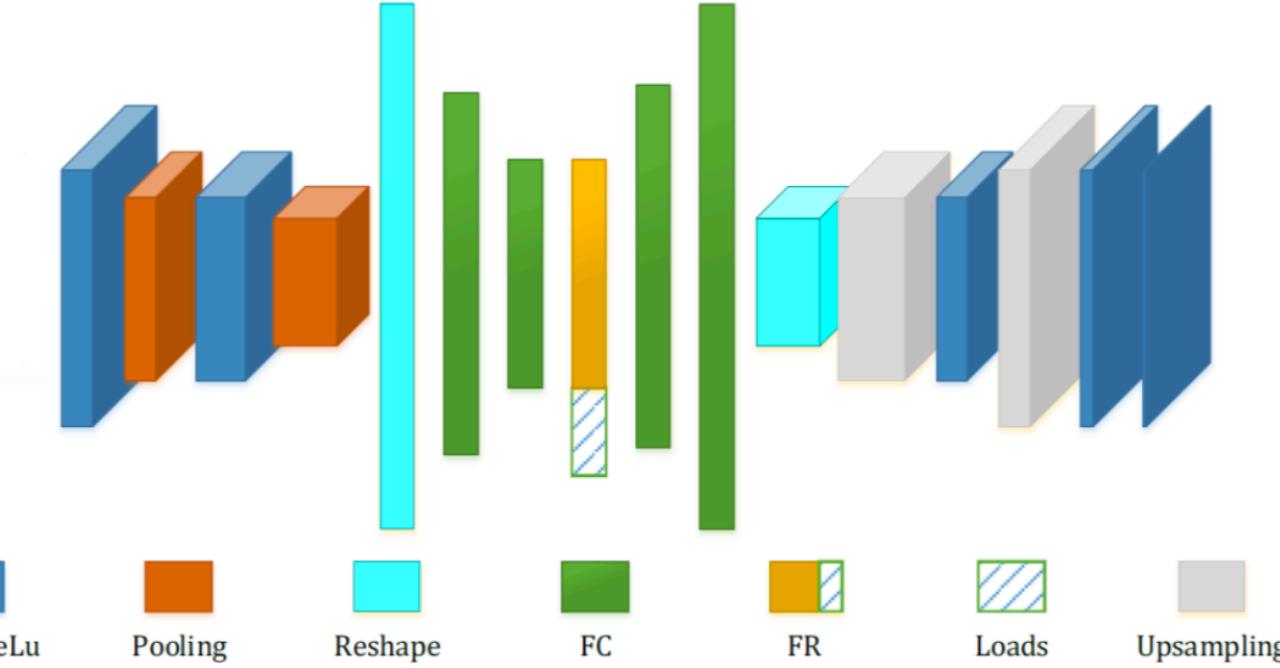


Synthetic Flow Simulations: flowfakes

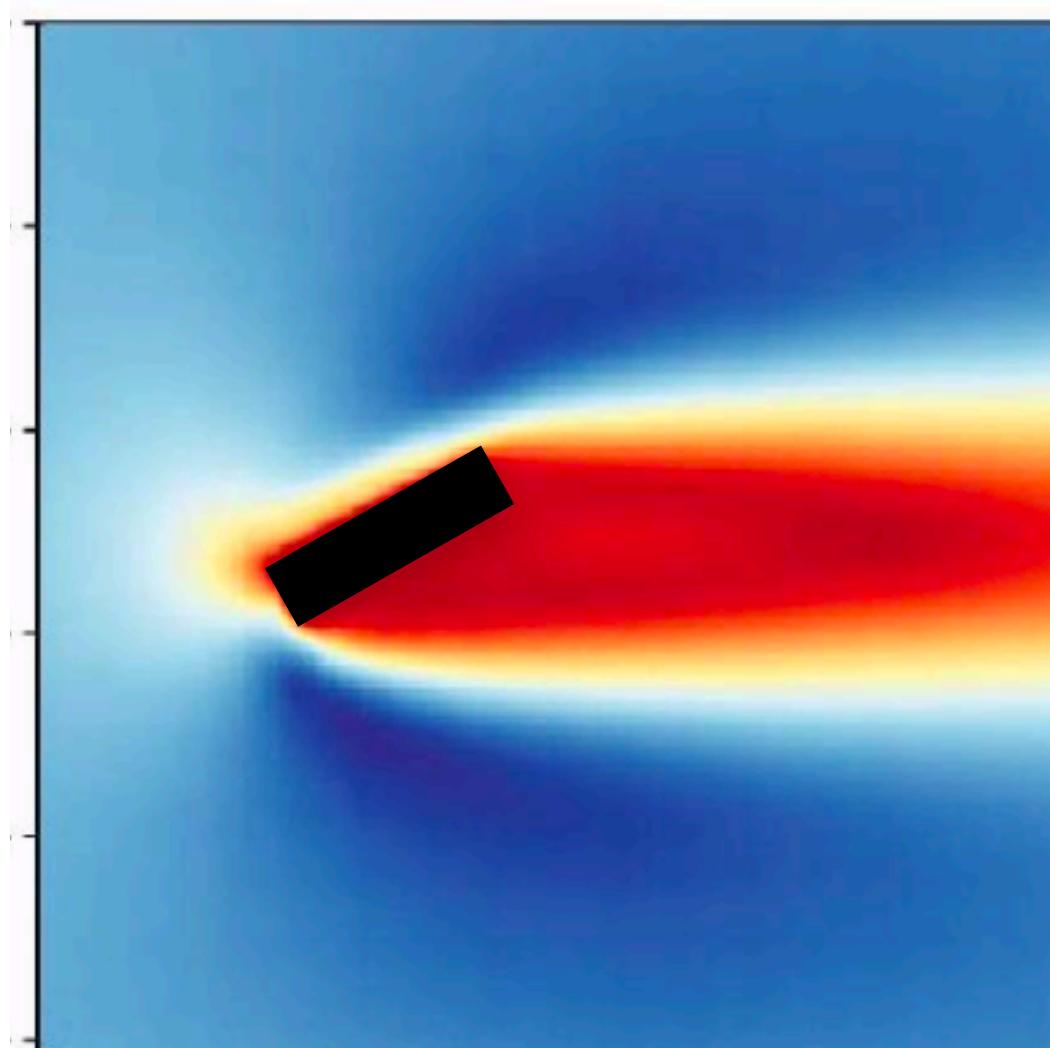
CNN for Linear Stress Predictions

Nie, Jiang, Kara @ Carnegie Mellon, 2018

Machine Learning Training
(convolutional autoencoders)



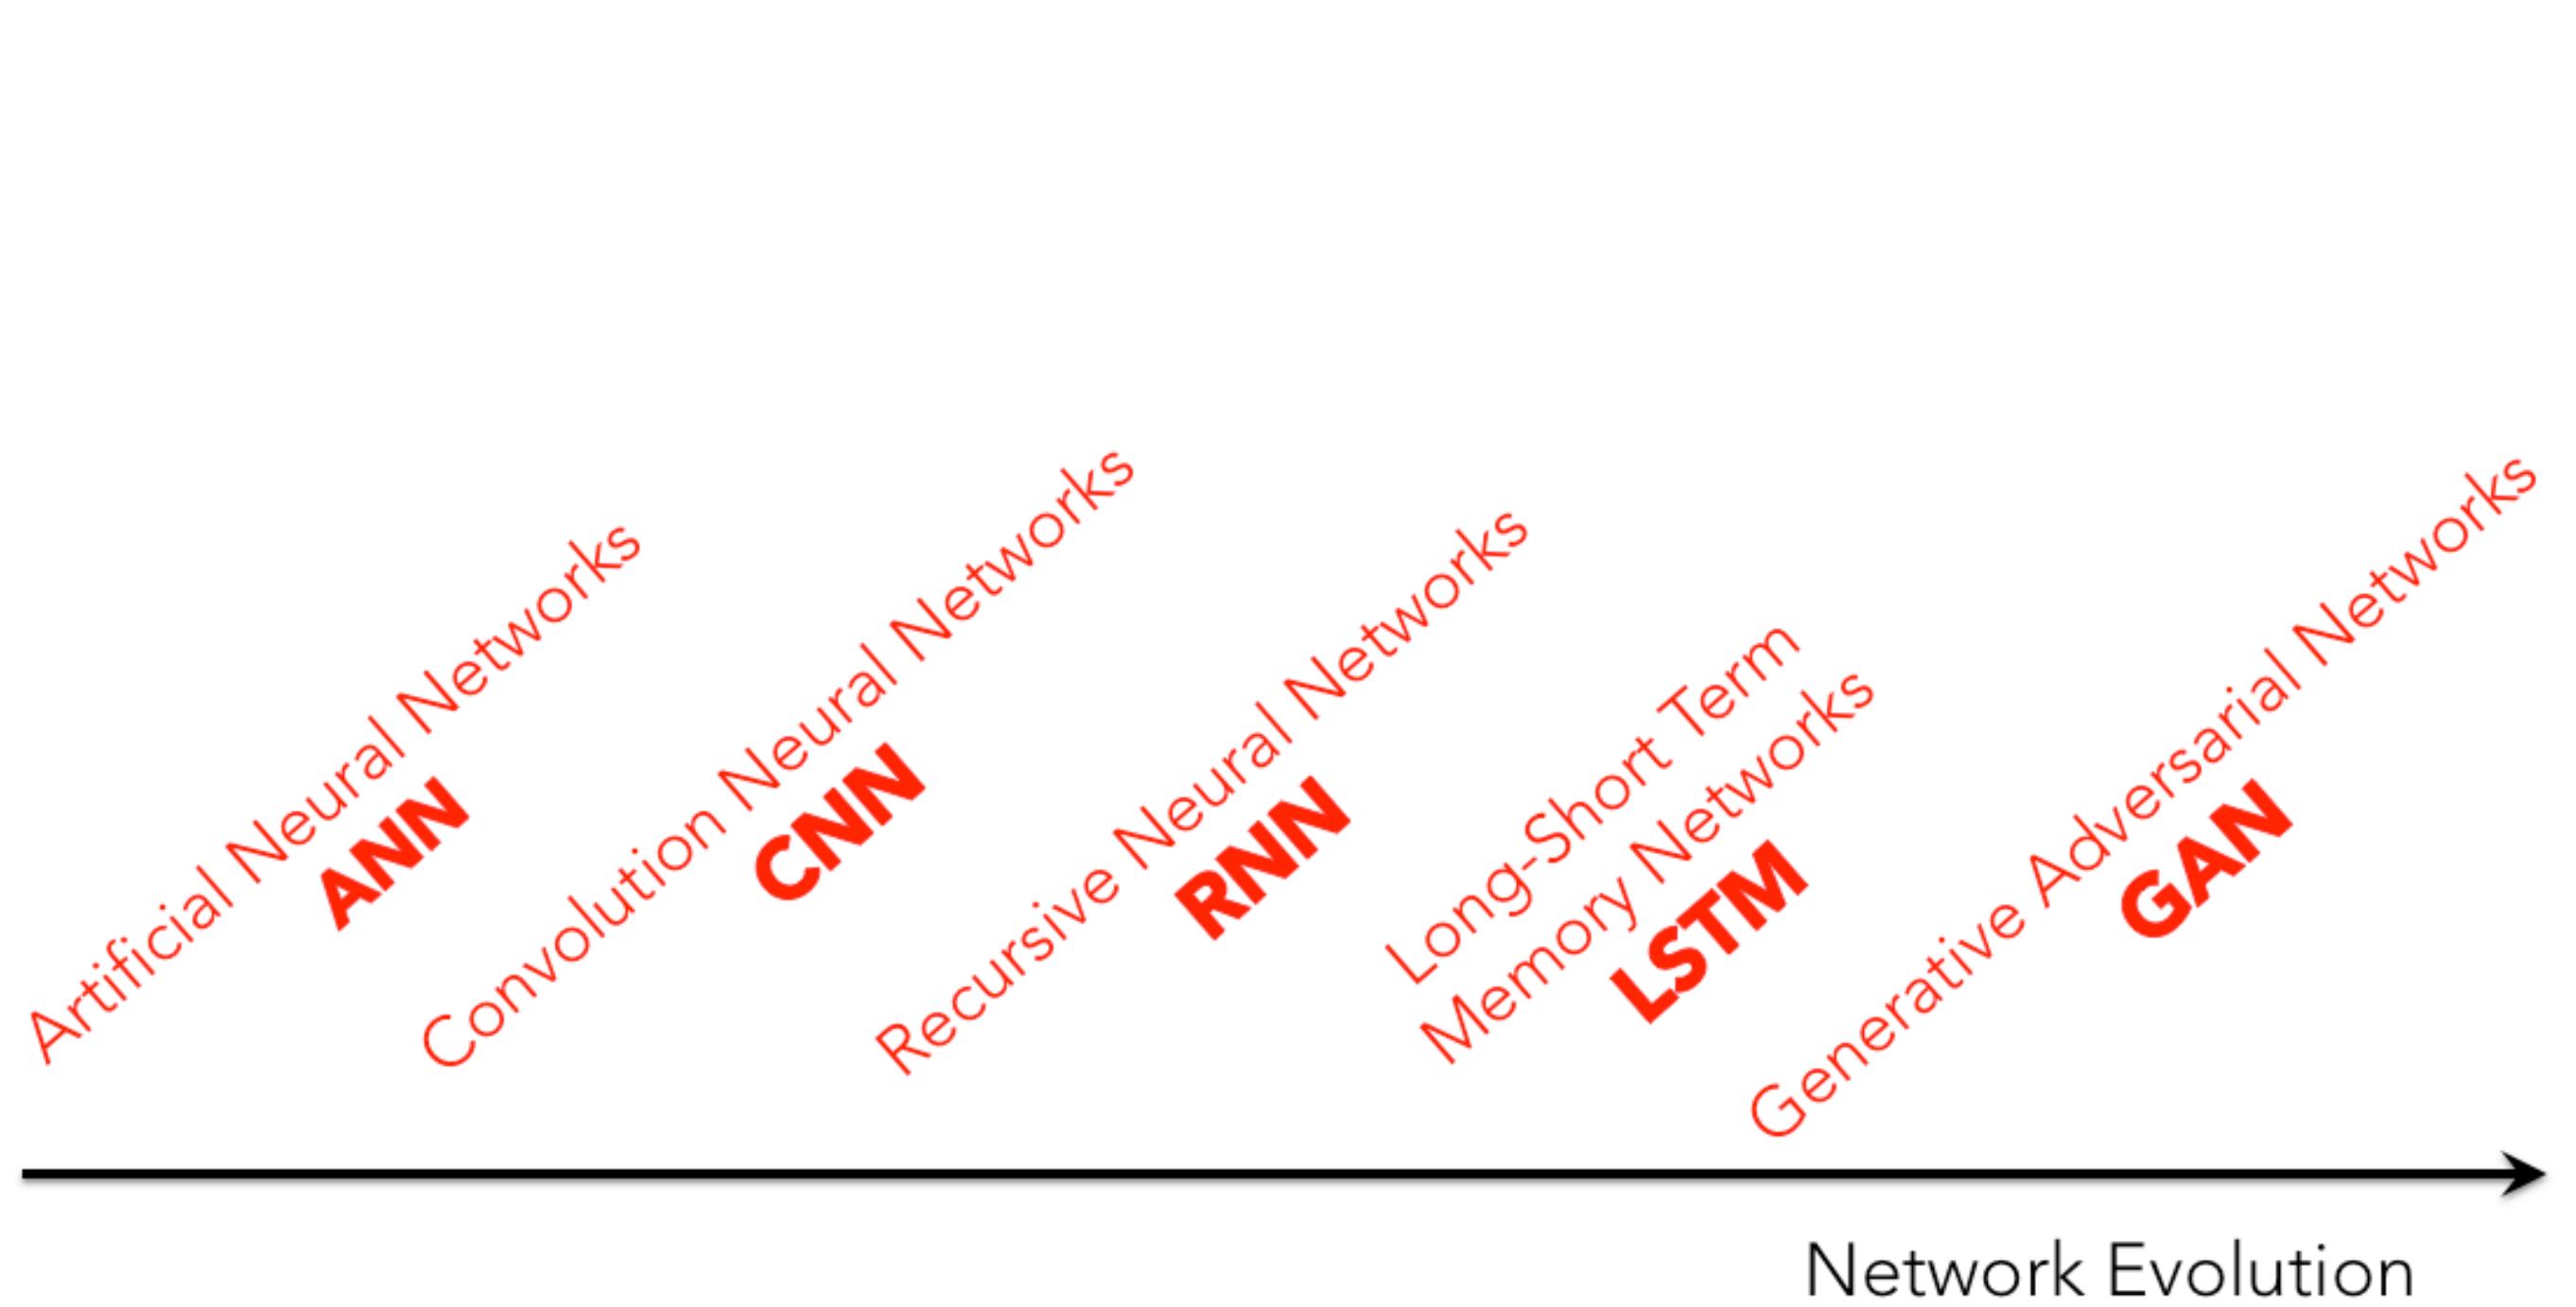
@ fixed training cost



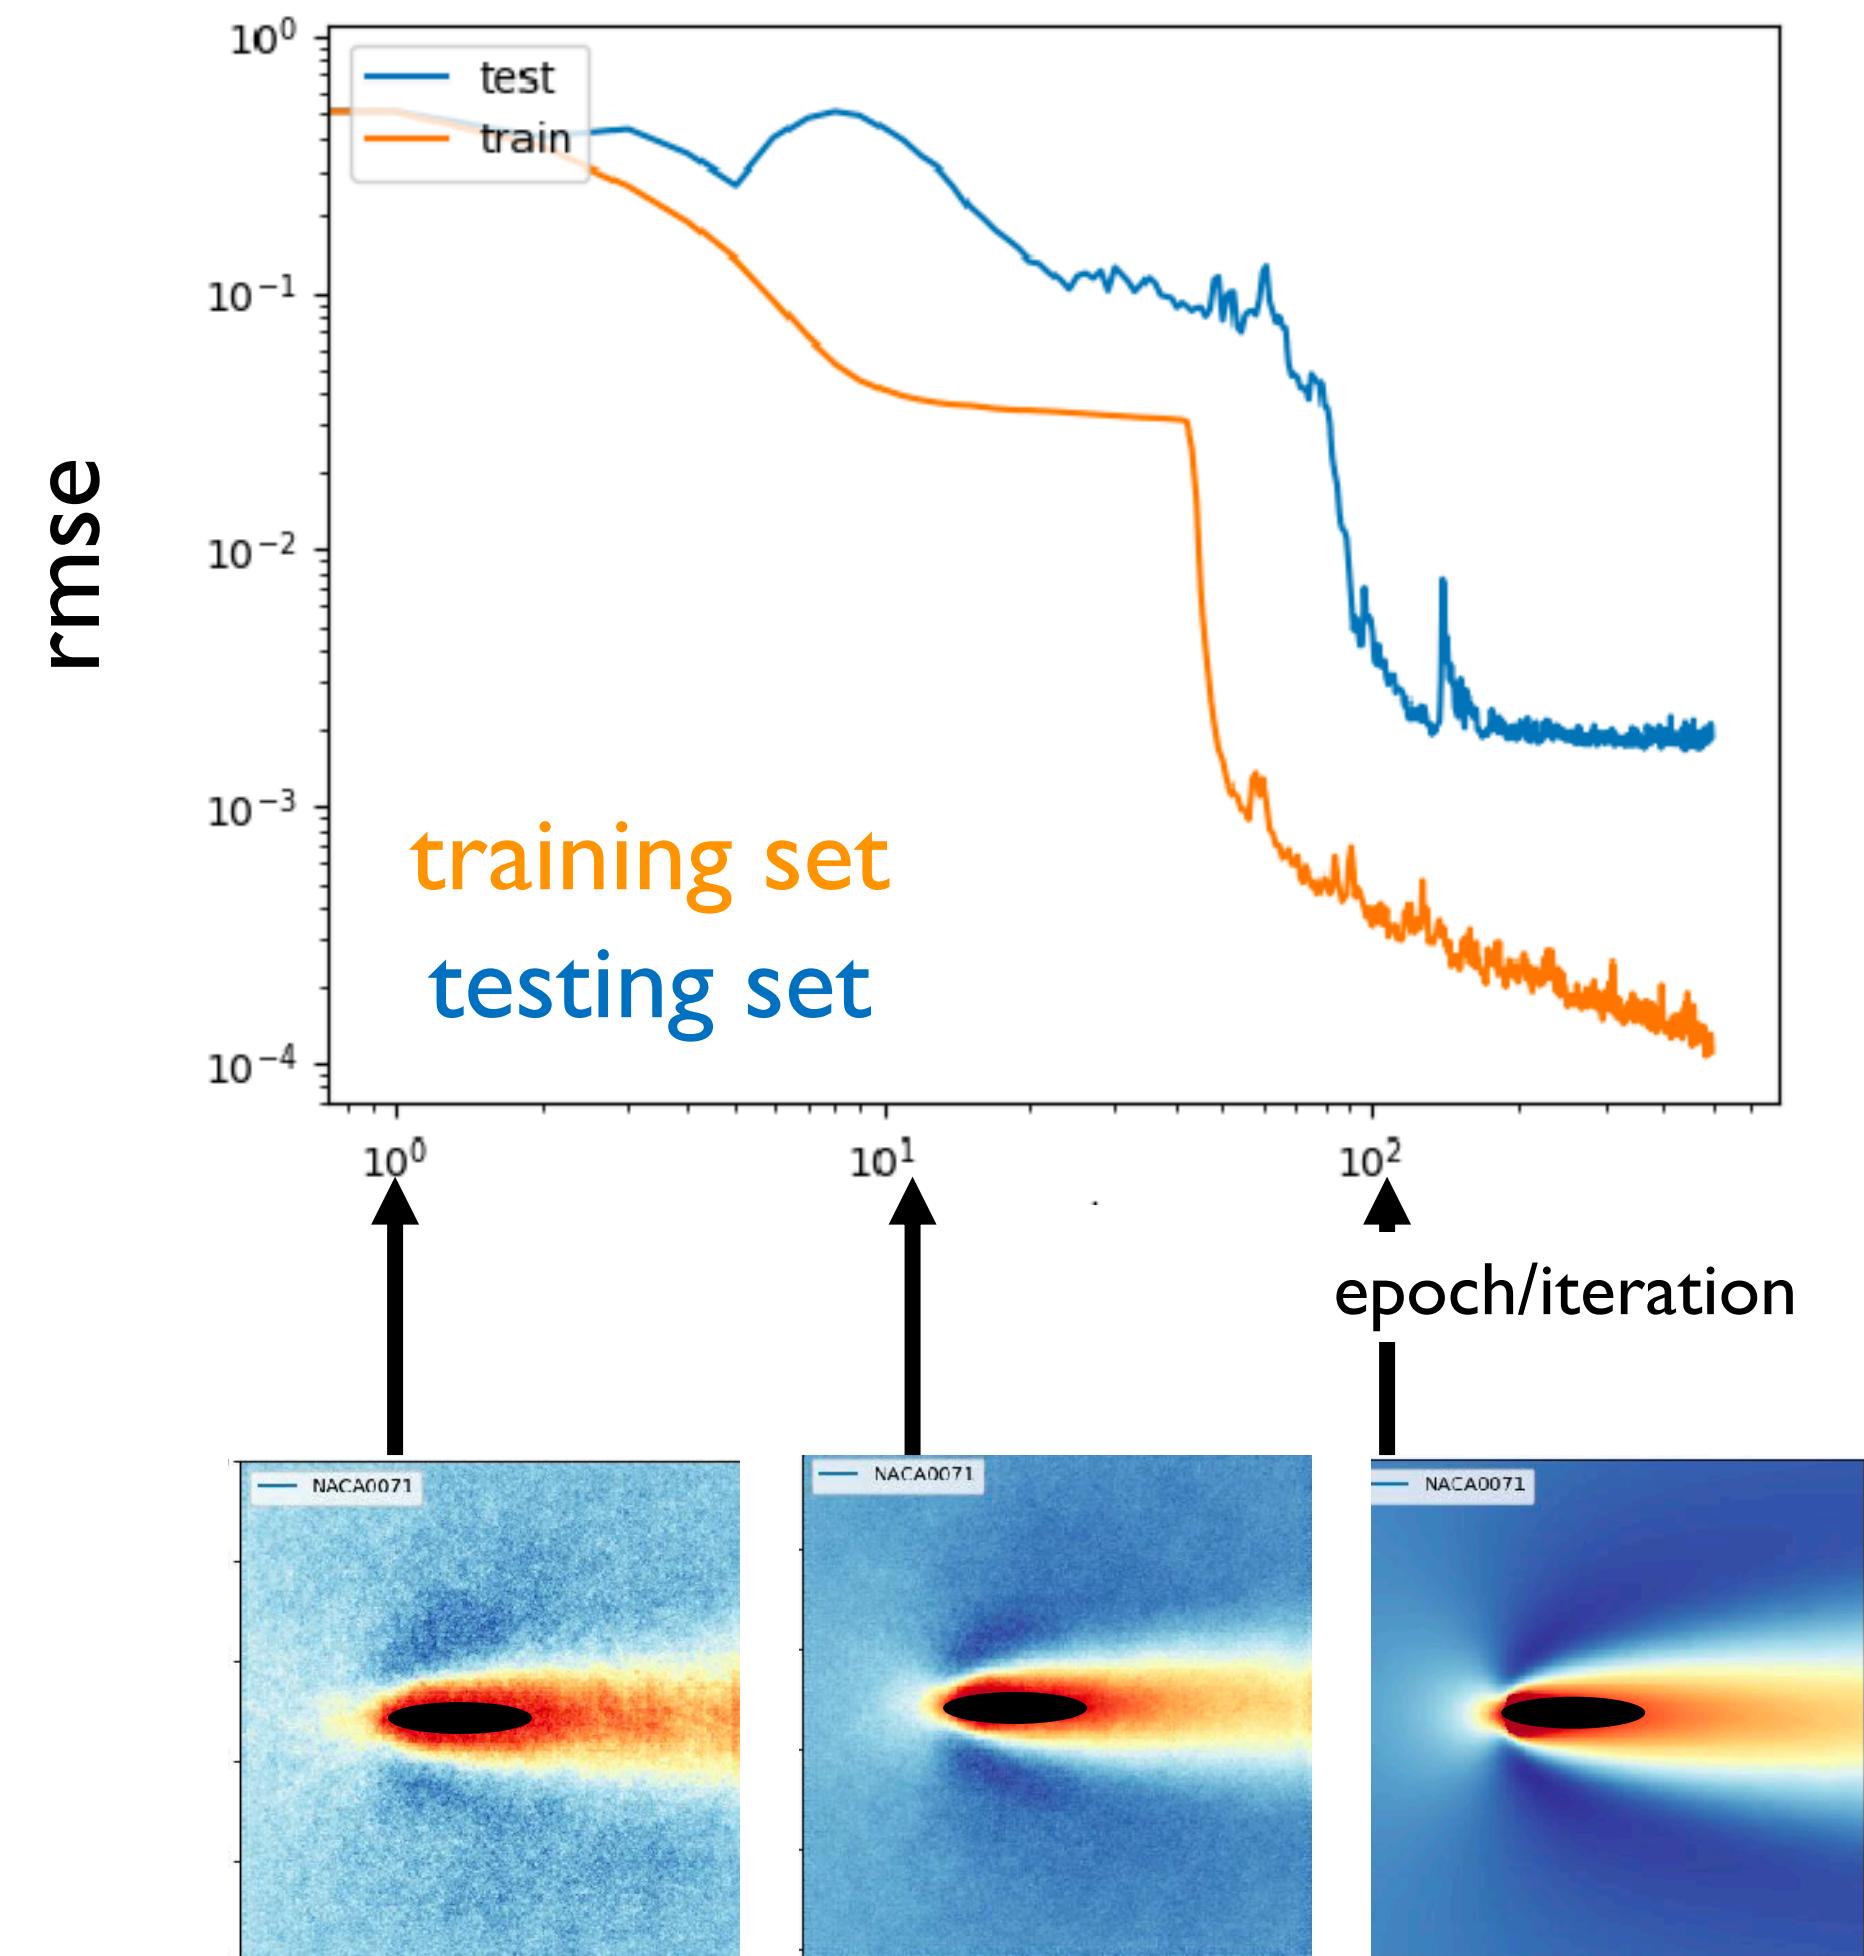


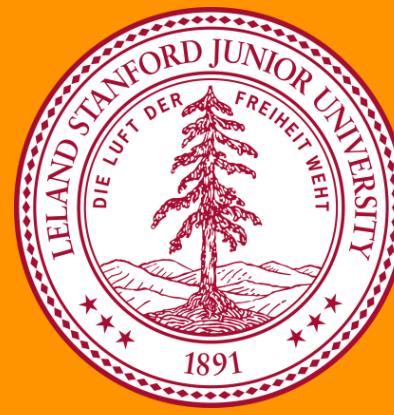
Pacing Items

99% of time to construct the dataset and train the network



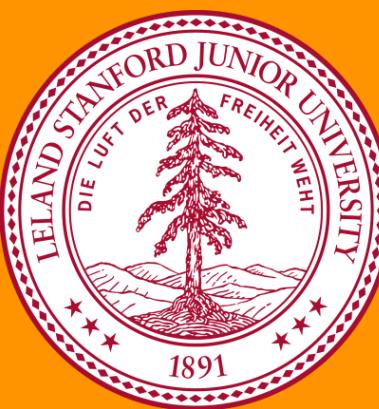
The choice of network architecture has strong influence on the success metric





Outline

1. Blind Experiment >> ML-driven Predictions
2. Physics First >> ML-driven Modeling
3. Reset >> ML-driven Numerical Methods



Turbulence Modeling in the Age of Data

Annual Review of Fluid Mechanics

Vol. 51:357-377 (Volume publication date January 2019)

First published as a Review in Advance on September 19, 2018

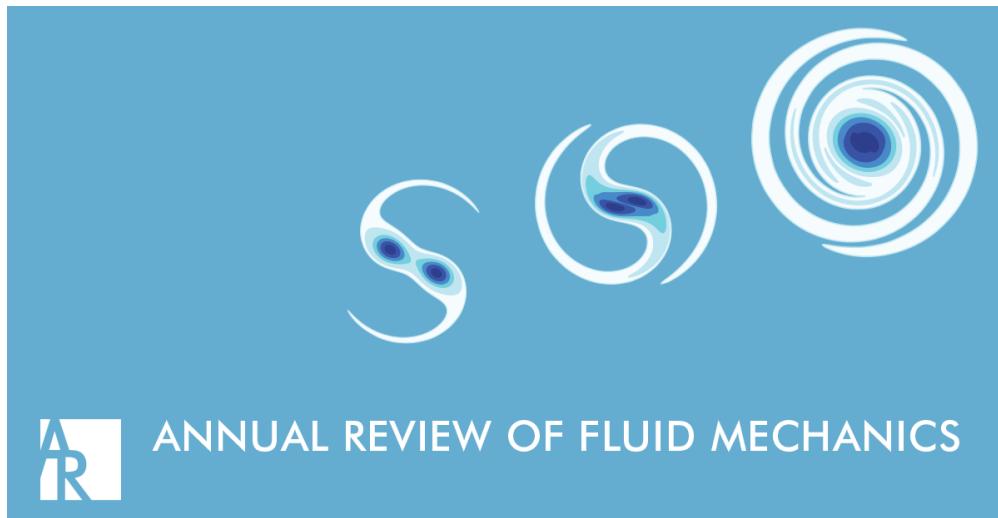
<https://doi.org/10.1146/annurev-fluid-010518-040547>

Karthik Duraisamy,^{1,*} Gianluca Iaccarino,^{2,*} and Heng Xiao^{3,*}

¹Department of Aerospace Engineering, University of Michigan, Ann Arbor, Michigan 48109, USA; email: kdur@umich.edu

²Department of Mechanical Engineering, Stanford University, Stanford, California 94305, USA; email: jops@stanford.edu

³Kevin T. Crofton Department of Aerospace and Ocean Engineering, Virginia Tech, Blacksburg, Virginia 24060, USA; email: hengxiao@vt.edu



[Full Text HTML](#)

[Download PDF](#)

[Article Metrics](#)

[Permissions](#) | [Reprints](#) | [Download Citation](#) | [Citation Alerts](#)



Random Forests and Turbulence Models

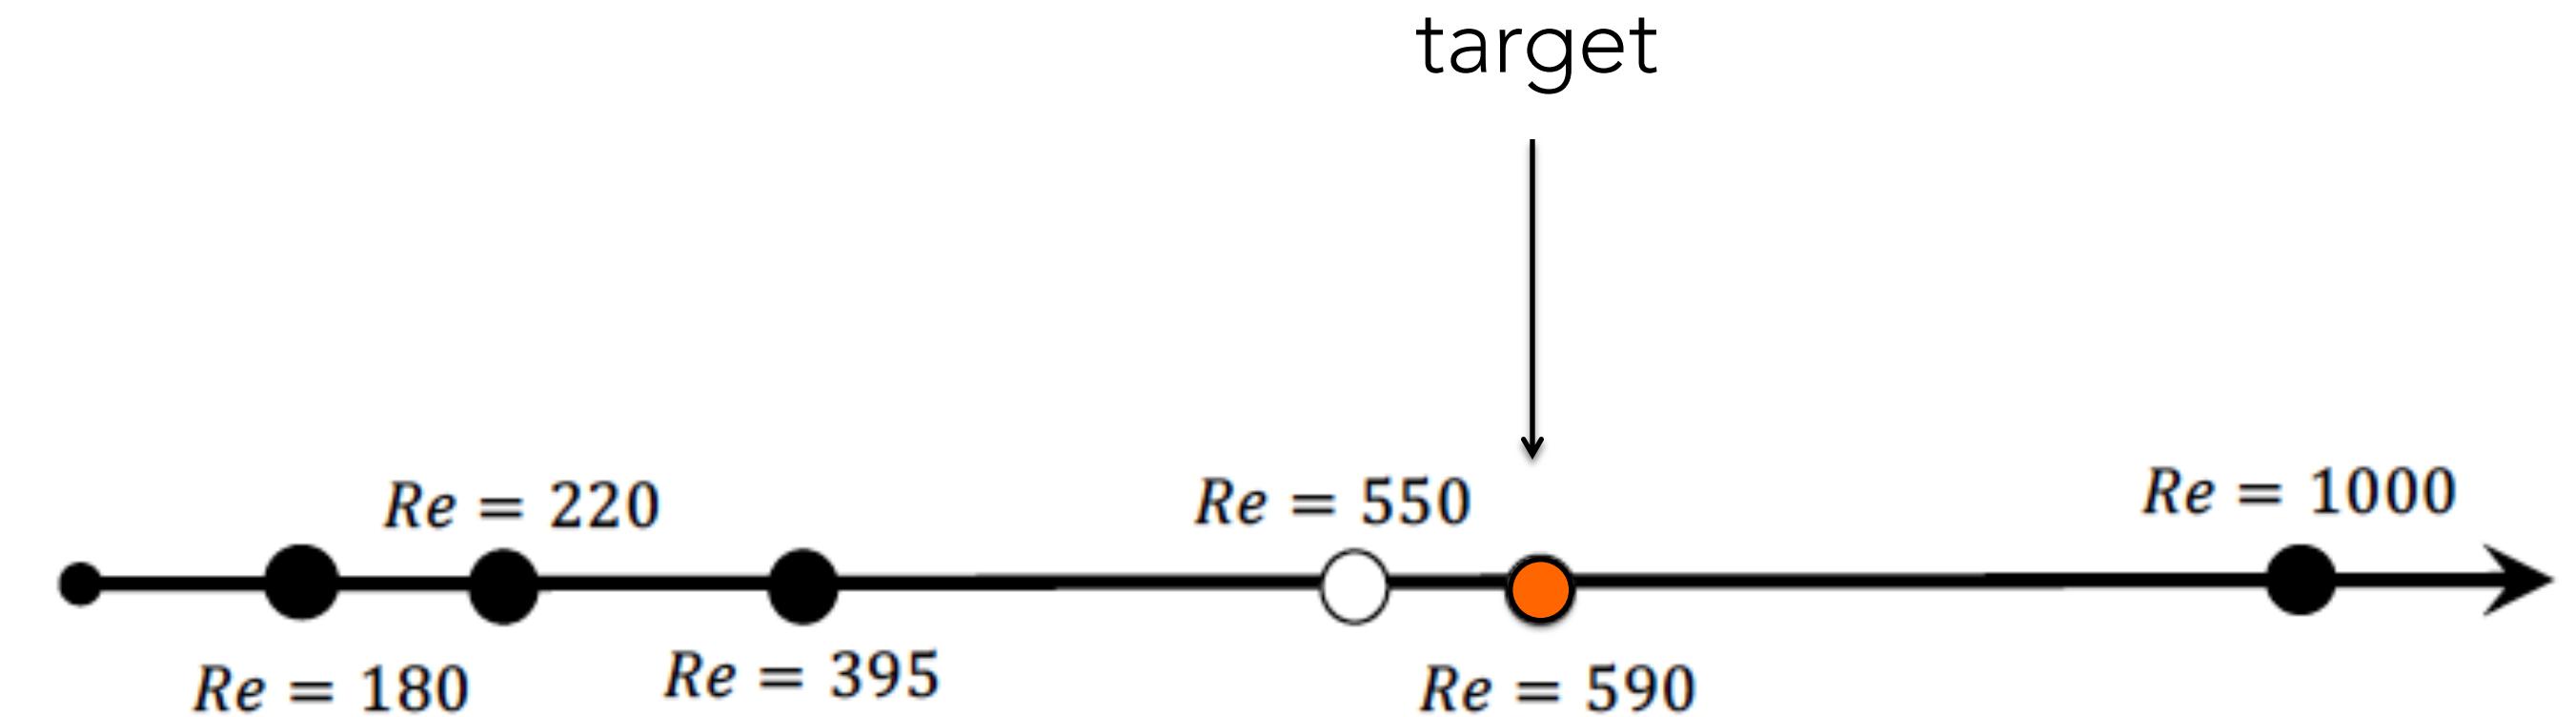
- **Goal:** learn a correction for the turbulence kinetic energy computed by the $k-\varepsilon$ model @ $Re_\tau = 590$

$$k_{\text{DNS}} = k_{k-\varepsilon} + \delta k(\eta)$$

- **Data:** DNS of channel flows from J. Jimenez et al.

- **Decisions:** dictionary of **28 known turbulence features**

- Turbulent Reynolds number (based on k)
- Time scale ratio
- Wall distance
- Ratio of rotation/strain
- Ratio of production and dissipation
- [...]



- **Forest:** collection of **25** decision trees
 - Max Depth: 10
 - Max Features: 6 (with repetition)
 - Decision split based on variance reduction
 - Training dataset: $Re_\tau = 180, 220, 395, 1000$
 - Test dataset: $Re_\tau = 590$

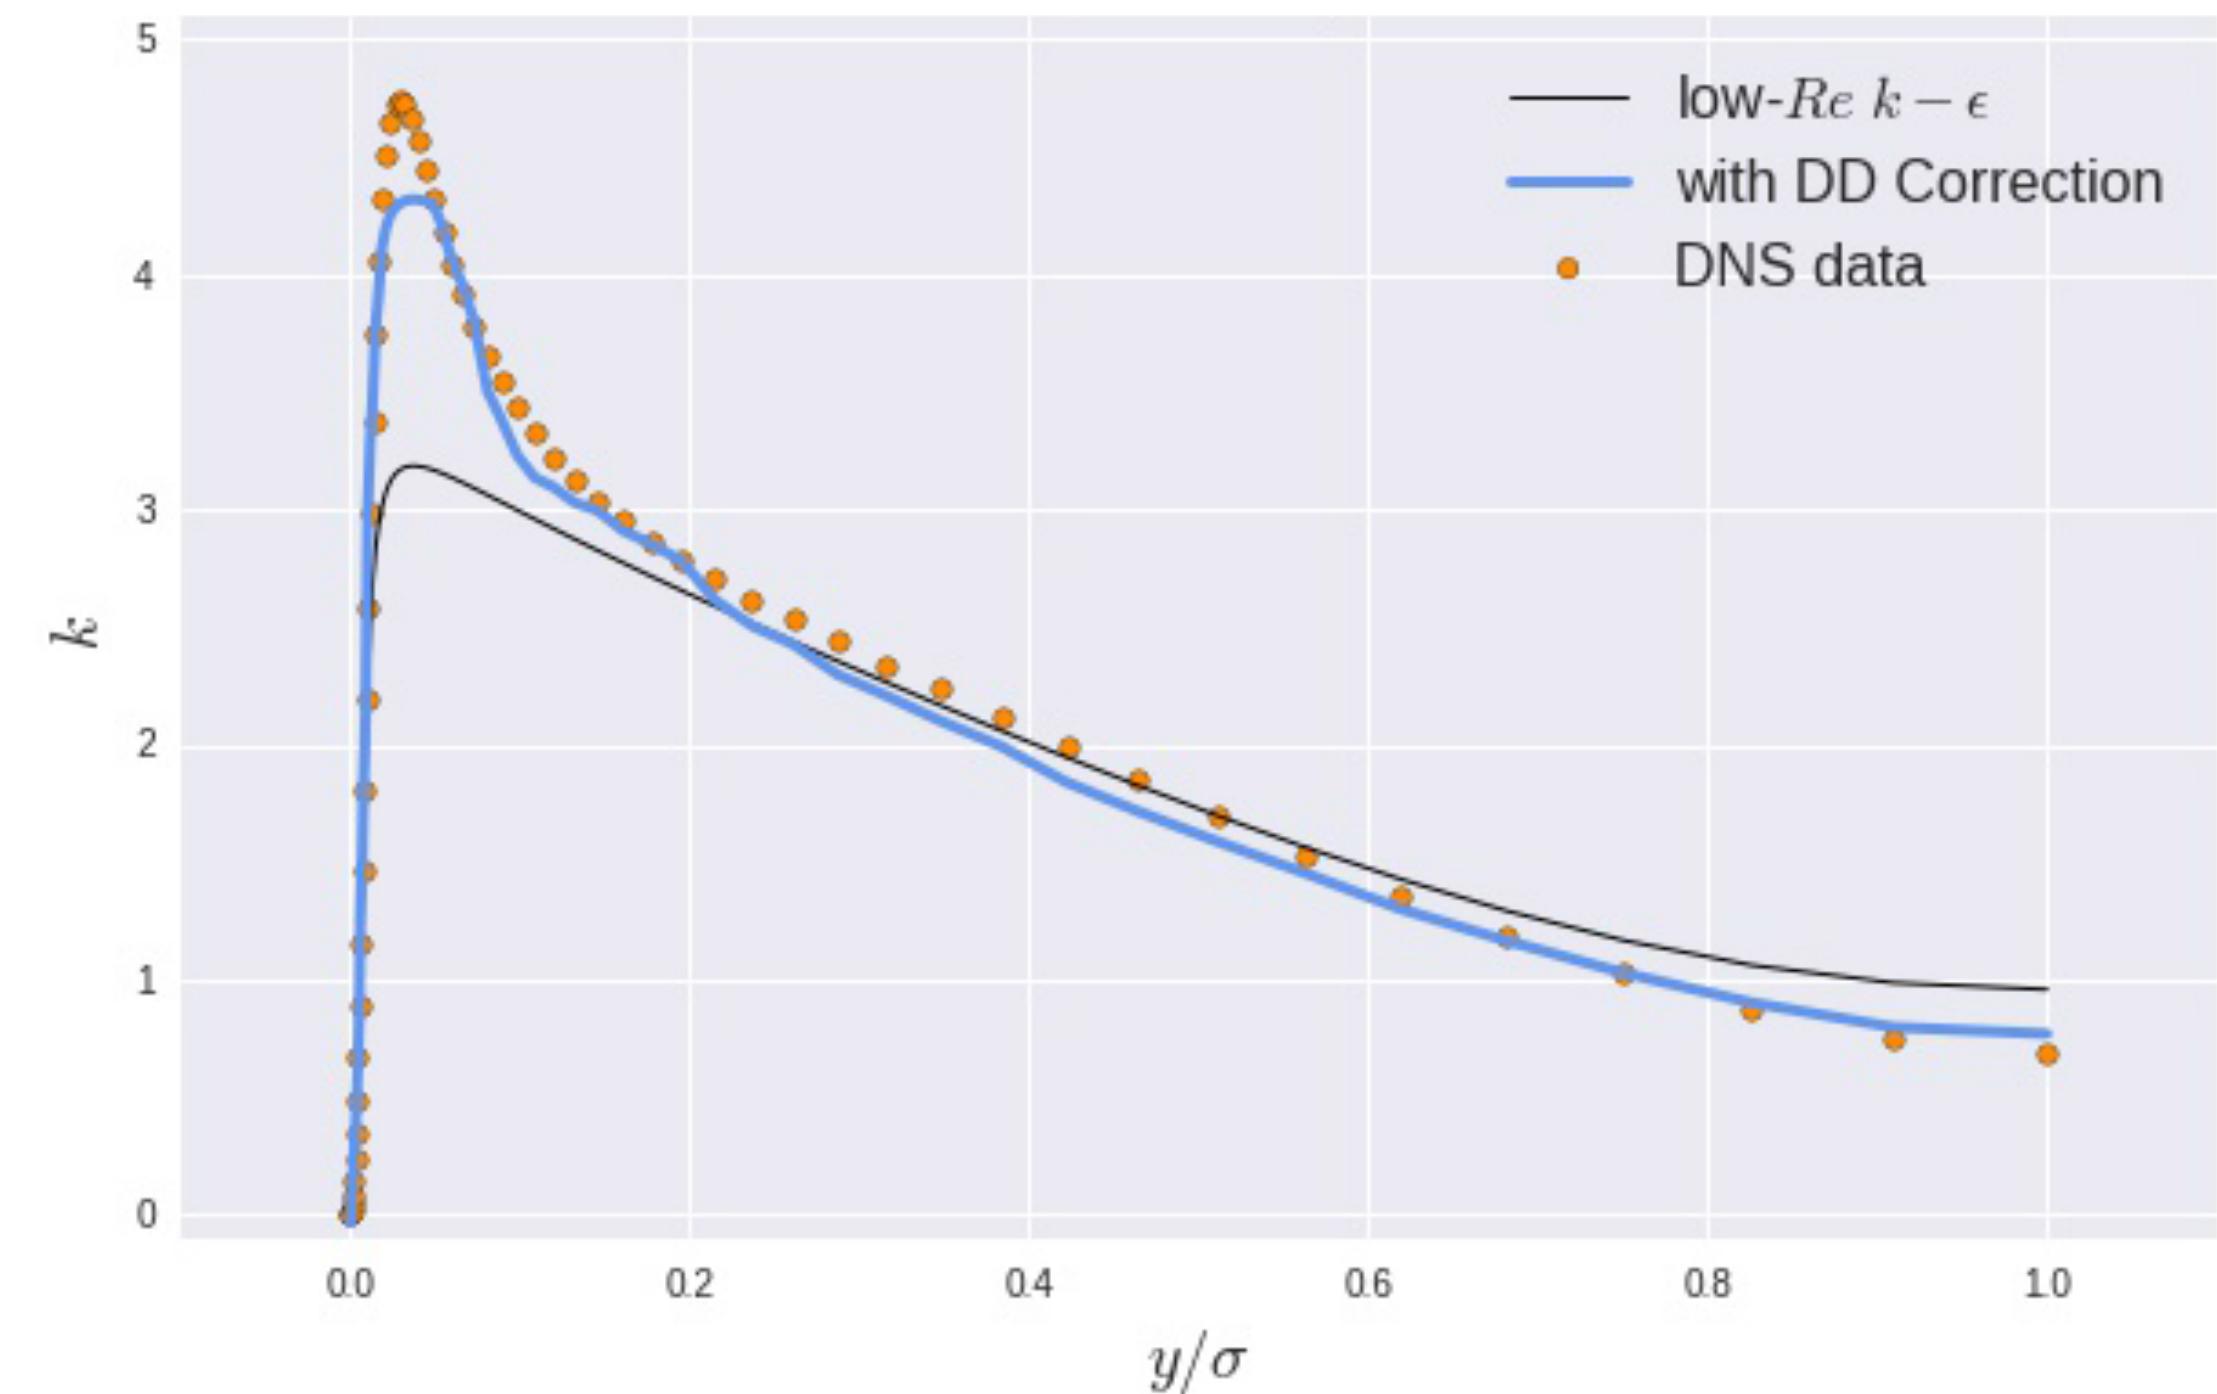
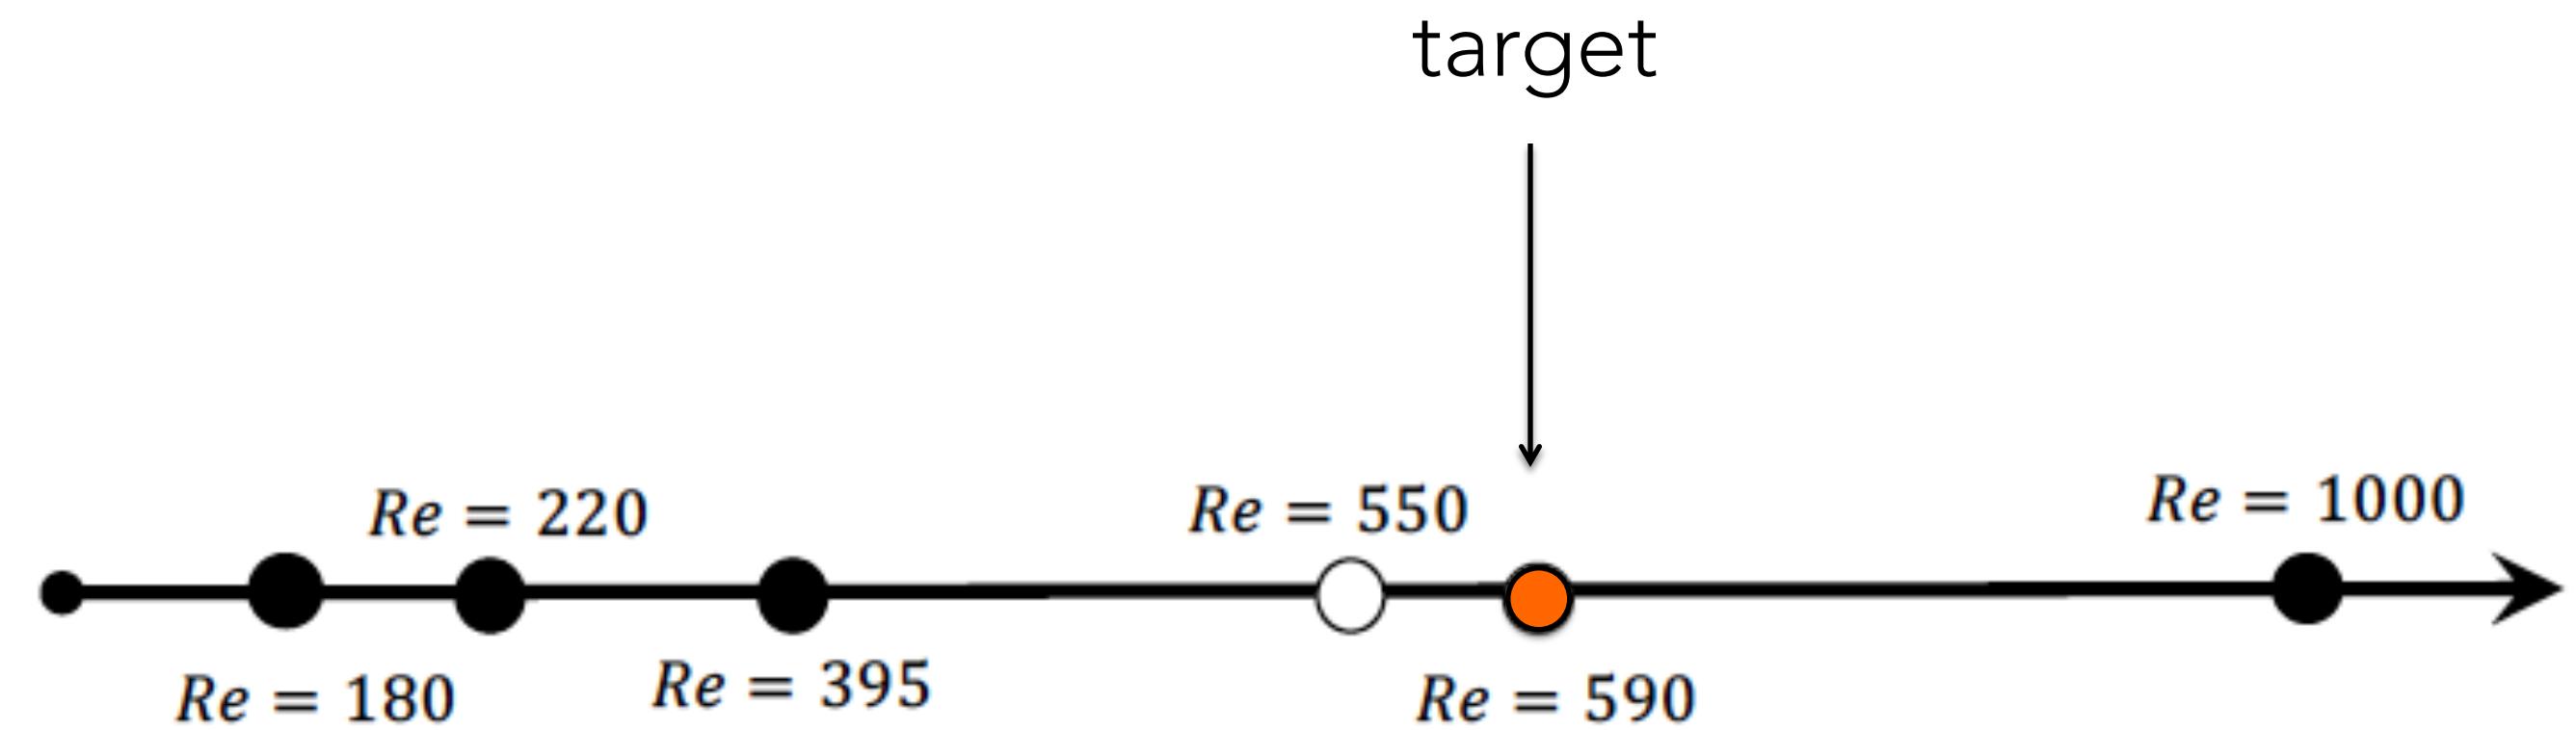


Random Forests and Turbulence Models

- **Goal:** learn a correction for the turbulence kinetic energy computed by the $k-\varepsilon$ model @ $Re_\tau = 590$

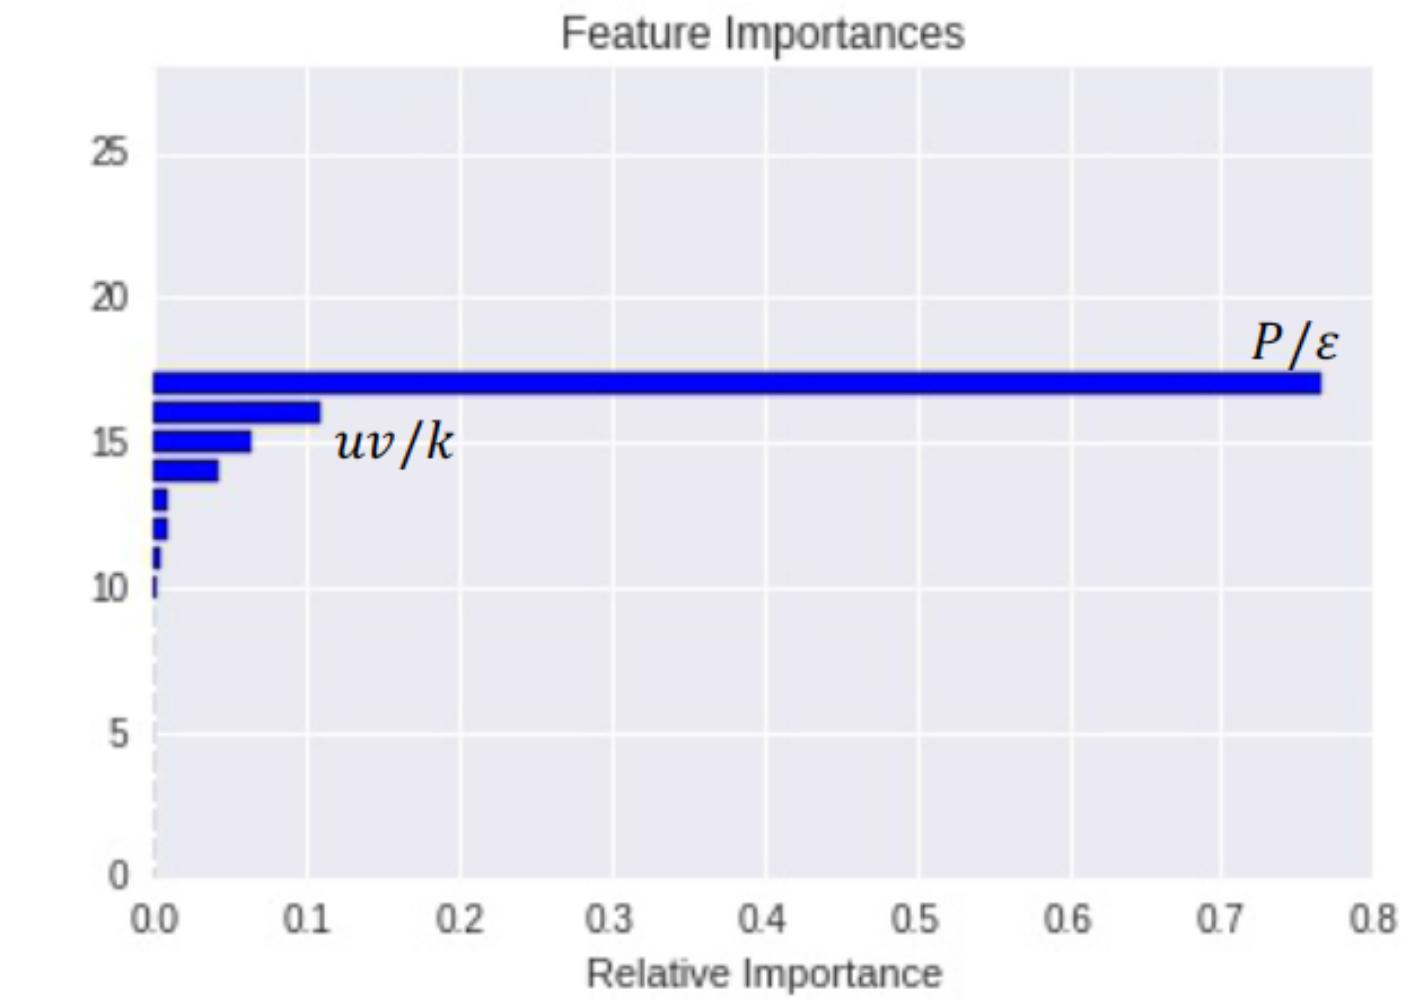
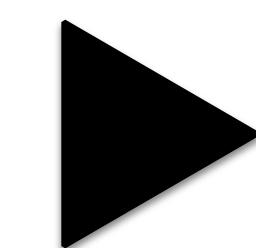
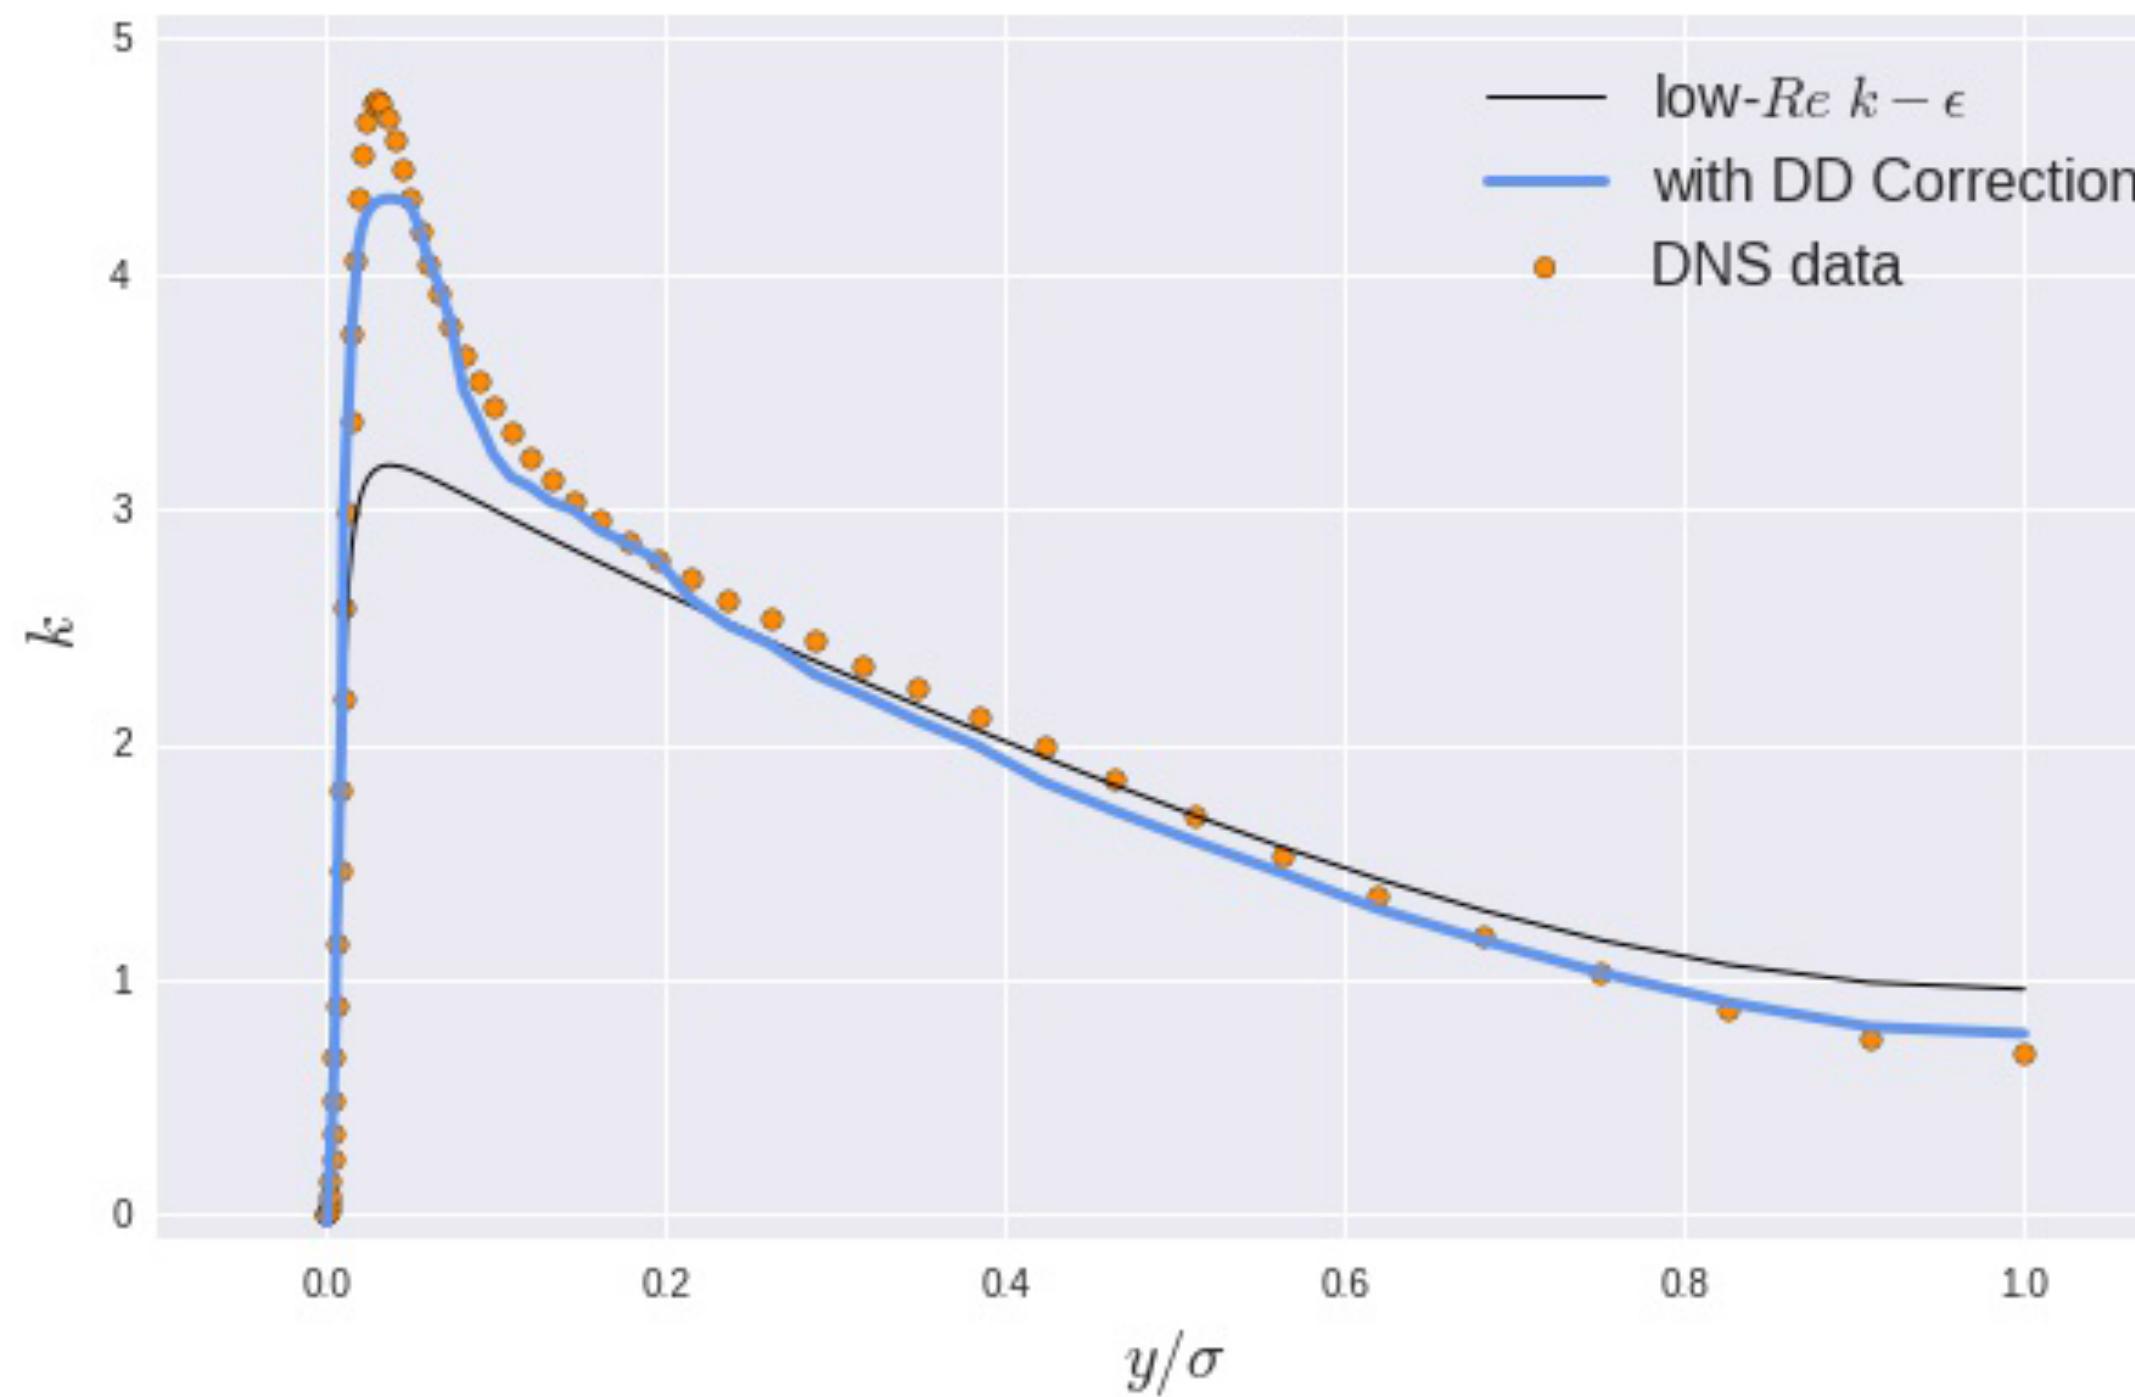
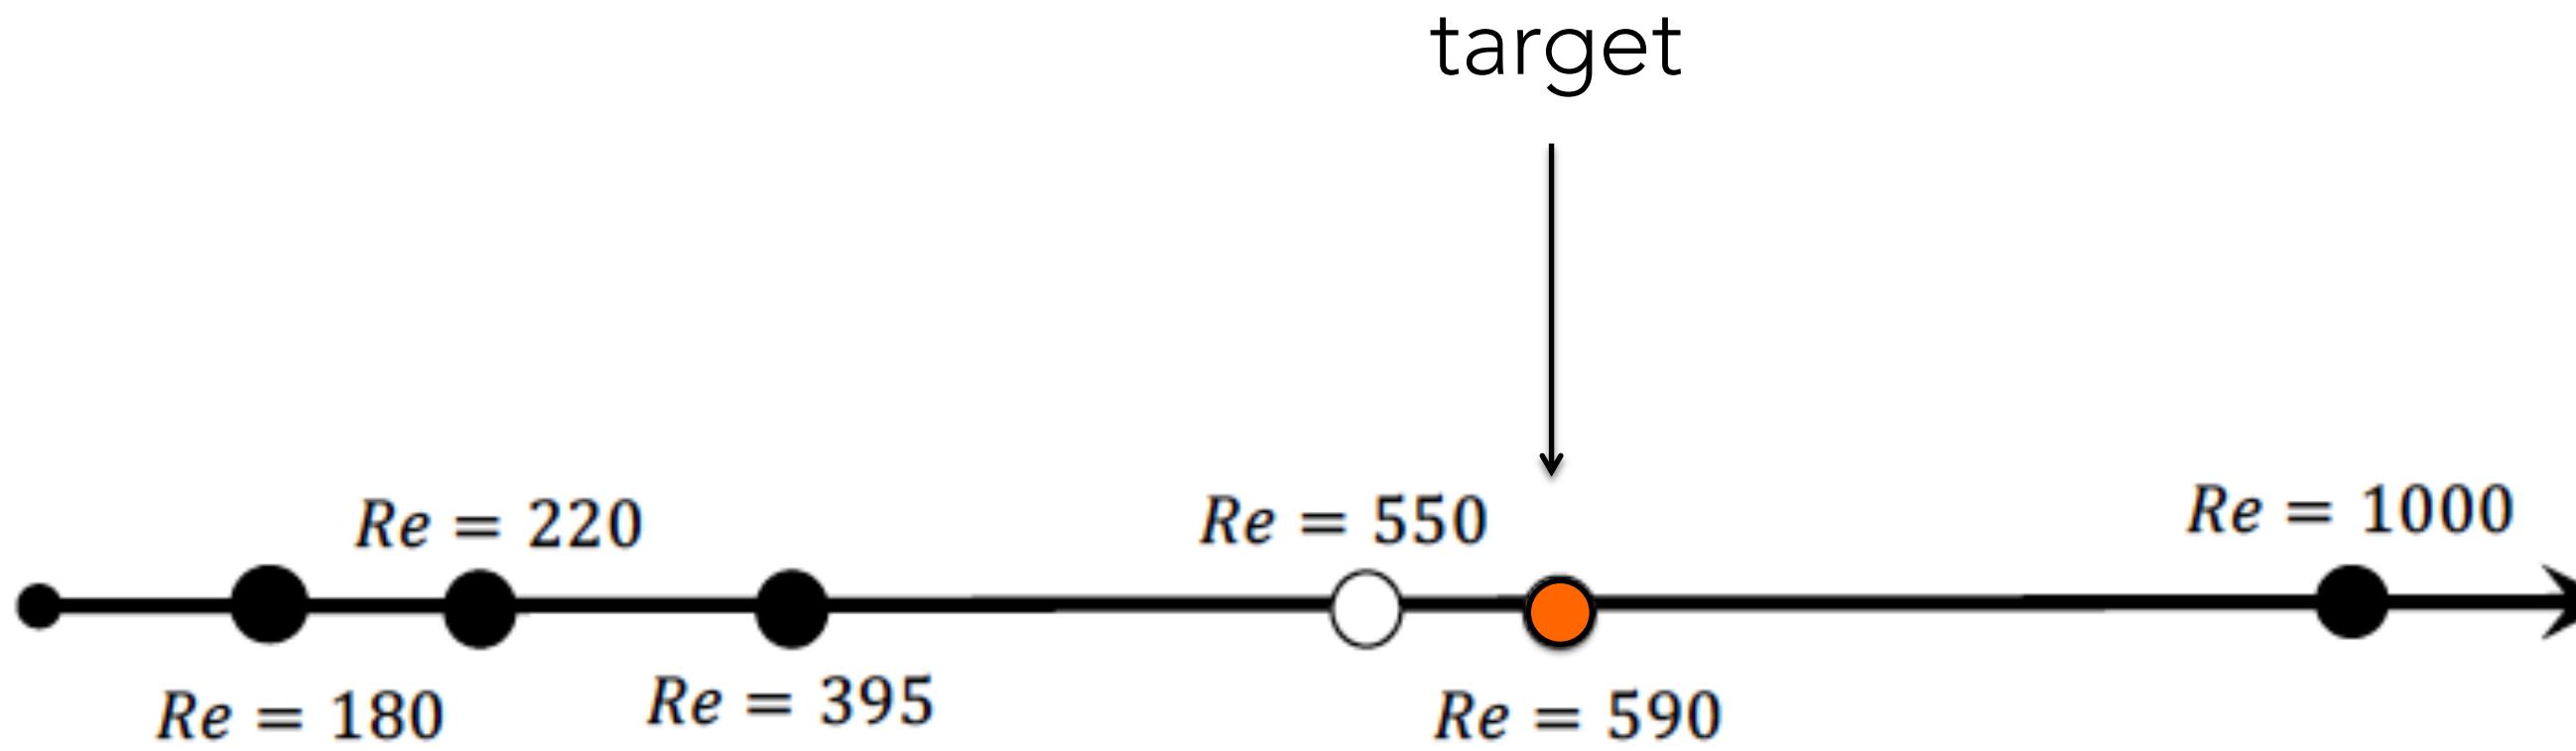
$$k_{\text{DNS}} = k_{k-\varepsilon} + \delta k(\eta)$$

- **Data:** DNS of channel flows from J. Jimenez et al.

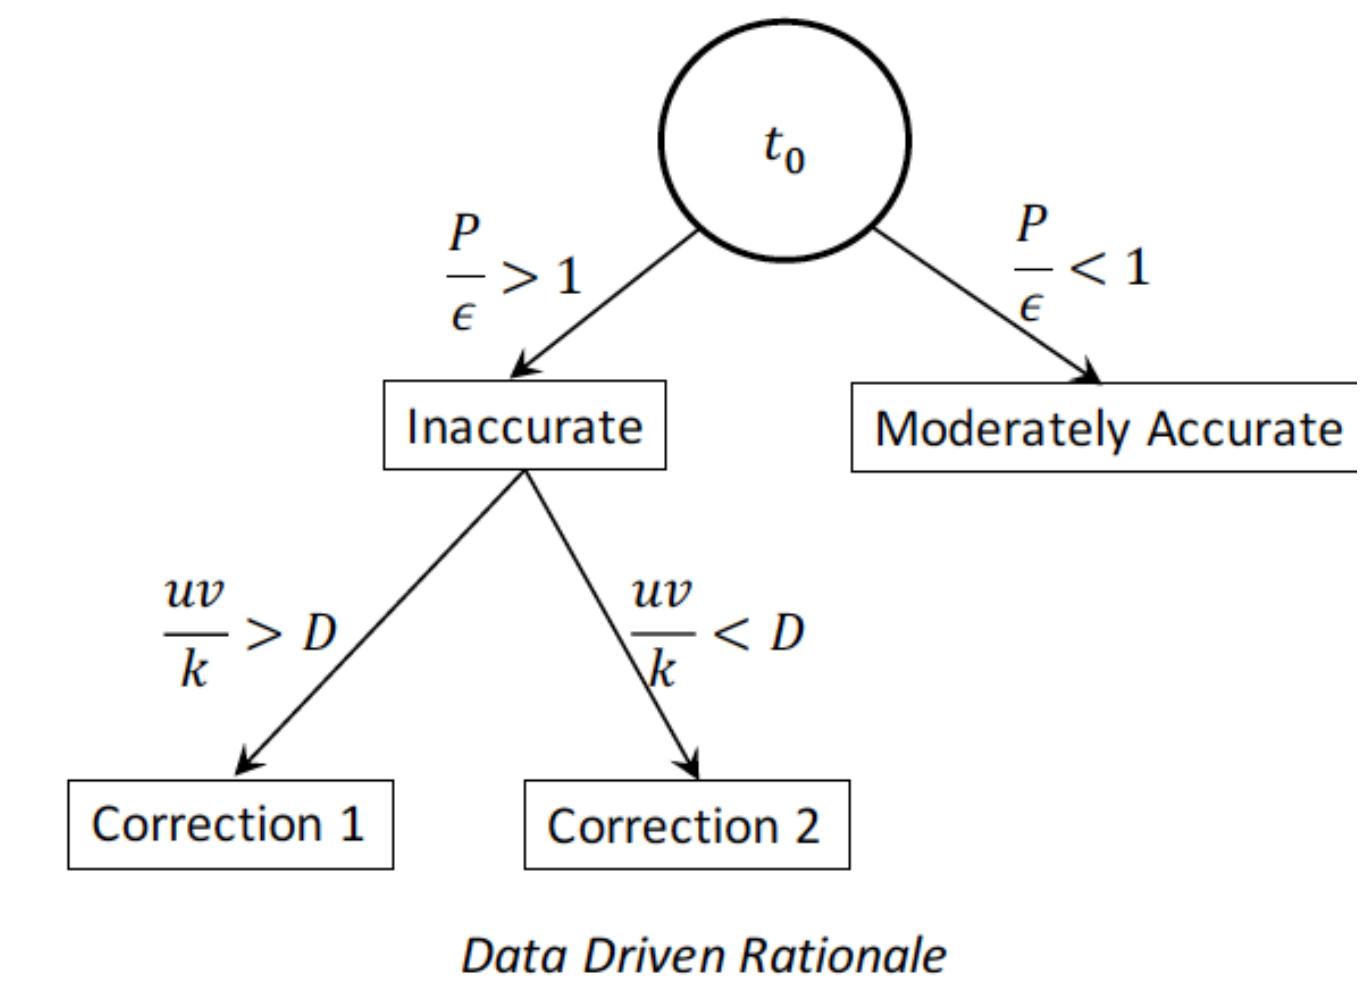


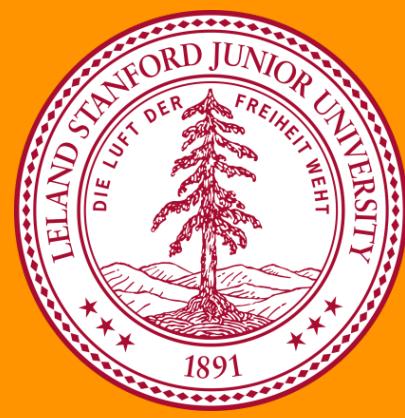


Random Forests and Turbulence Models



Data-driven rationale & insight

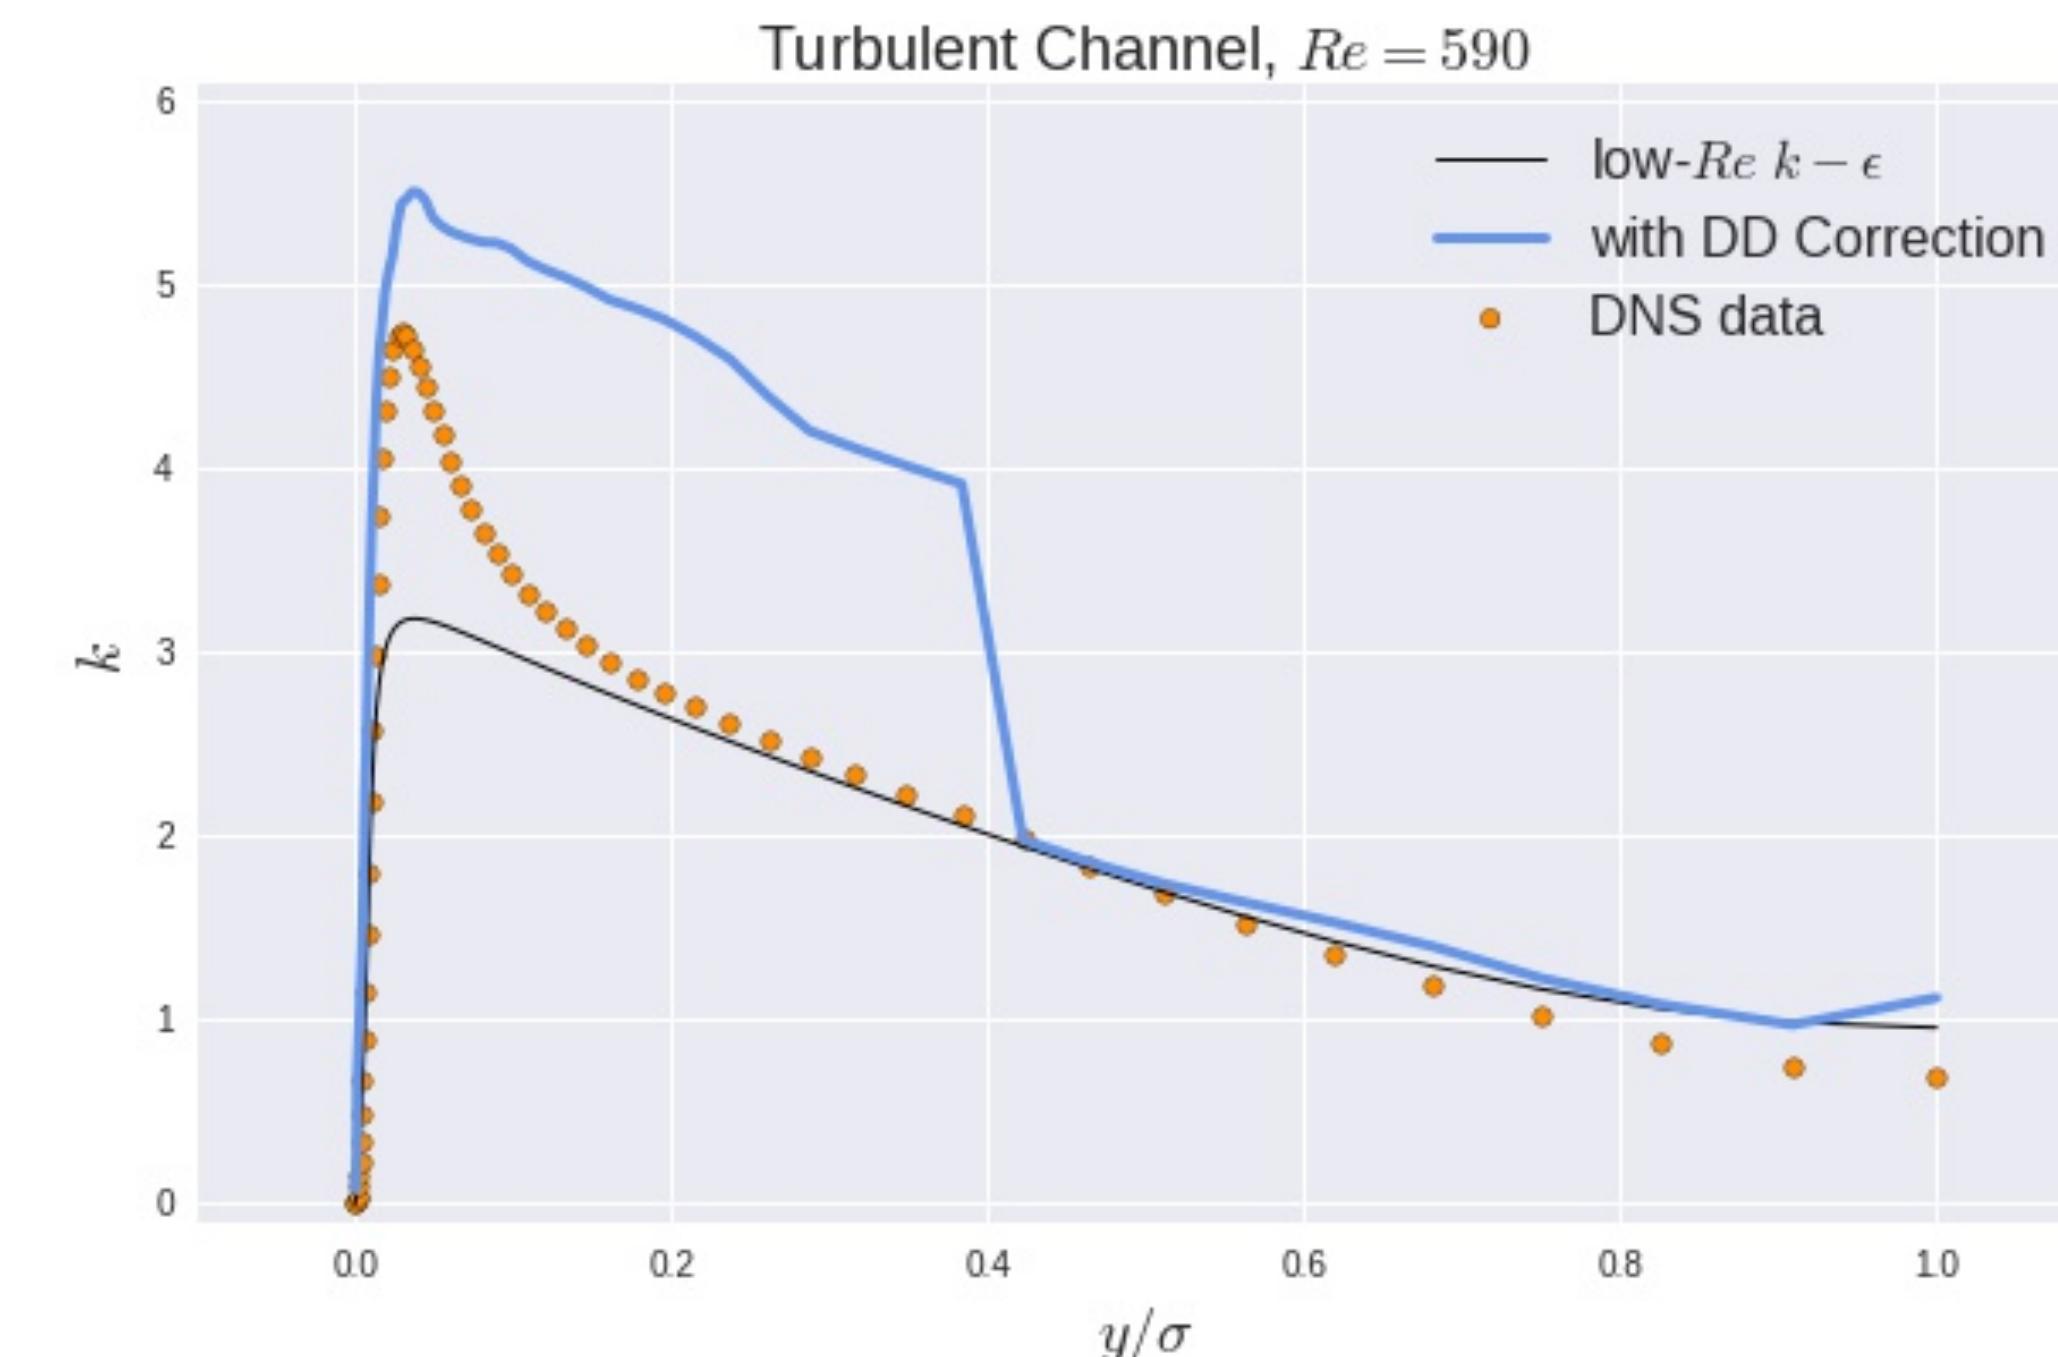
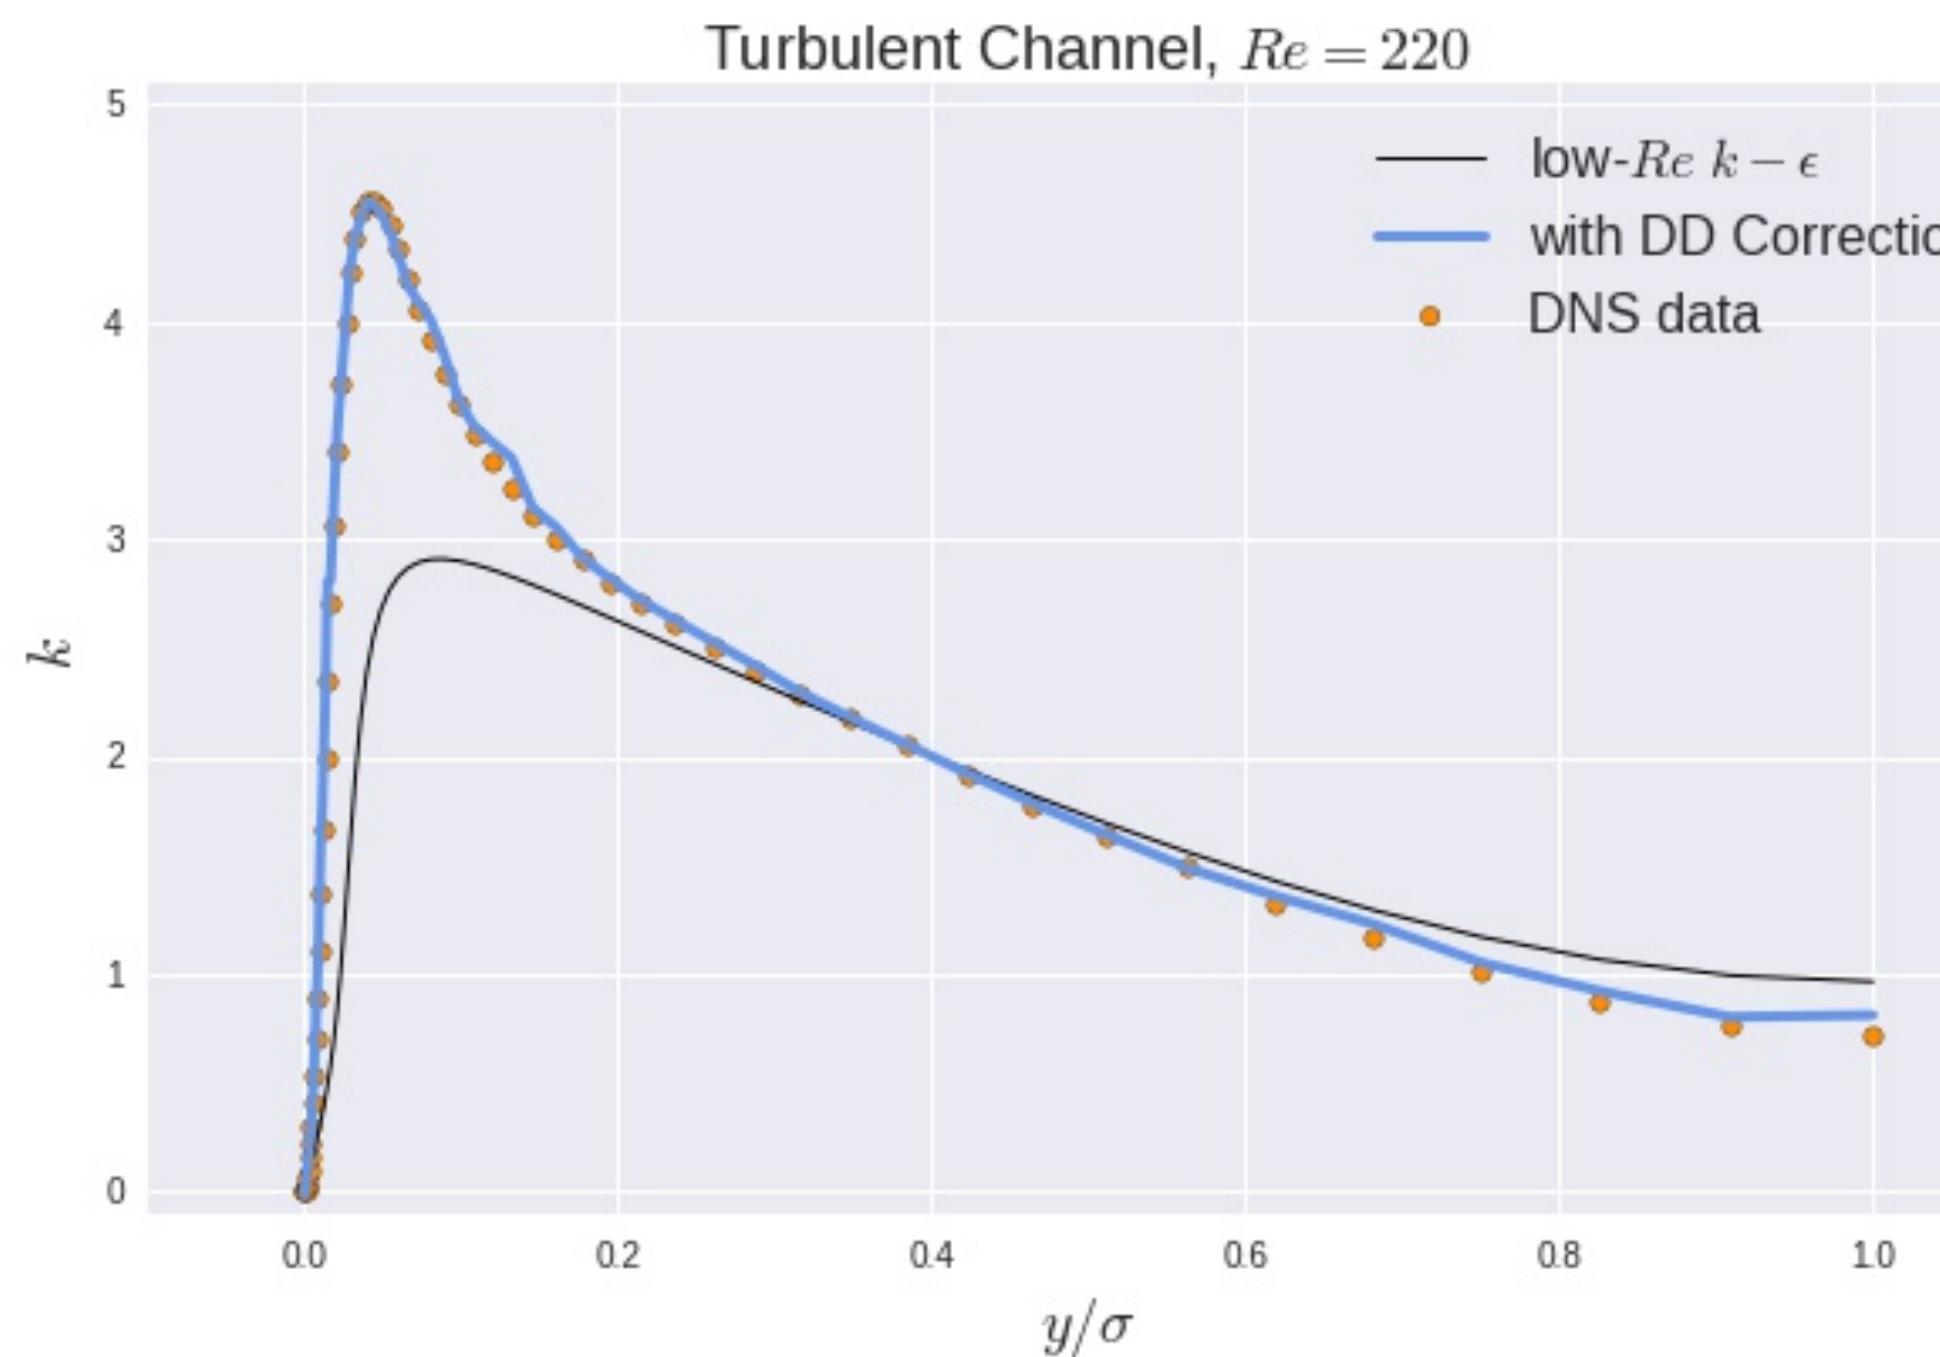
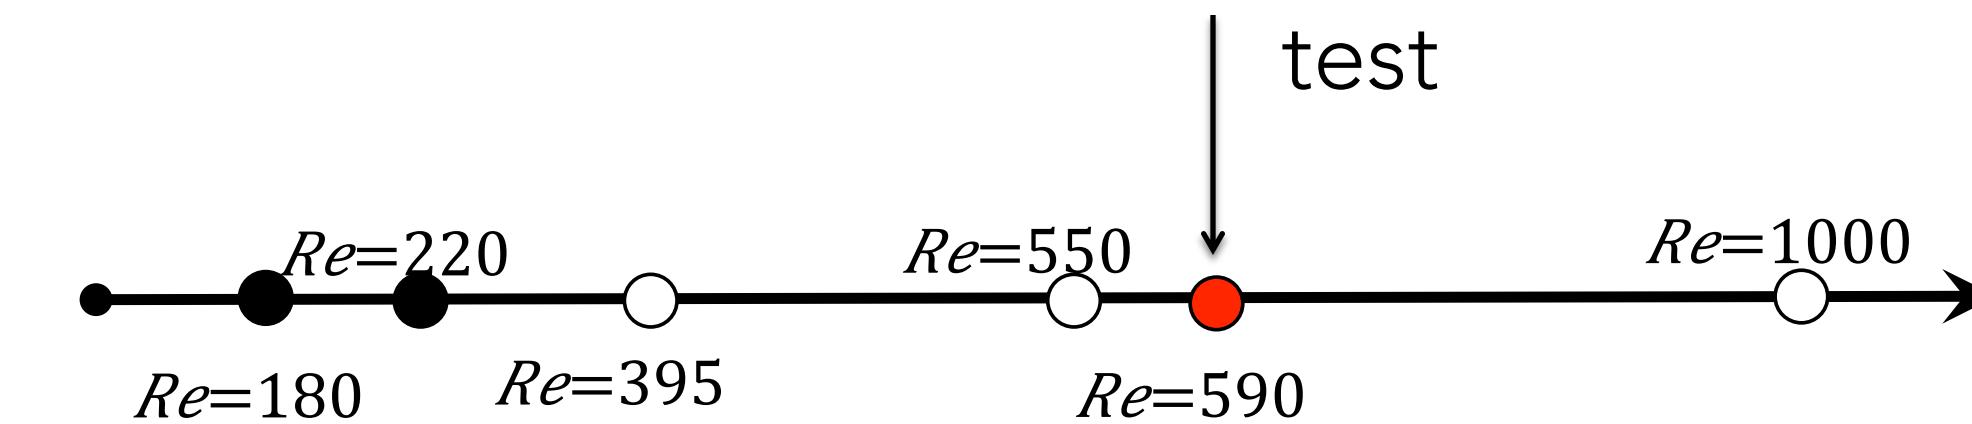
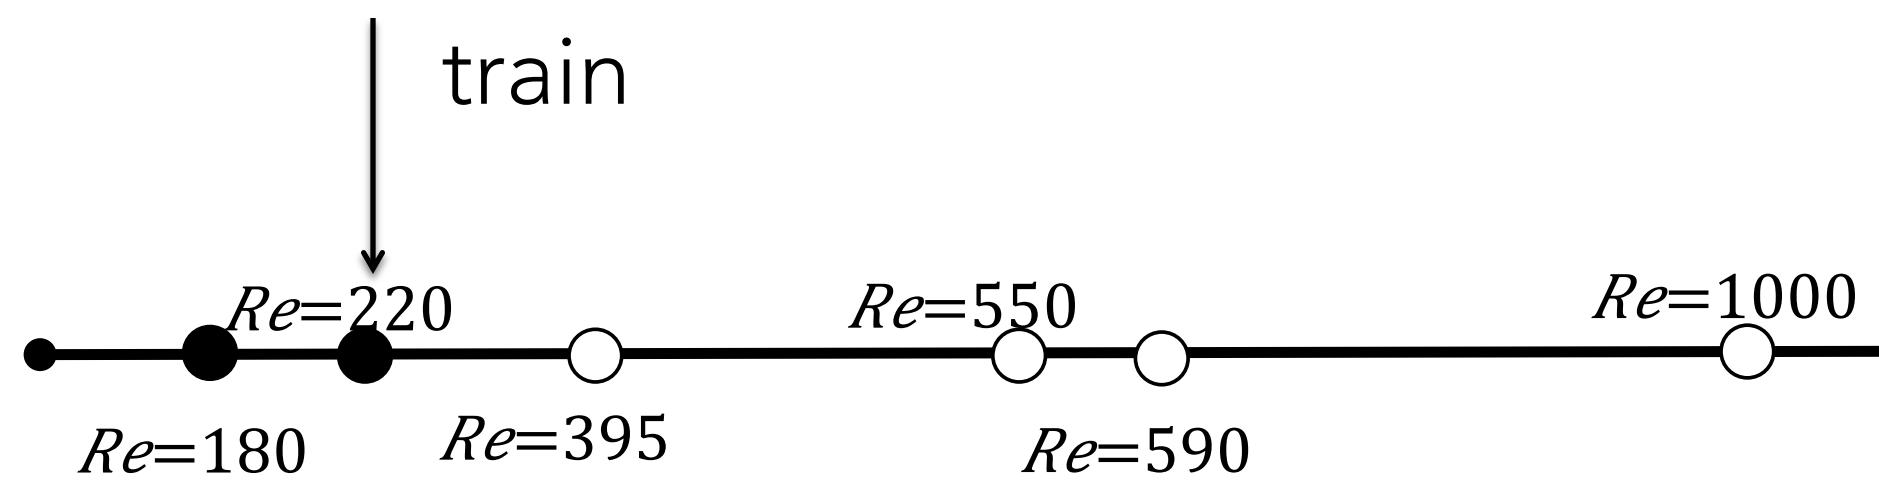


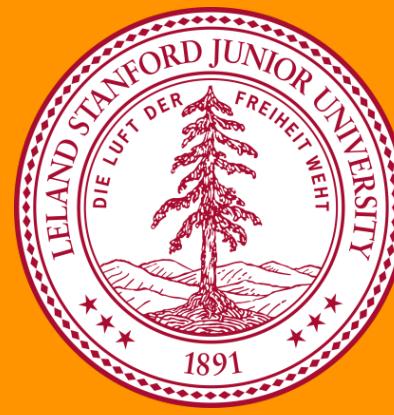


Random Forests and Turbulence Models

Insufficient size of the dataset: **overfitting**

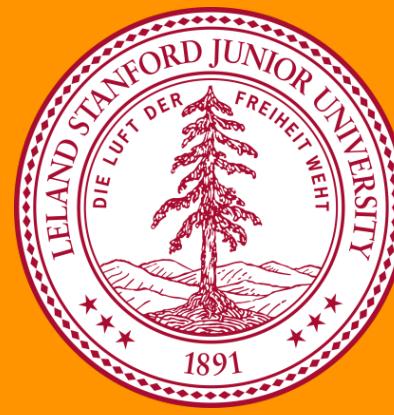
Inappropriate choice of dataset range: **extrapolation**





Outline

1. Blind Experiment >> ML-driven Predictions
2. Physics First >> ML-driven Modeling
3. Reset >> ML-driven Numerical Methods



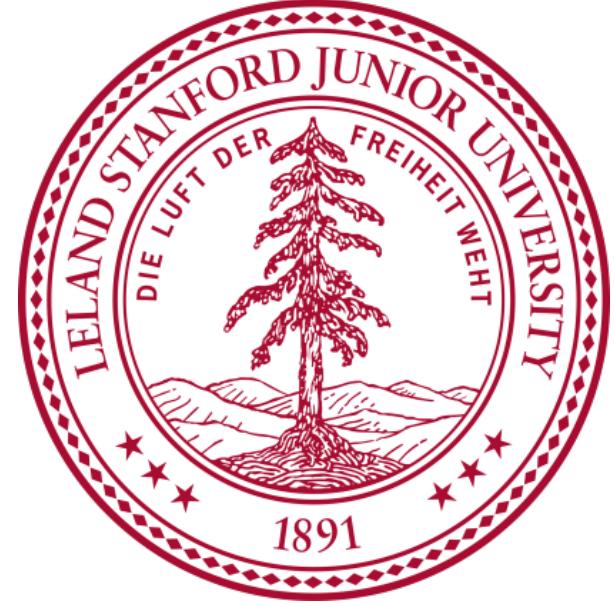
Pacing Items

Which one is the right ML (data-driven) strategy?

What (and how much) data is really necessary?

How can we systematically improve the prediction of ML algorithms?

How to construct network architectures that embed prior knowledge (e.g. physical principles) rather than enforcing them as constraints in the optimization problem?



(Machine) Learning to Differentiate

Oskar Alund, Jan Nordstrom

Department of Mathematics
Linköping University

oskar.alund@liu.se jan.norstrom@liu.se

Gianluca Iaccarino

Mechanical Engineering Department &
Institute for Computational Mathematical Engineering
Stanford University
jops@stanford.edu

**NSF Workshop on Exuberance of Machine
Learning in Transport Phenomena**
February 10 — 11, 2020
Dallas, TX



(Machine) Learning to Differentiate

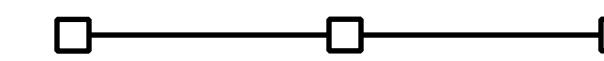
Can we build an ML-based strategy to construct a numerical method for solving
Partial Differential Equations?



Differentiate Polynomials

Consider a 3-points mesh

$$\Omega_h = \{-1, 0, 1\}$$



and polynomials evaluated on the mesh points

$$\mathbf{1} = [1 \quad 1 \quad 1]^\top, \quad \mathbf{x} = [-1 \quad 0 \quad 1]^\top, \quad \mathbf{x}^2 = [1 \quad 0 \quad 1]^\top$$

We can write the differentiation process as a matrix operation

$$D\mathbf{1} = 0 \quad D\mathbf{x} = 1 \quad D\mathbf{x}^2 = 2\mathbf{x}$$

Leading to:

$$D \begin{bmatrix} 1 & -1 & 1 \\ 1 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0 & 1 & -2 \\ 0 & 1 & 0 \\ 0 & 1 & 2 \end{bmatrix} \implies D = \begin{bmatrix} -1.5 & 2 & -0.5 \\ -0.5 & 0 & 0.5 \\ 0.5 & -2 & 1.5 \end{bmatrix}$$

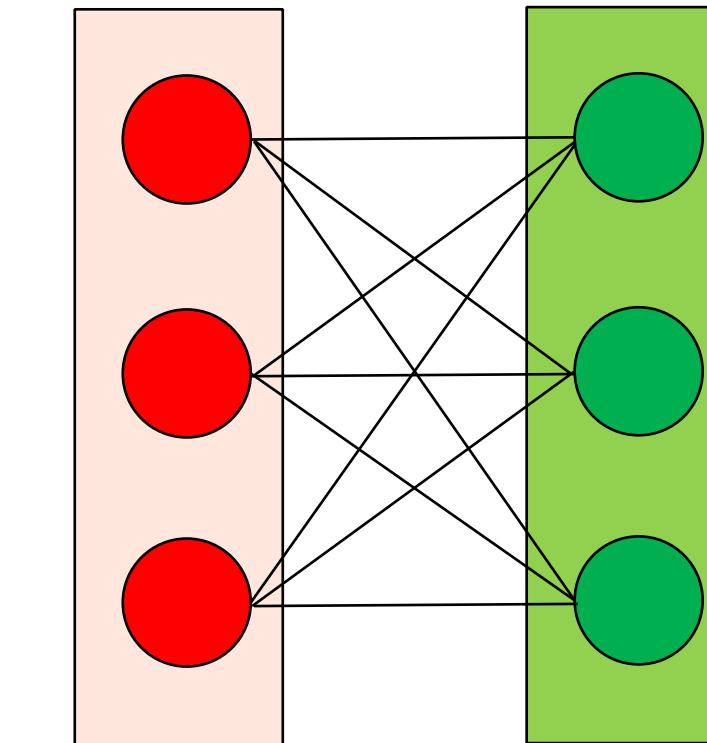


ML to Differentiate Polynomials

The simplest network that can accomplish the same task is:

\mathbf{p}_i

function
values at
nodes



\mathbf{p}'_i

derivative
at nodes

with weights $\mathbf{W} = (w_{ij}) \in \mathbb{R}^{3 \times 3}$

obtained minimizing

$$\sum_{i=0}^2 \|\mathbf{p}'_i - \mathbf{W}\mathbf{p}_i\|_2^2$$

If we assume linear activation functions and no-bias $\mathbf{W} \implies \mathbf{D}$



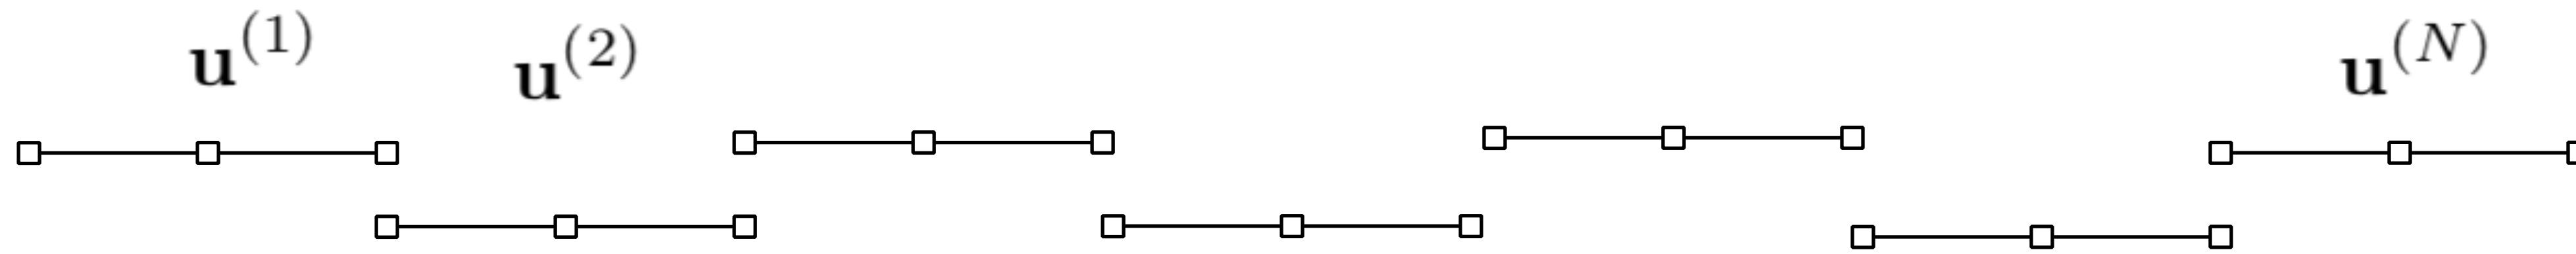
A Numerical Technique for PDEs



Consider linear transport

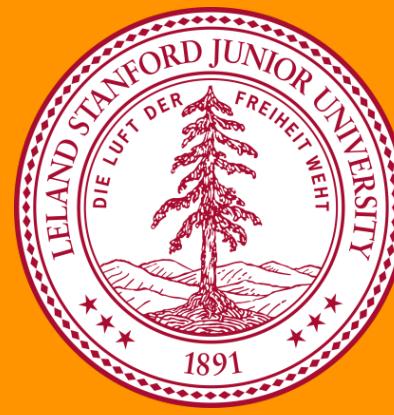
$$u_t + u_x = 0, \quad x \in (0, 1), \quad t > 0,$$

and a mesh with 3-point stencil in each mesh element (i.e. a nodal DG scheme)



$$\begin{aligned}\mathbf{u}_t^{(1)} + \mathbf{D}^{(1)}\mathbf{u}^{(1)} &= \sigma_1(u_0^{(1)} - g)\mathbf{e}_0 \\ \mathbf{u}_t^{(2)} + \mathbf{D}^{(2)}\mathbf{u}^{(2)} &= \sigma_2(u_0^{(2)} - u_2^{(1)})\mathbf{e}_0 \\ &\vdots \\ \mathbf{u}_t^{(N)} + \mathbf{D}^{(N)}\mathbf{u}^{(N)} &= \sigma_N(u_0^{(N)} - u_2^{(N-1)})\mathbf{e}_0\end{aligned}$$

the RHS corresponds to coupling terms between elements and the boundary conditions with coefficient chosen to achieve **stability**



A Stable Numerical Technique for PDEs



if $D = P^{-1}Q$

with $\mathbf{1}^\top P \mathbf{x}^k = \int_{-1}^1 x^k dx$ and $Q + Q^\top = \text{diag}(-1, 0, 1)$

the discretization satisfies the summation-by-part property (SBP) and the scheme is provably **stable!**

H. O. Kreiss and J. Oliger, Comparison of accurate methods for the integration of hyperbolic equations, Tellus, 24 (1972), pp. 199-215.

We can train **2 networks**, one for P and one for Q and build D combining weights



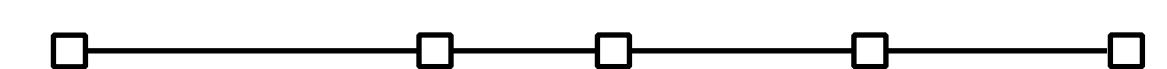
A Stable Numerical Technique for PDEs



We can train **2 networks**, one for **P** and one for **Q** and build **D** combining weights

Remarks:

- We can use **any function to train** the networks for **P** and **Q** as long as we can compute integrals (for **P**) and derivatives (for **Q**)
- Enforcing $\mathbf{Q} + \mathbf{Q}^\top = \text{diag}(-1, 0, 1)$ is not straightforward we can use a **constraint on the loss function** or **design a network** that automatically enforces a **constraint on the weights**
- We can use any point-distribution within each mesh element and achieve high-order accuracy, e.g. $\Omega_h = \{-1, -0.3, 0, 0.5, 1\}$





A Stable Numerical Technique for PDEs



CO learning_to_diff.ipynb ☆

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text Reconnect ▾

>

▼ Learning to Differentiate

In this notebook we show the basic procedure of constructing SBP operators using Tensorflow/Keras.

Let us first import a few necessary libraries.

```
[ ] import numpy as np
import tensorflow as tf
import numpy.polynomial.polynomial as poly
import matplotlib.pyplot as plt
tf.keras.backend.set_floatx('float64')
np.set_printoptions(precision=4, suppress=True)

num_copies=20000
noise_variance=0.0
```

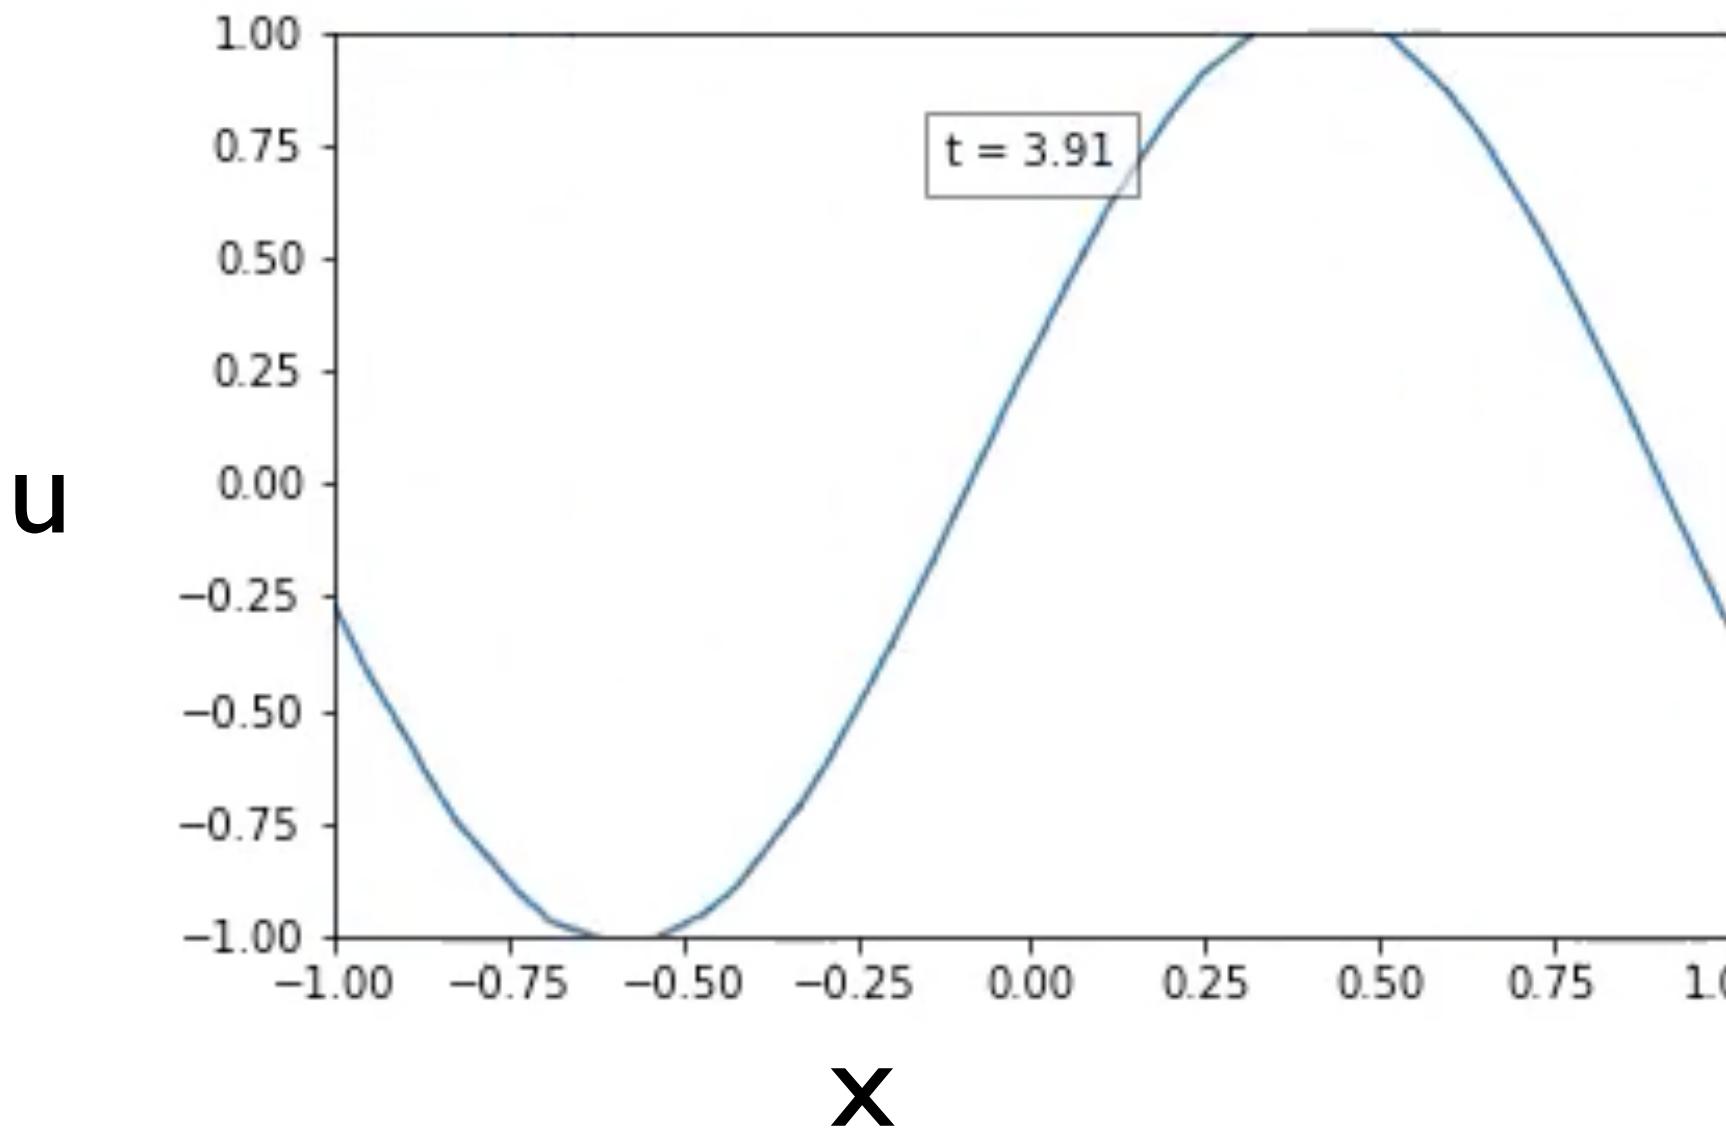
▼ Training the quadrature.



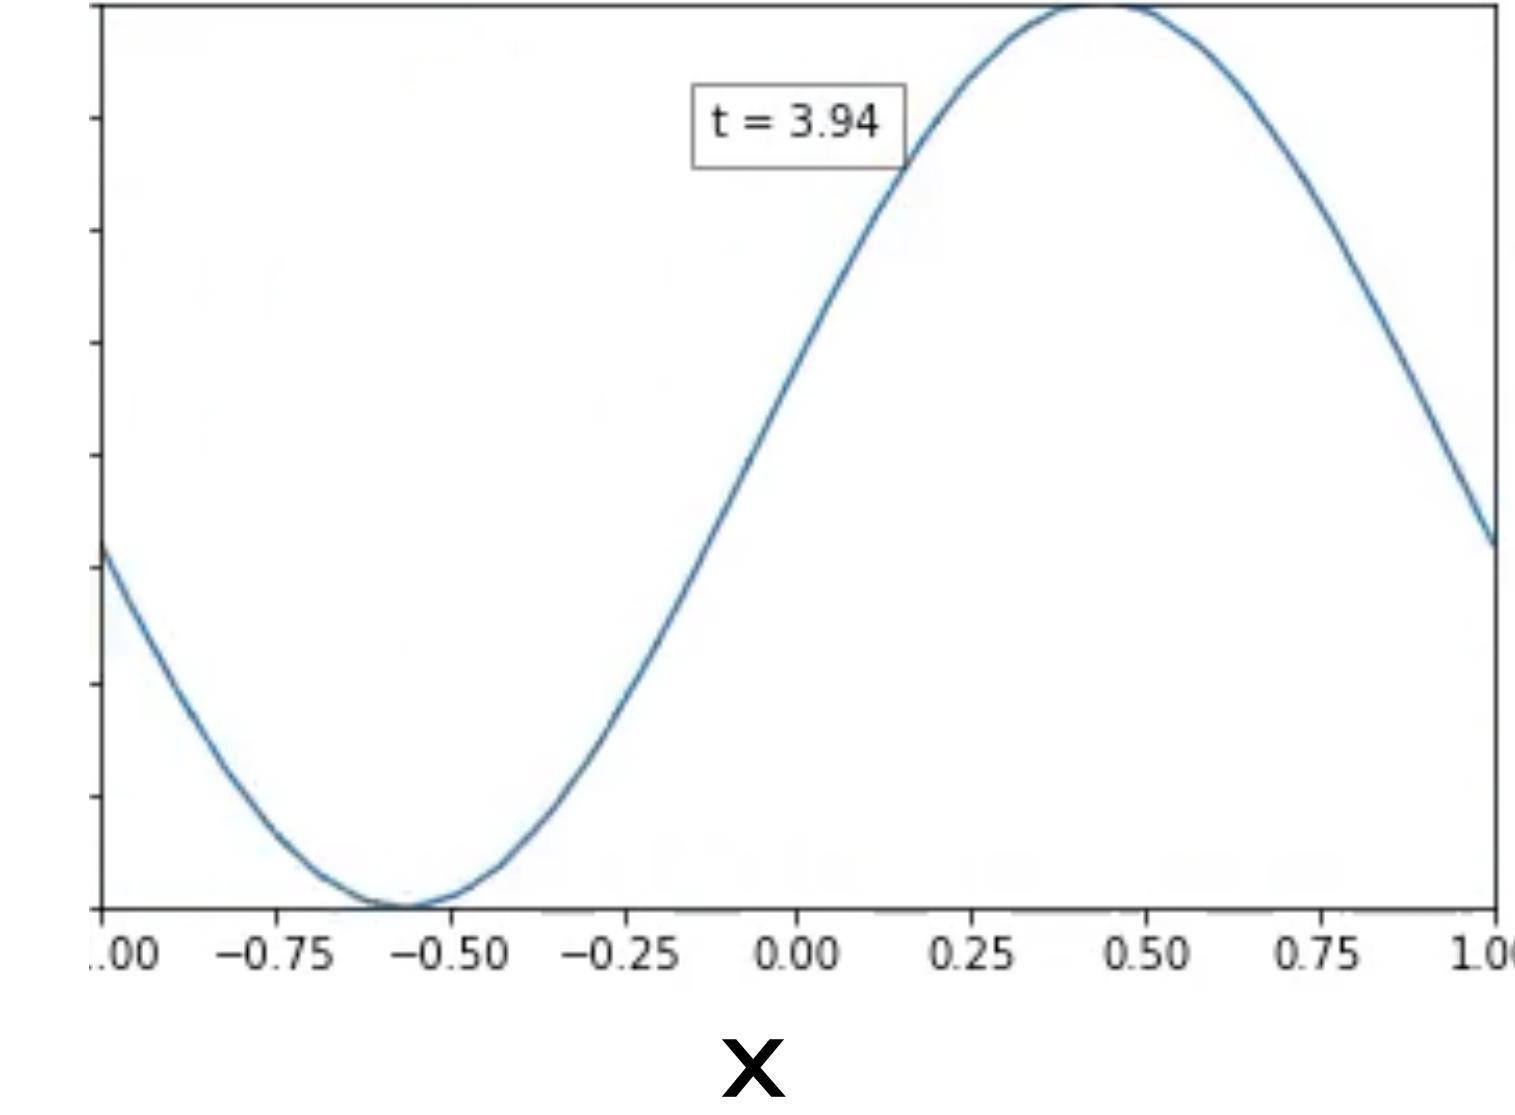
(Machine) Learning to Differentiate

Discretization for Linear Advection
 $N=15$ Mesh Elements

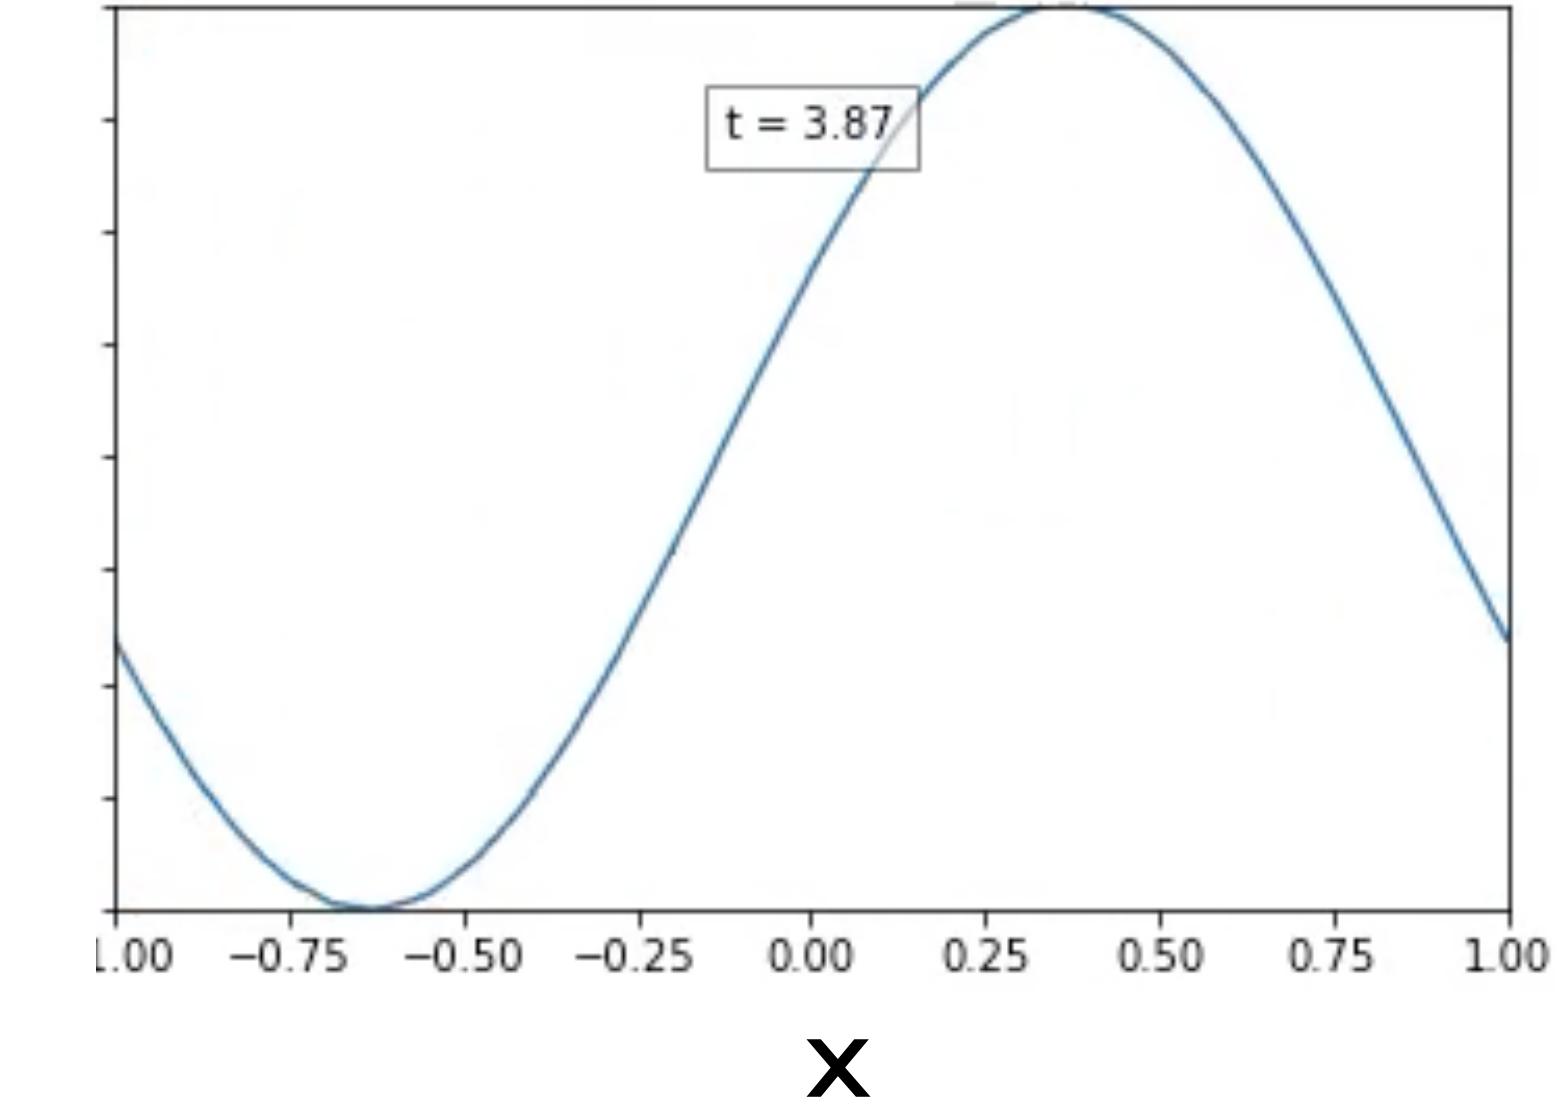
No Constraints



SBP Constraint via
Loss Function



SBP Constraint via
Network Structure

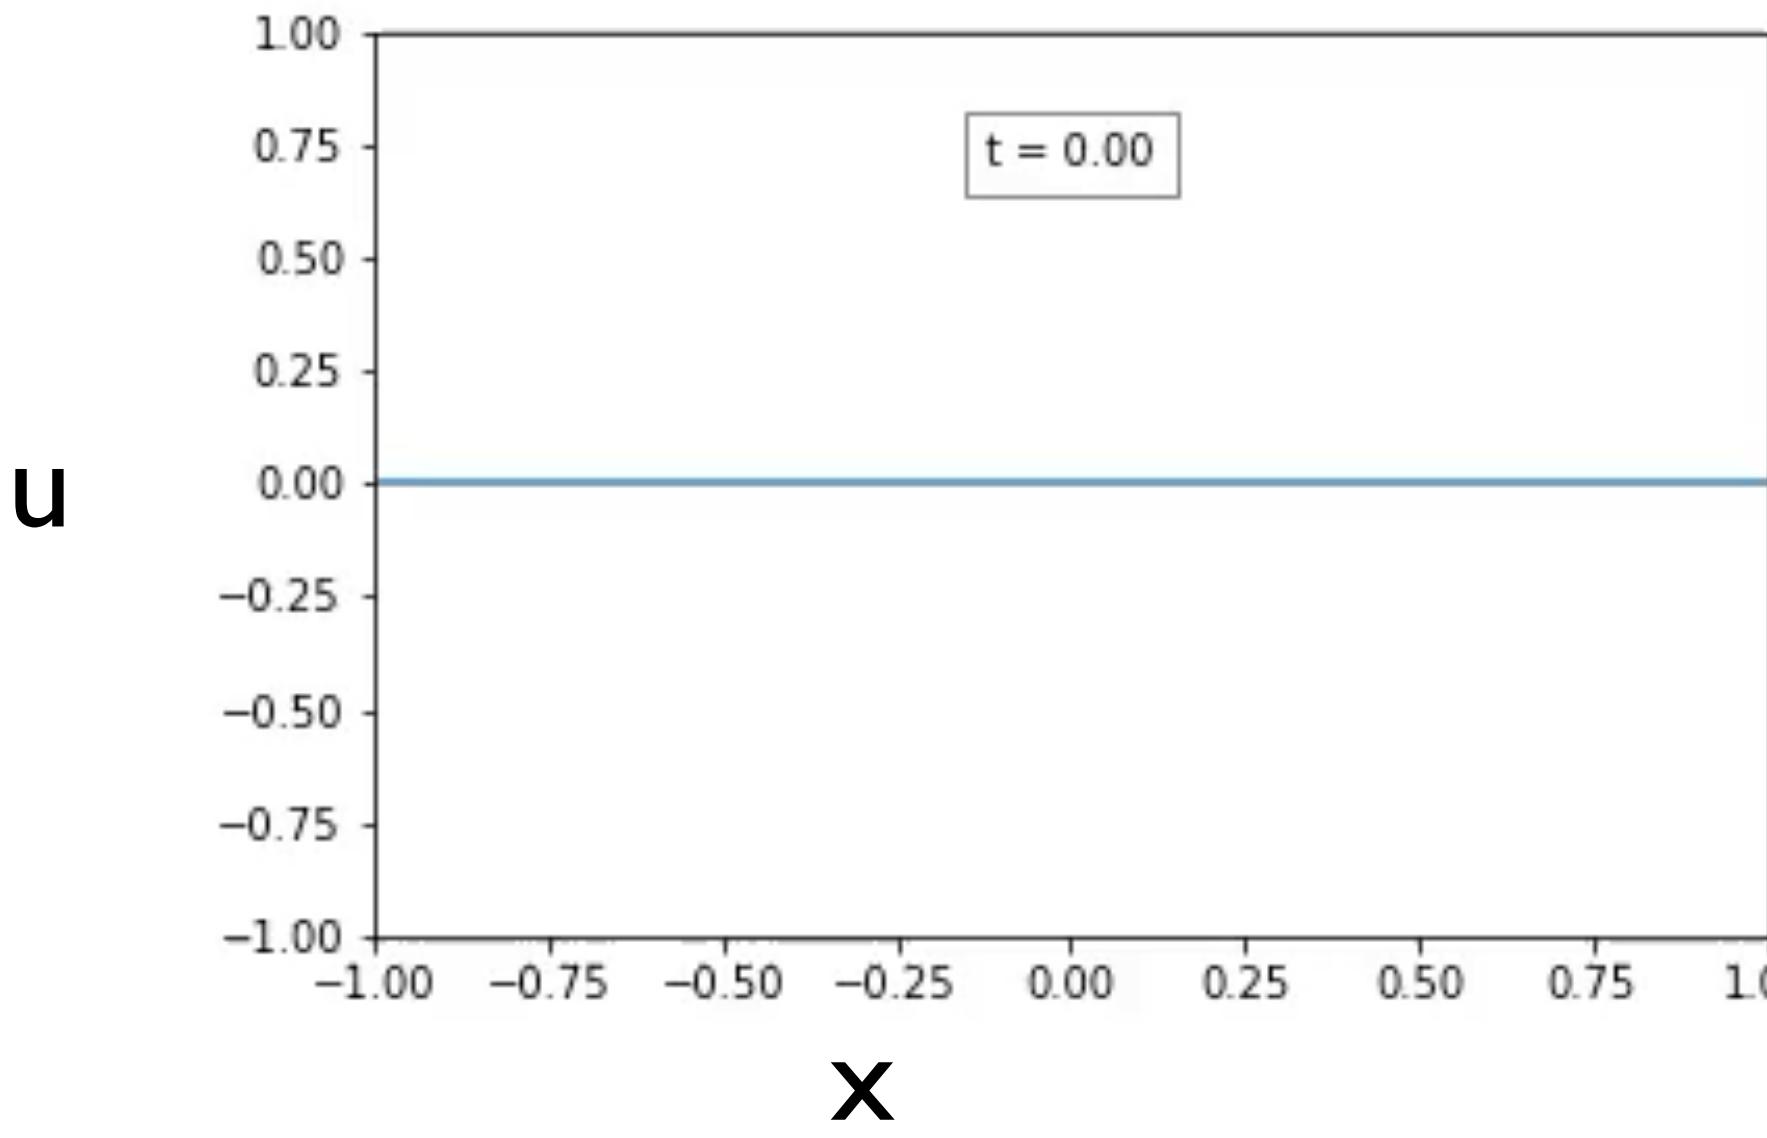




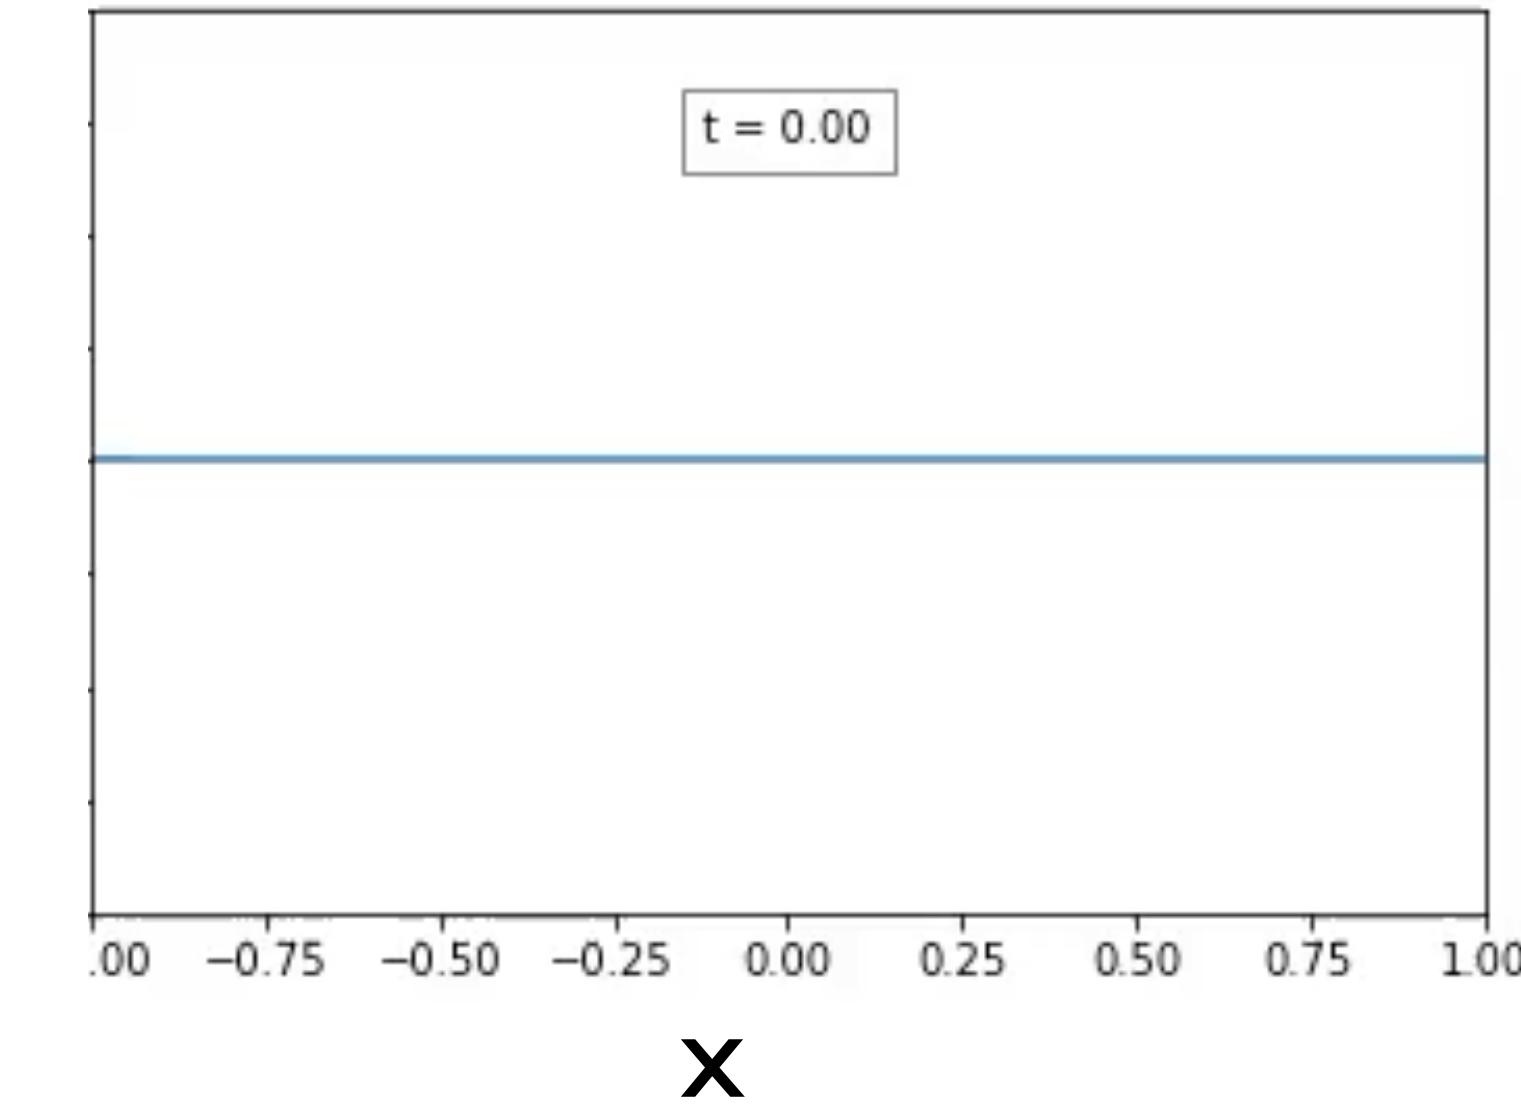
(Machine) Learning to Differentiate

Discretization for Linear Advection
N=50 Mesh Elements

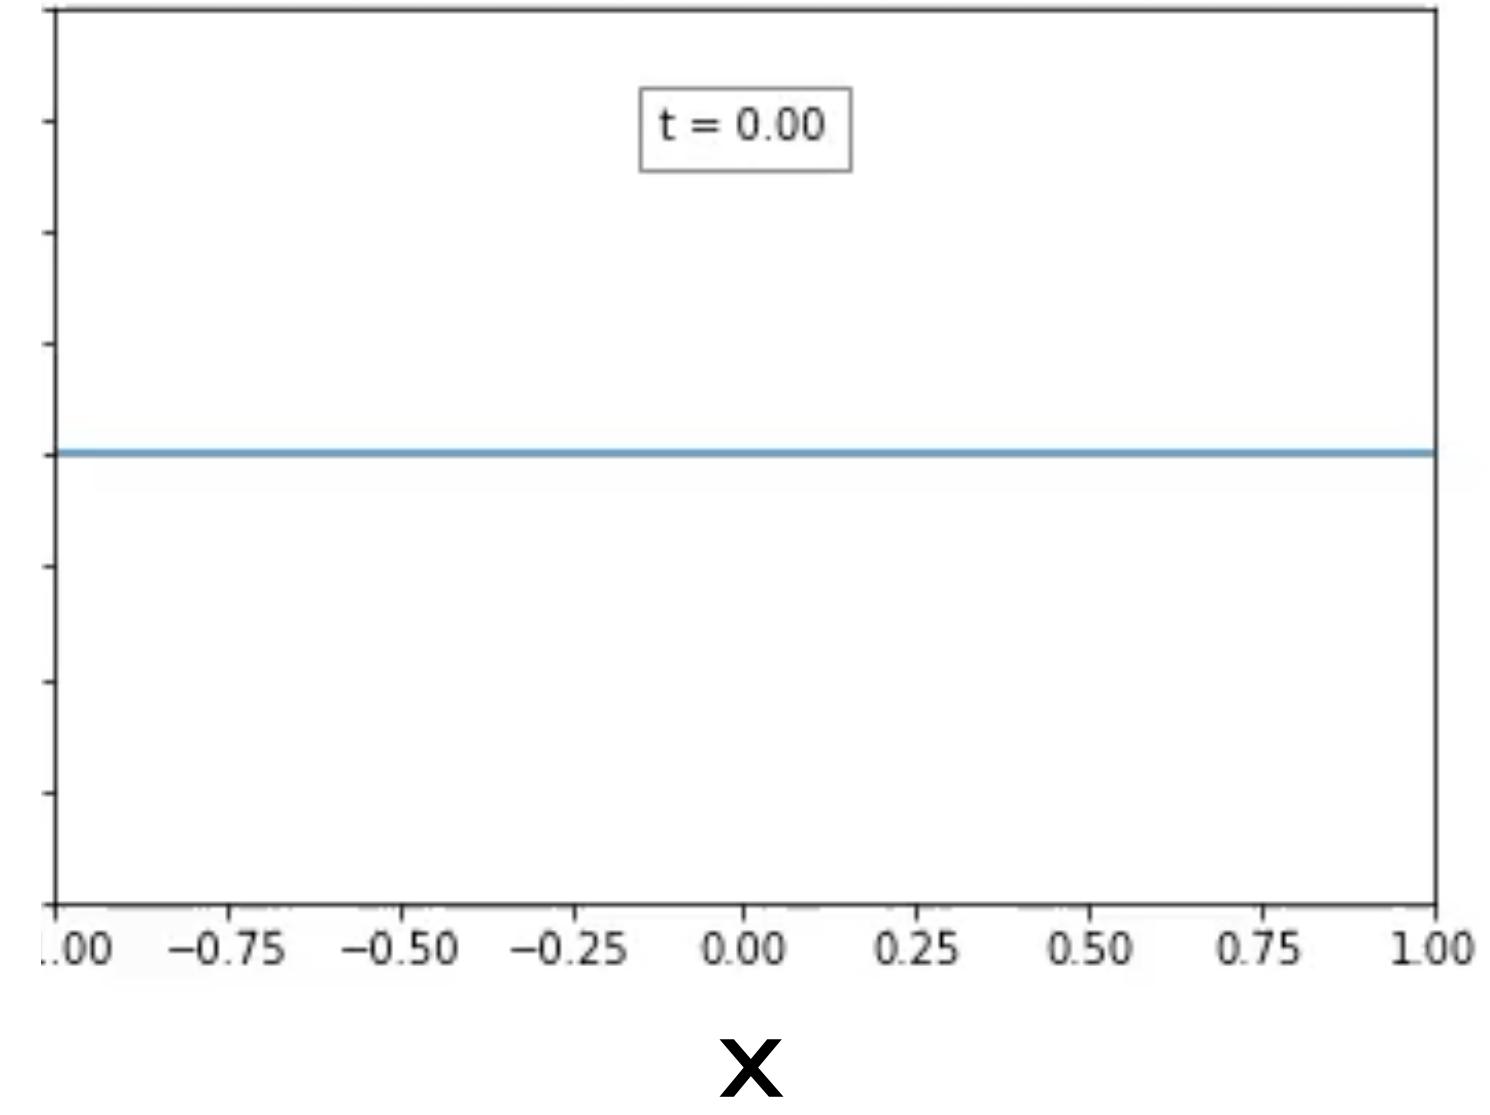
No Constraints



SBP Constraint via
Loss Function



SBP Constraint via
Network Structure



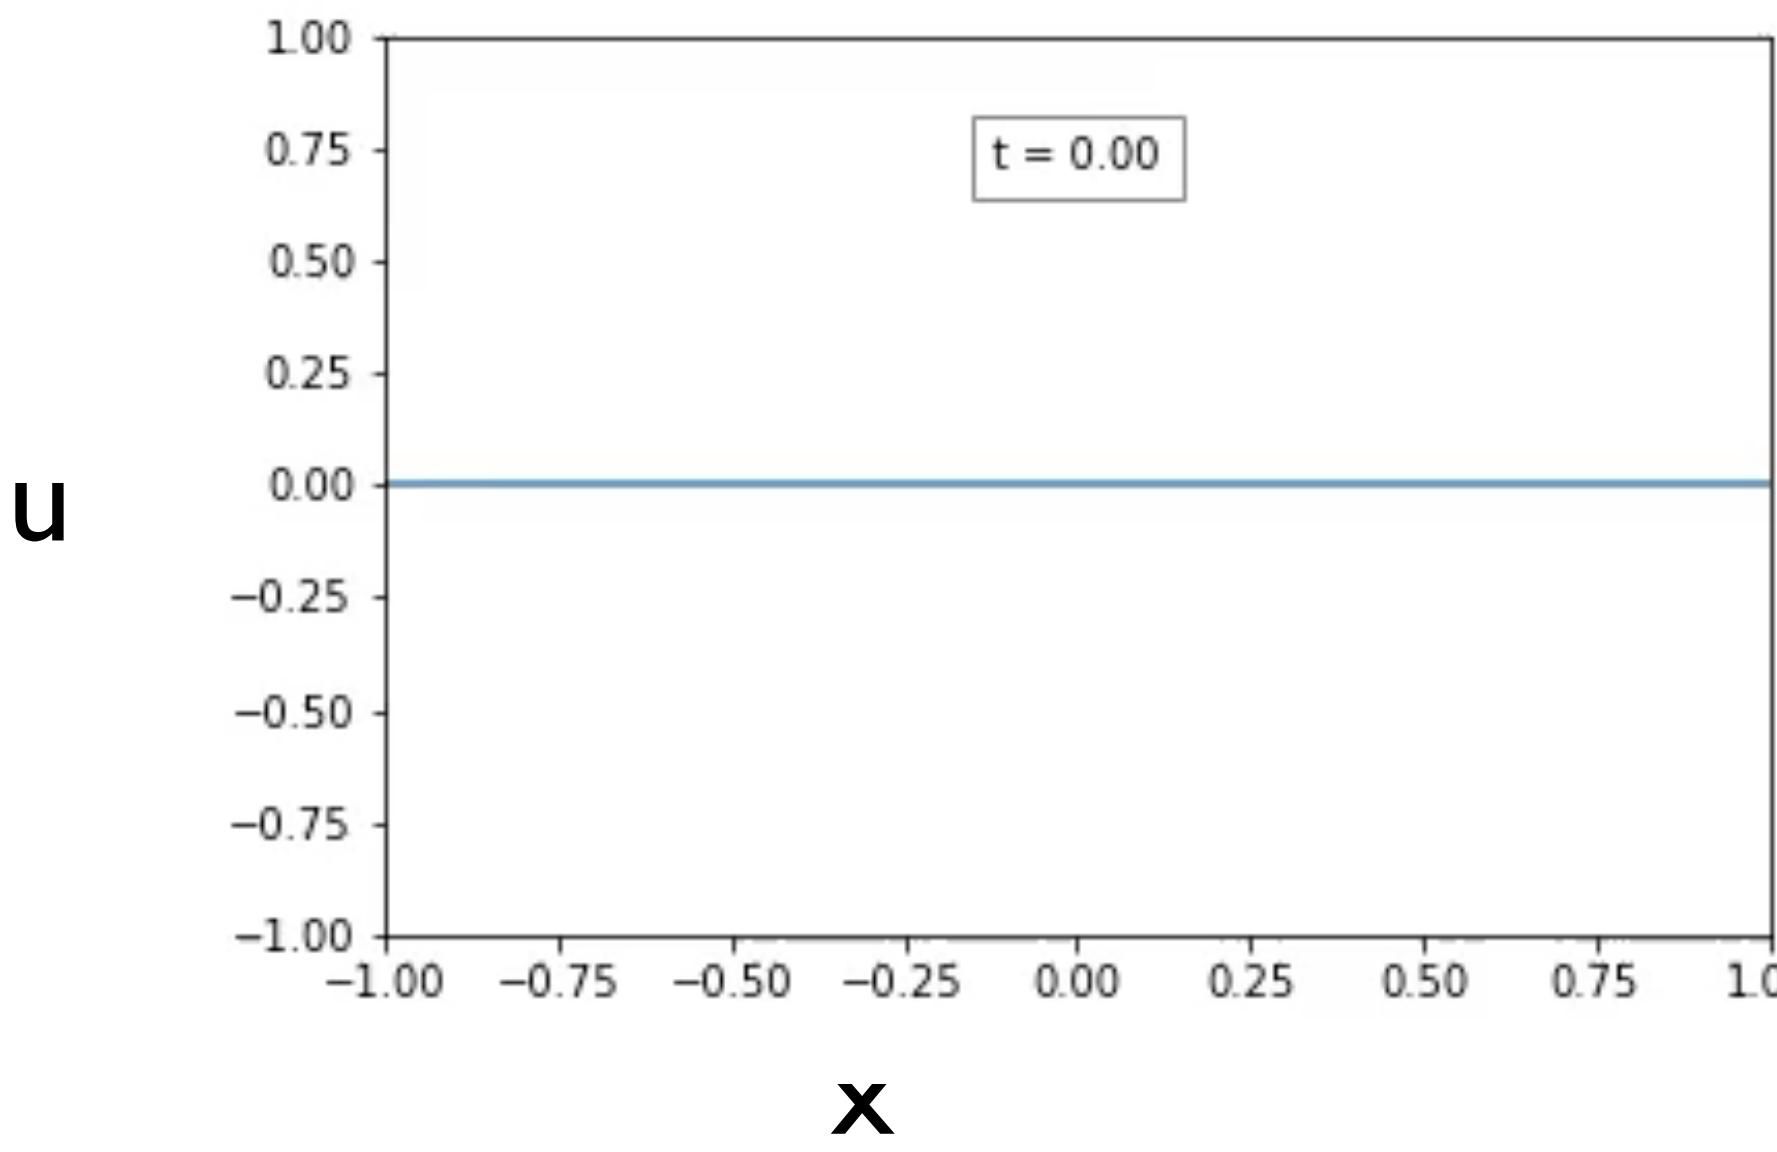


(Machine) Learning to Differentiate

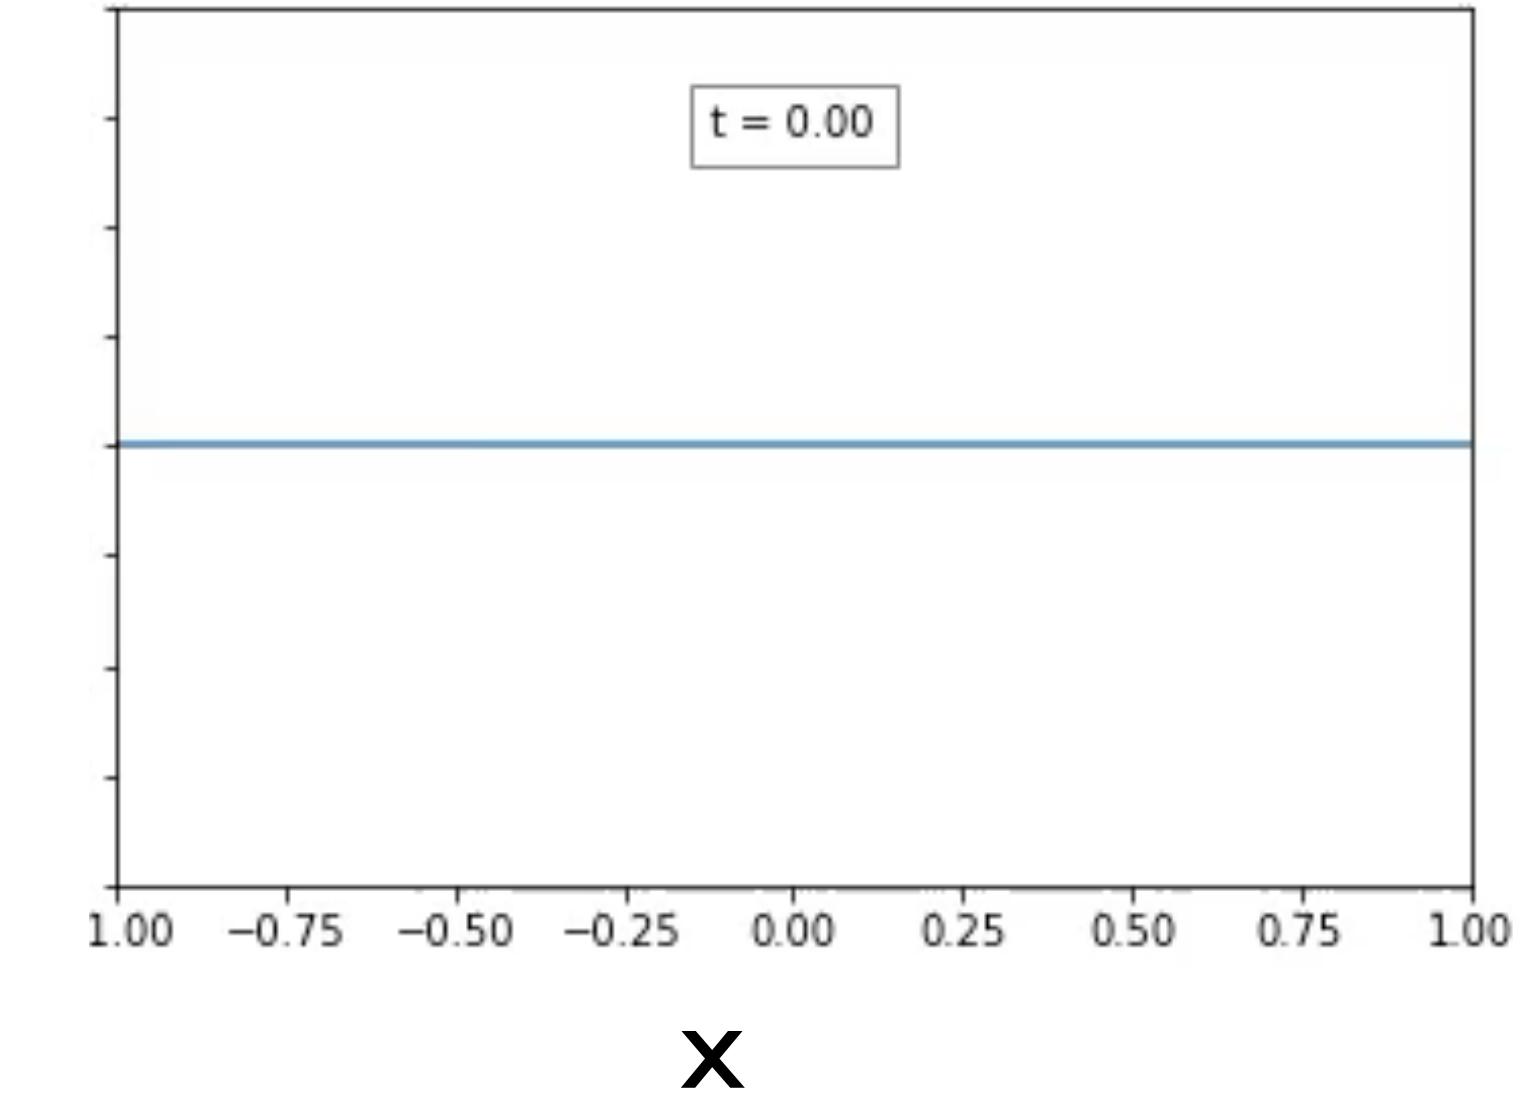


Discretization for Linear Advection - training with noisy functions
N=15 Mesh Elements

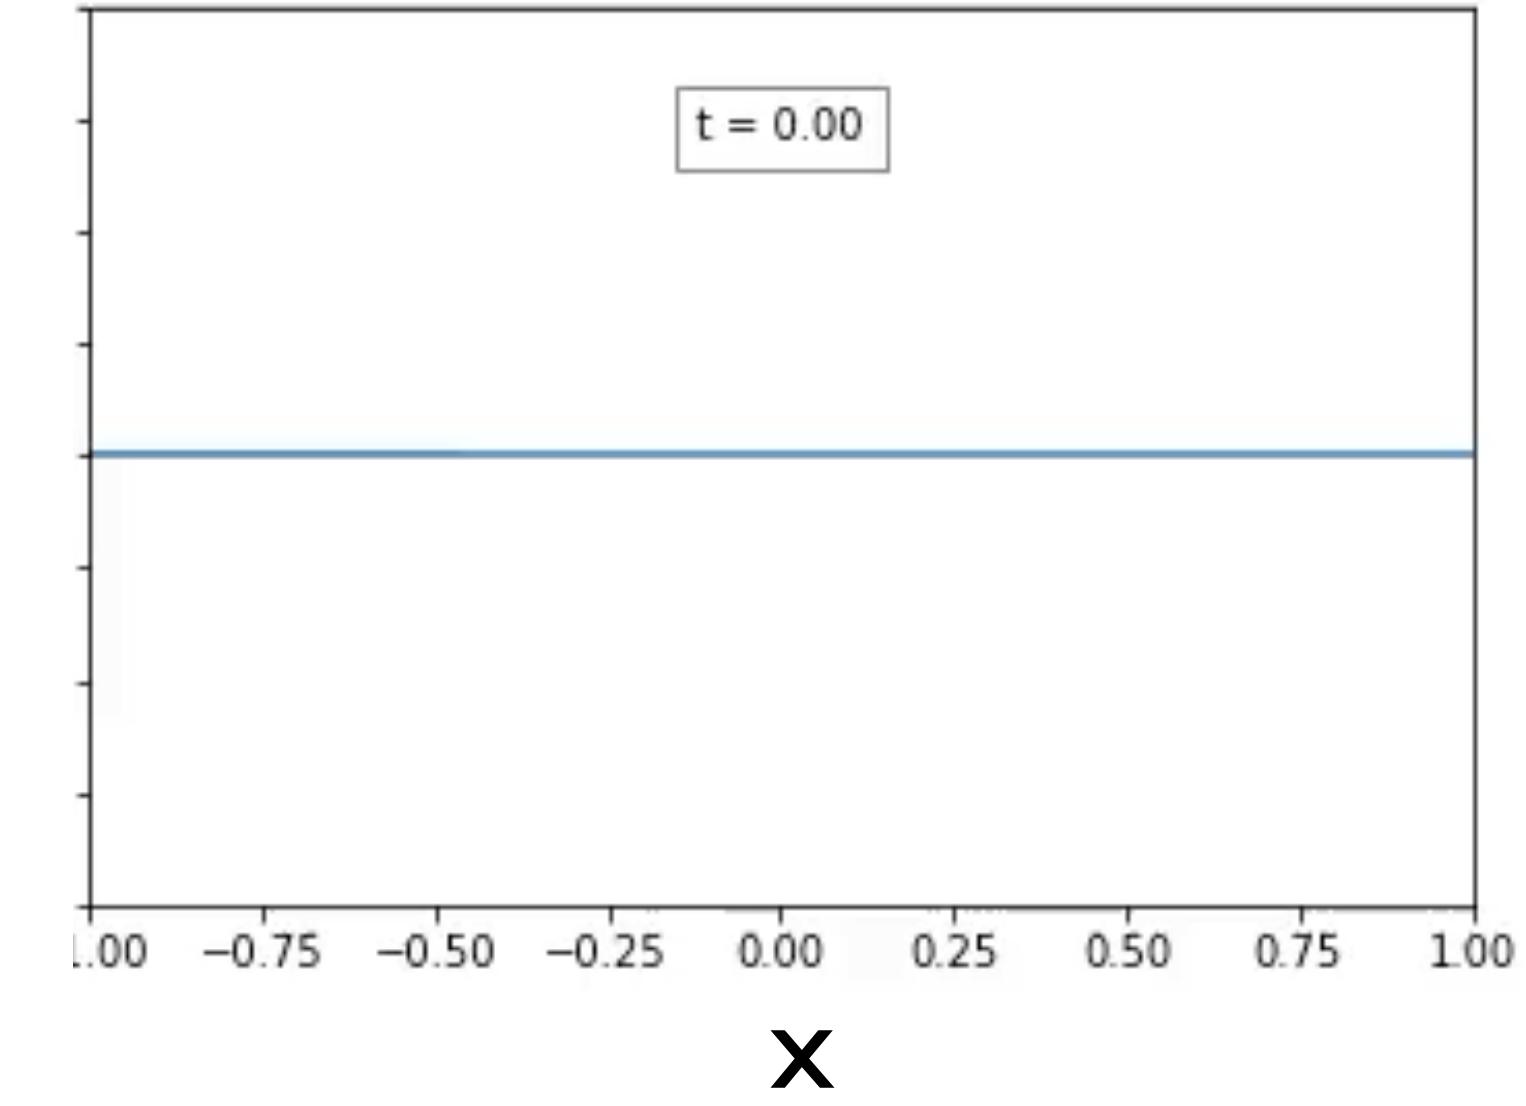
No Constraints



SBP Constraint via
Loss Function



SBP Constraint via
Network Structure





(Machine) Learning to Differentiate

Discretization for Linear Advection

No Constraints

unstable

SBP Constraint via
Loss Function

Regularization		
N	error	rate
10	4.3548e-02	-
20	4.4266e-03	3.30
40	8.5910e-04	2.37
80	2.0676e-04	2.05
160	4.4616e-05	2.21

SBP Constraint via
Network Structure

Weight restriction		
N	error	rate
10	3.7098e-02	-
20	3.4367e-03	3.43
40	3.8367e-04	3.16
80	4.6386e-05	3.05
160	5.7488e-06	3.01



A Numerical Technique for PDEs

Consider **non-linear transport** $u_t + uu_x = 0, \quad x \in (-1, 1), \quad t > 0$

building $\mathbf{D} = \mathbf{P}^{-1}\mathbf{Q}$ using ML we cannot ensure stability!

However we can use the transformation $u_t + \frac{1}{3}(uu_x + (u^2)_x) = 0,$

and the operator $\mathcal{L}(\mathbf{u}) = \frac{1}{3}\mathbf{P}^{-1}\mathbf{Q}(\mathbf{u})$ with $\mathbf{U} = \text{diag}(\mathbf{u})$ and $\mathbf{Q}(\mathbf{u}) = \mathbf{U}\mathbf{Q} + \mathbf{Q}\mathbf{U}$

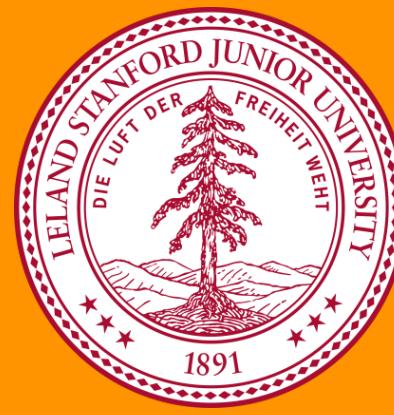
$$\mathbf{u}_t^{(1)} + \mathcal{L}(\mathbf{u}^{(1)})\mathbf{u}^{(1)} = \sigma_1 \mathbf{P}^{-1}(u_0^{(1)} - g)\mathbf{e}_0$$

$$\mathbf{u}_t^{(2)} + \mathcal{L}(\mathbf{u}^{(2)})\mathbf{u}^{(2)} = \sigma_2 \mathbf{P}^{-1}(u_0^{(2)} - u_2^{(1)})\mathbf{e}_0$$

⋮

$$\mathbf{u}_t^{(N)} + \mathcal{L}(\mathbf{u}^{(N)})\mathbf{u}^{(N)} = \sigma_N \mathbf{P}^{-1}(u_0^{(N)} - u_2^{(N-1)})\mathbf{e}_0.$$

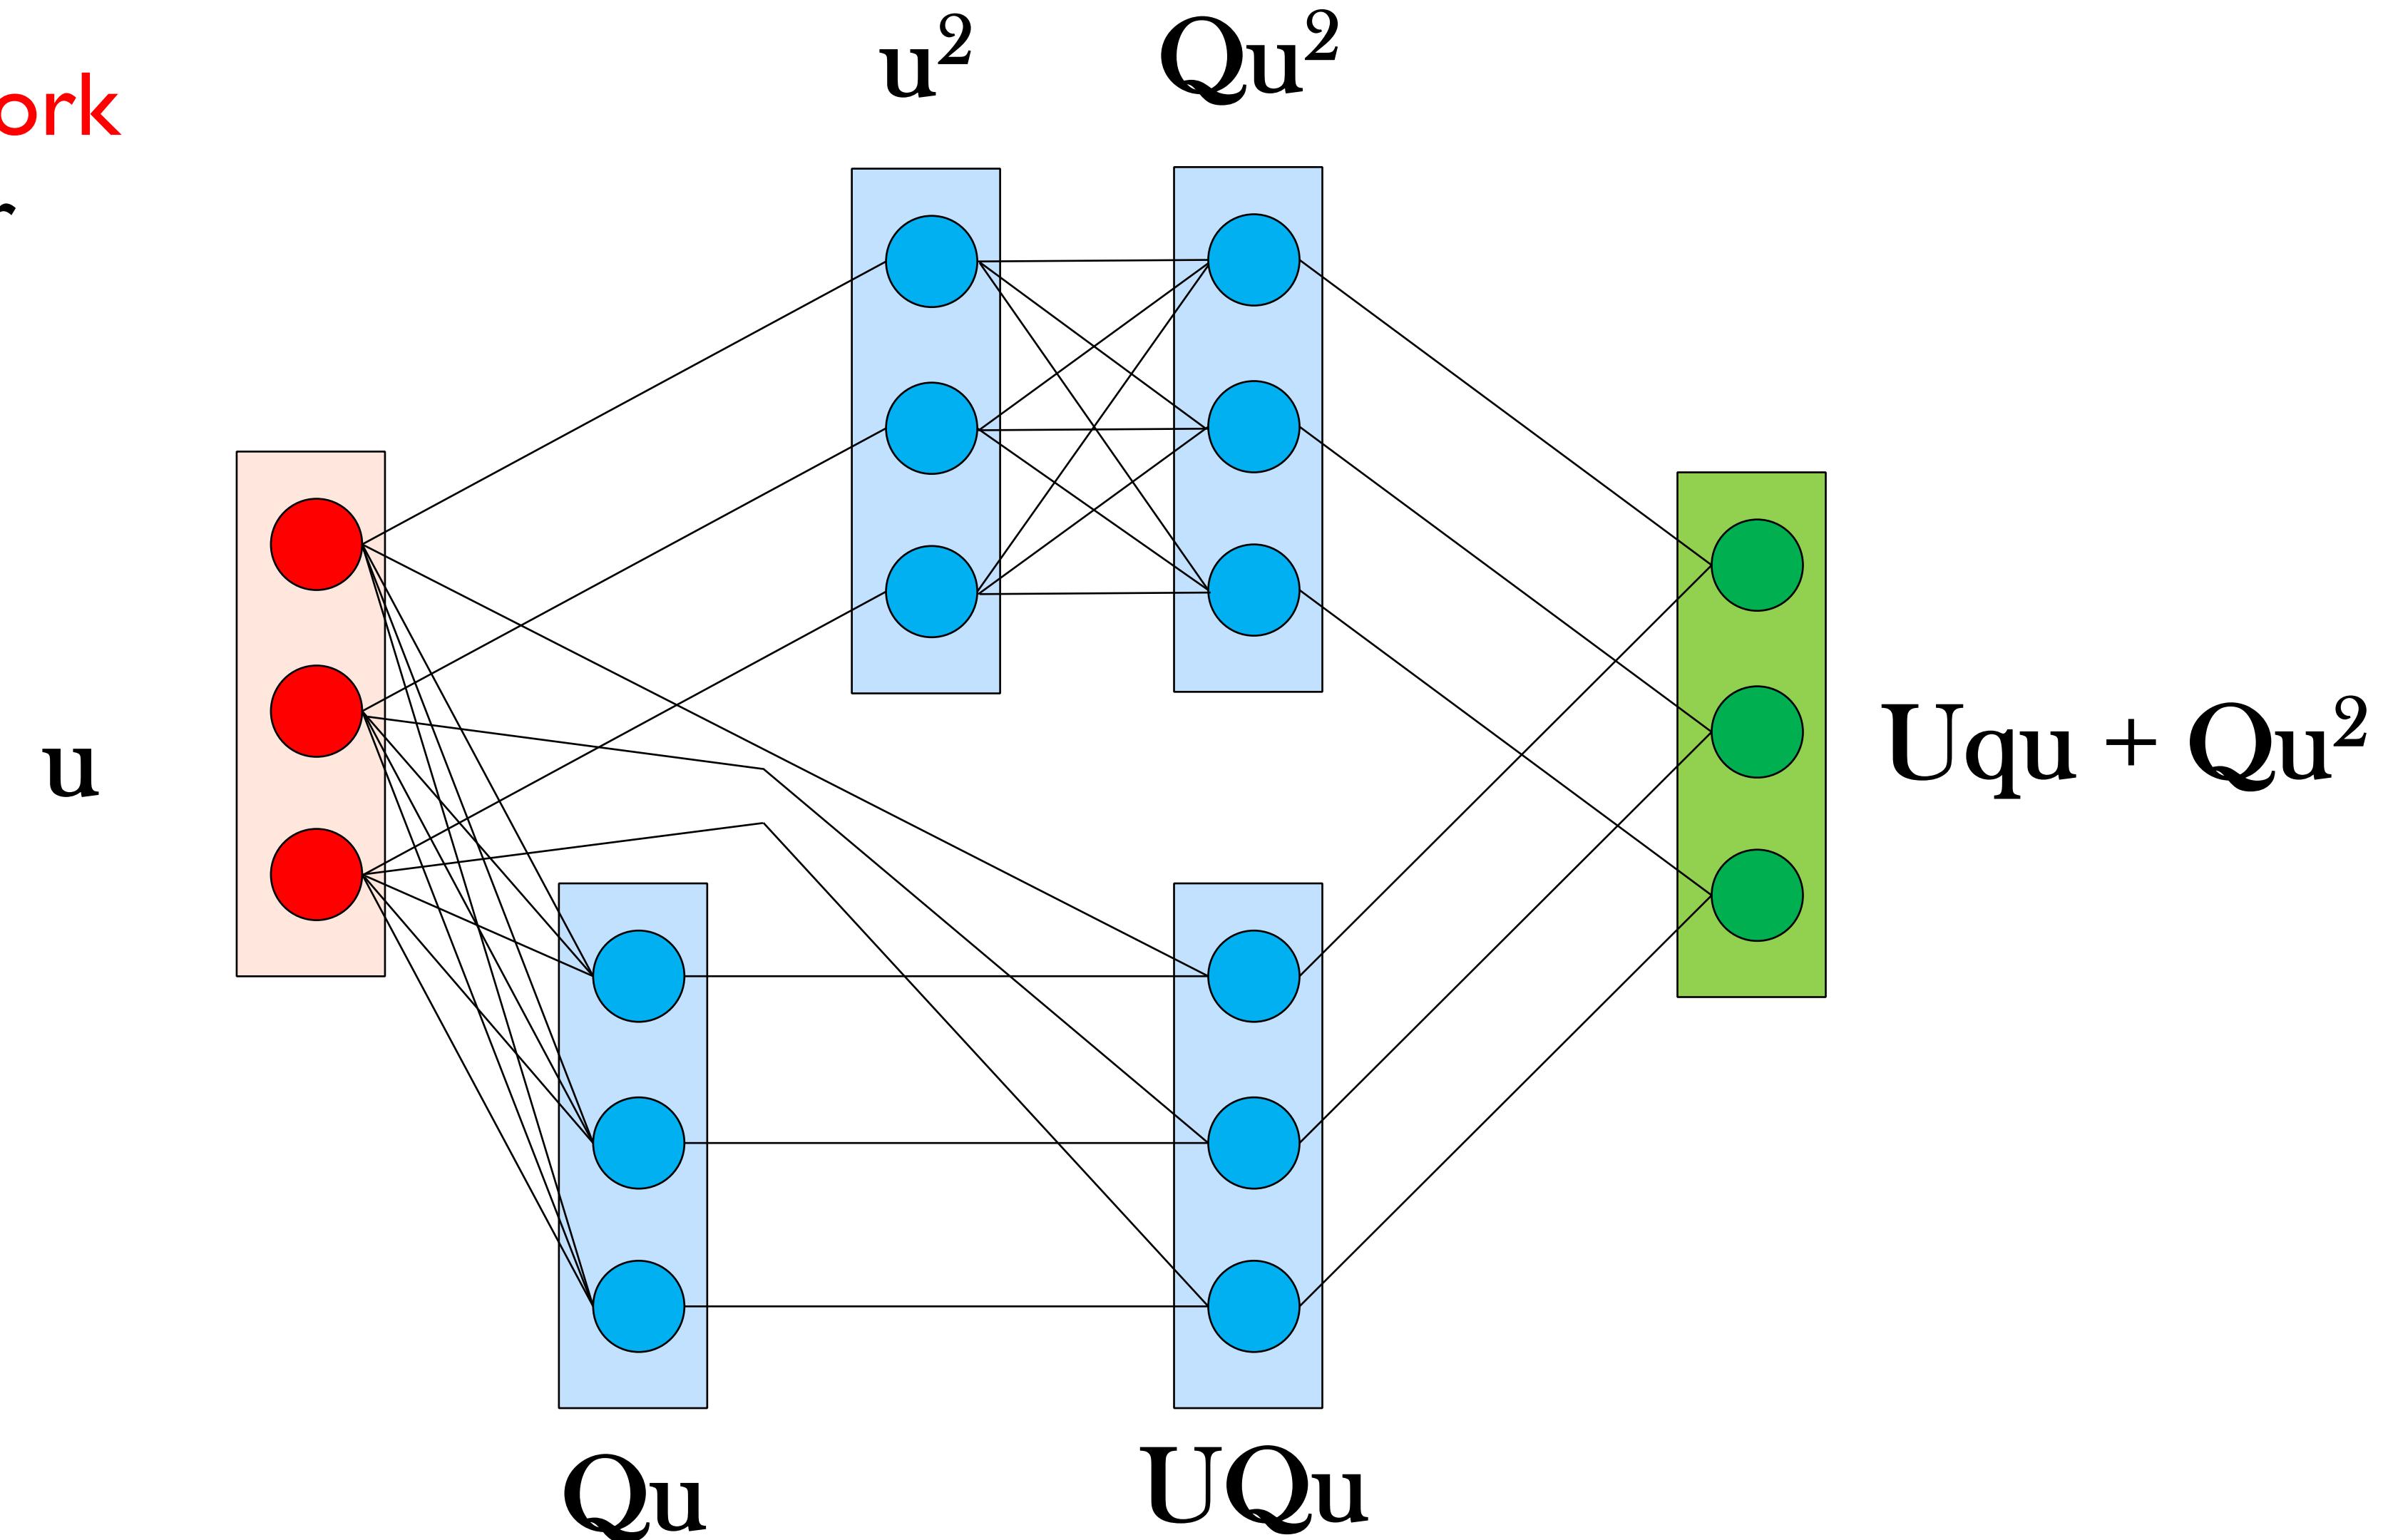
...similar to the linear case
we can prove **stability!**



A Numerical Technique for PDEs

We can build a **new network**
to construct the operator

$$\mathcal{L}(u) = \frac{1}{3}P^{-1}\mathcal{Q}(u)$$





(Machine) Learning to Differentiate



Discretization for **non-Linear** Advection

No Constraints

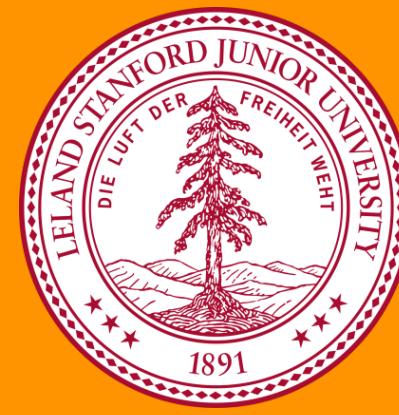
unstable

SBP Constraint via
Loss Function

Regularization		
N	error	rate
10	1.8121e-02	-
20	3.0915e-03	2.55
40	4.0387e-04	2.94
80	6.0469e-05	2.74
160	1.0986e-05	2.46

SBP Constraint via
Network Structure

Weight restriction		
N	error	rate
10	2.6565e-02	-
20	2.4962e-03	3.41
40	2.7021e-04	3.21
80	3.2516e-05	3.05
160	4.1301e-06	2.98



Summary & Conclusions

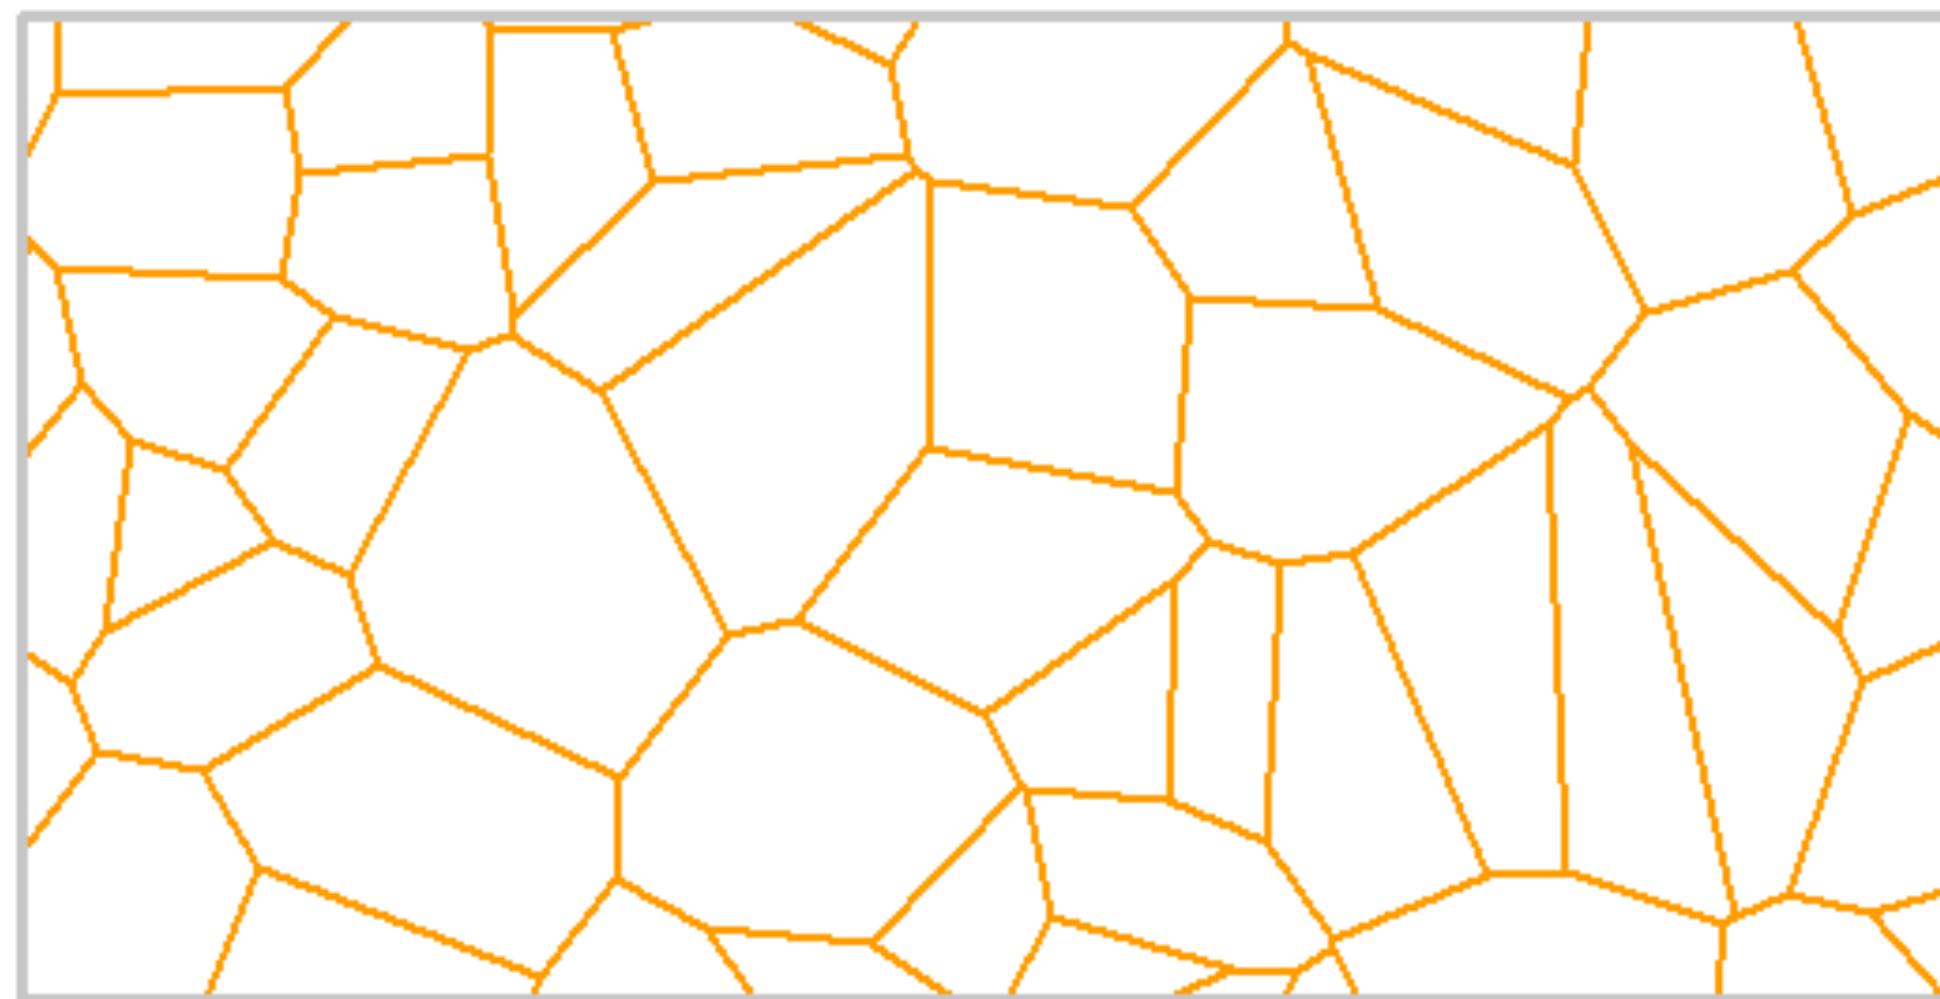


- We built discretization schemes (nodal discontinuous Galerkin) using neural networks
- We enforced stability constraints using the SBP property
- We investigated the possibility of enforcing (physical) properties either through regularization (loss function) or via network constraints (weights)
- We constructed a specific network architecture for non-linear operators
- We observe that training with specific function sets (polynomials) accelerates learning

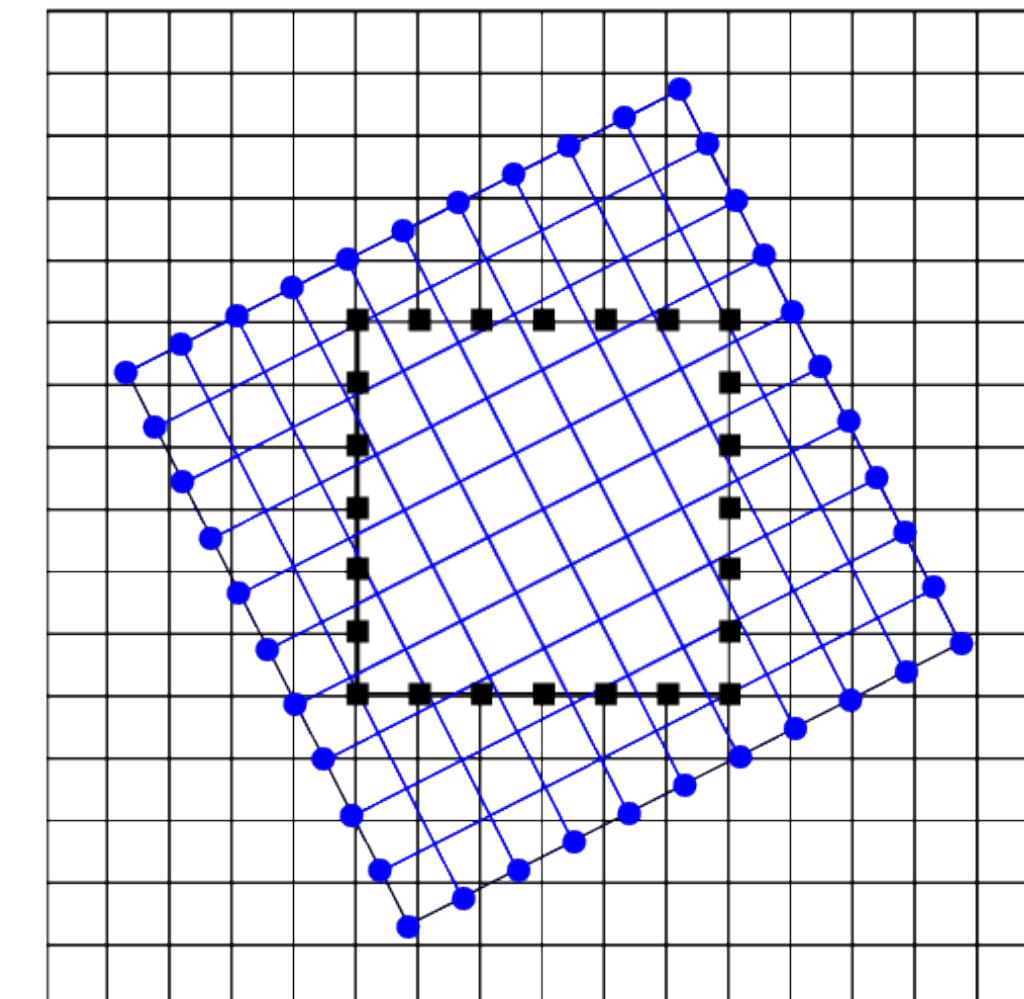


What's Next?

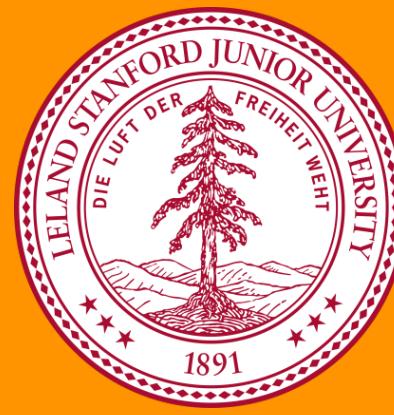
Genuine Multi-dimensional and Provably Stable Discretization Schemes for



Unstructured Polyhedral Grids



Irregular/Overlapping Grids



Final Thoughts



Data-driven techniques offer great potential, and we are just at the beginning...

Great diversity in techniques and objectives create difficulties in assessing where we are and how far we are moving

Low-hanging fruits: ML techniques are powerful generator of heuristics

Long term: need to consider different paradigm that blend numerical techniques and physics modeling to take full advantage of both HPC and ML



Questions?



thank you

Ali Kashefi, Aashwin Mishra (Part I & 2)

Oskar Ålund, Gianluca Iaccarino, Jan Nordström,
“Learning to Differentiate”, submitted to JCP 2020

